

# AI Robotics

## Perception



MISSISSIPPI STATE  
UNIVERSITY™

# Overview



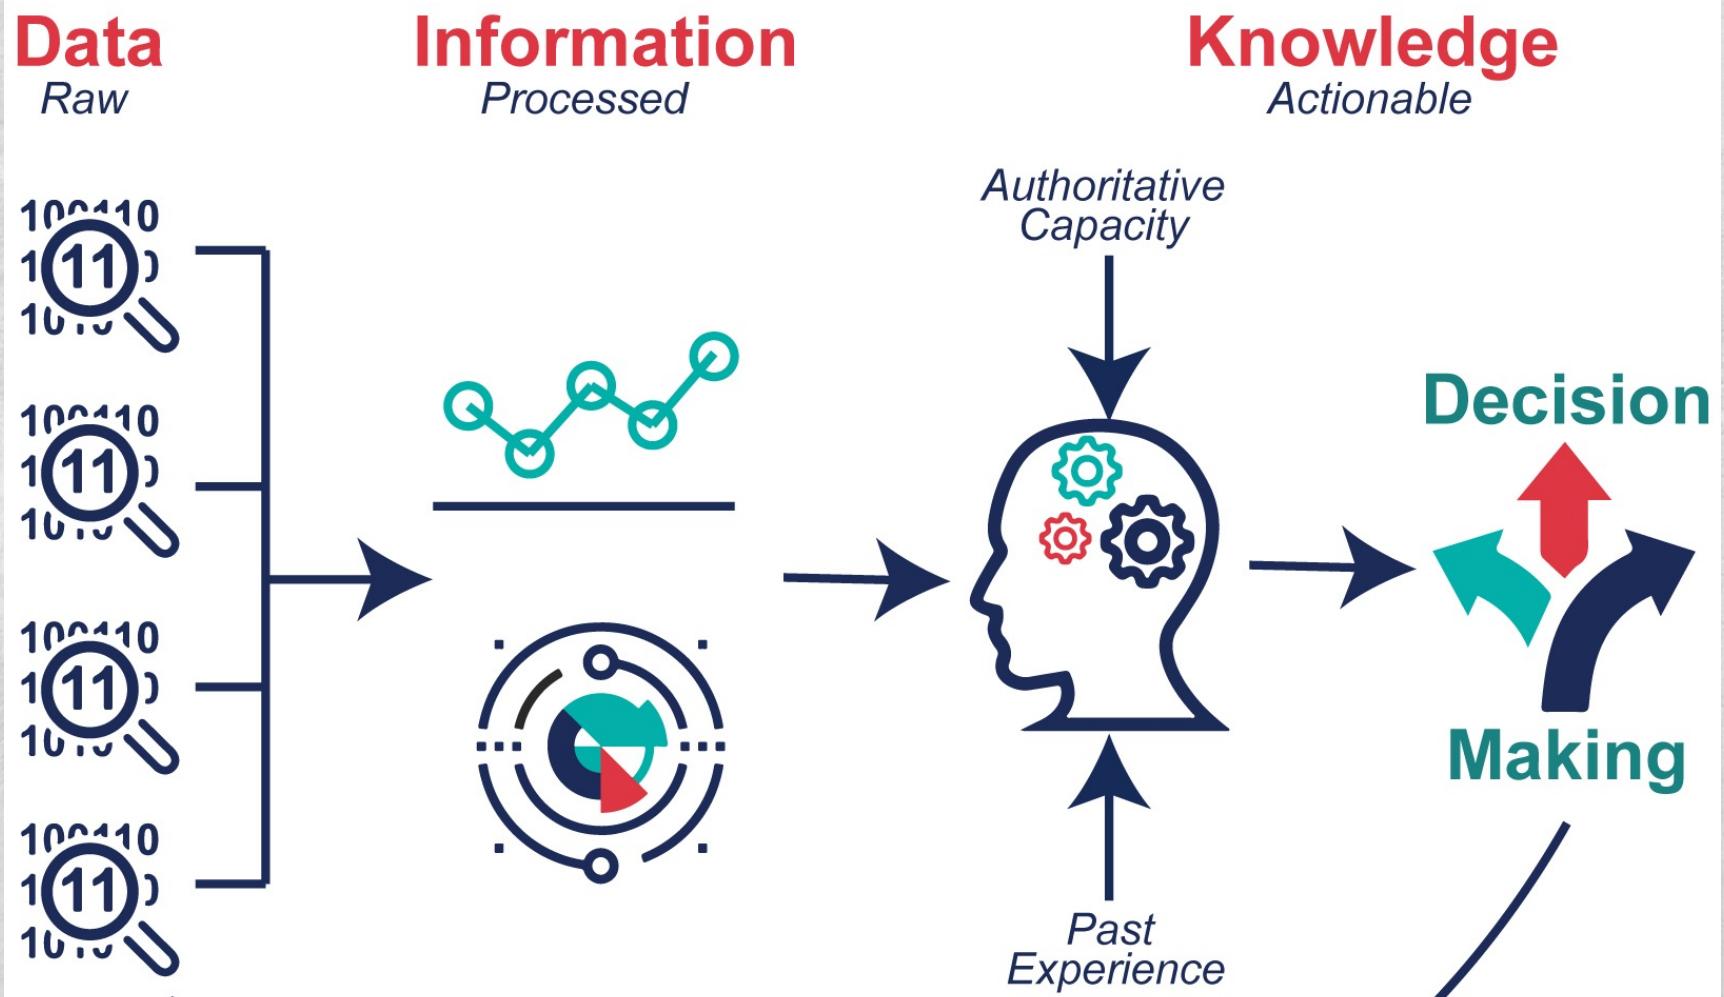
MISSISSIPPI STATE  
UNIVERSITY™

# Perception

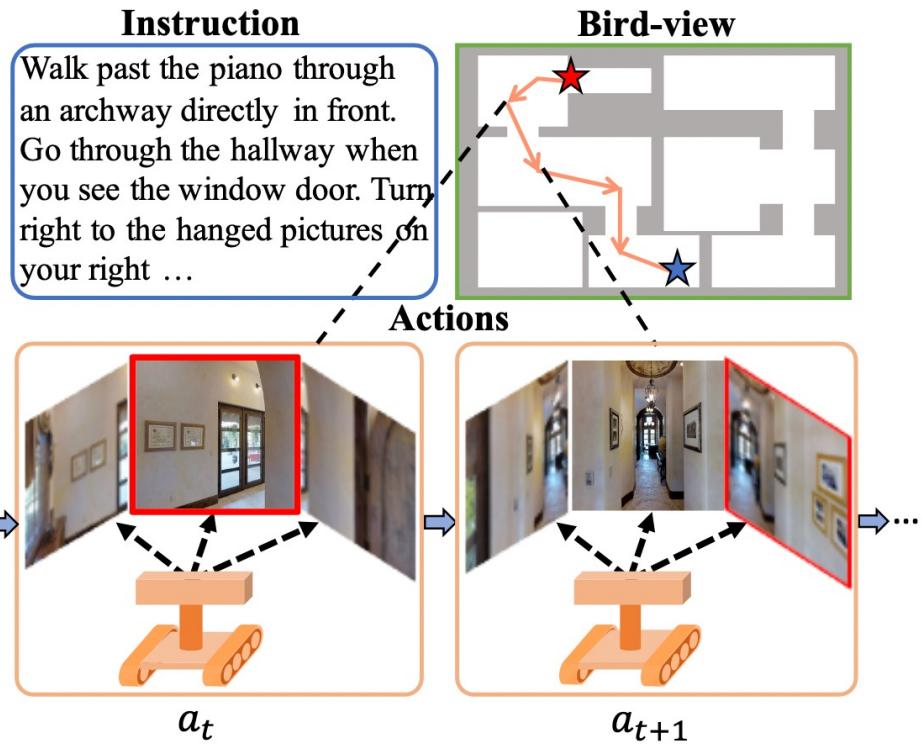
For robots to act and react more richly with the world around them, they require a deeper perception and understanding of the environment, often referred to as semantic understanding, of the world in which they operate



# Perception



# Perception in Indoor Navigation



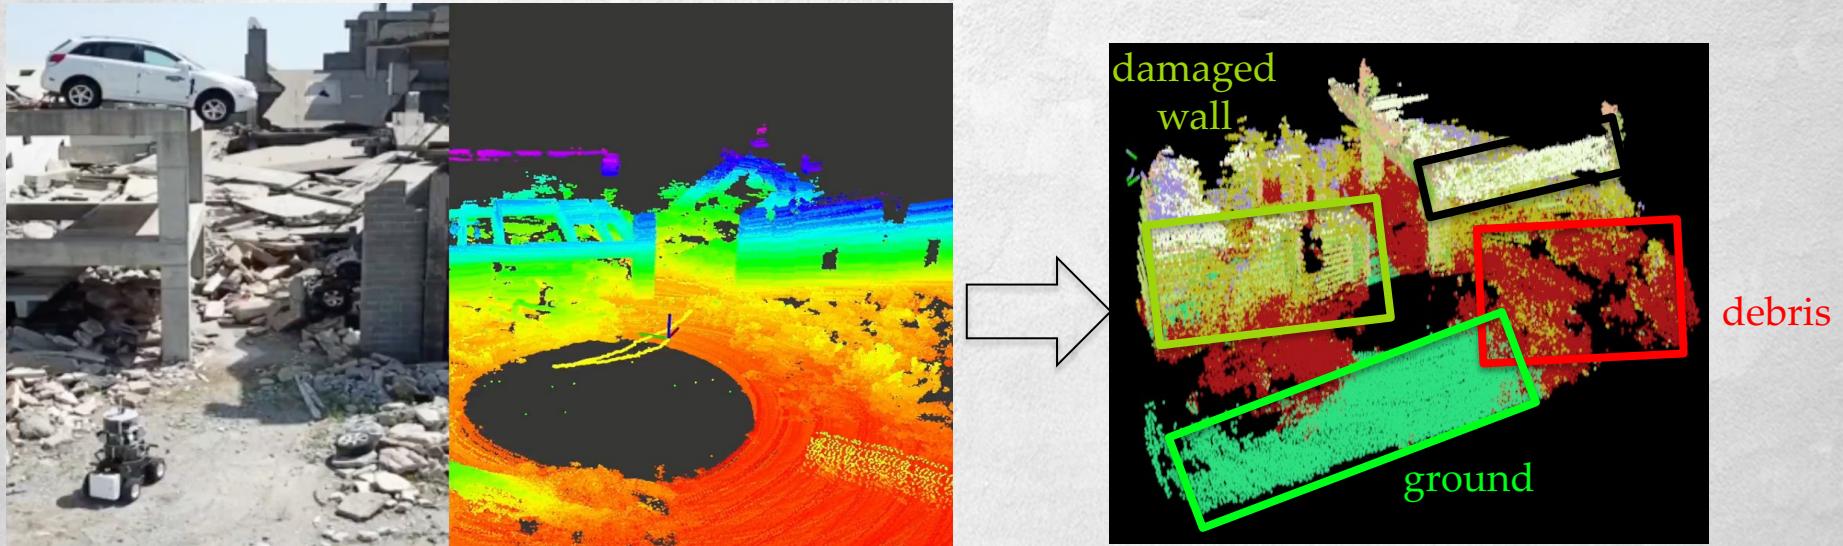
Exit the bedroom and go towards the table. Go to the stairs on the left of the couch. Wait on the third step.



MISSISSIPPI STATE  
UNIVERSITY™

# Perception in Disaster Relief

Improve the ability of mobile robots to carry out post-disaster reconnaissance tasks through automated segmentation and damage detection from acquired point cloud data



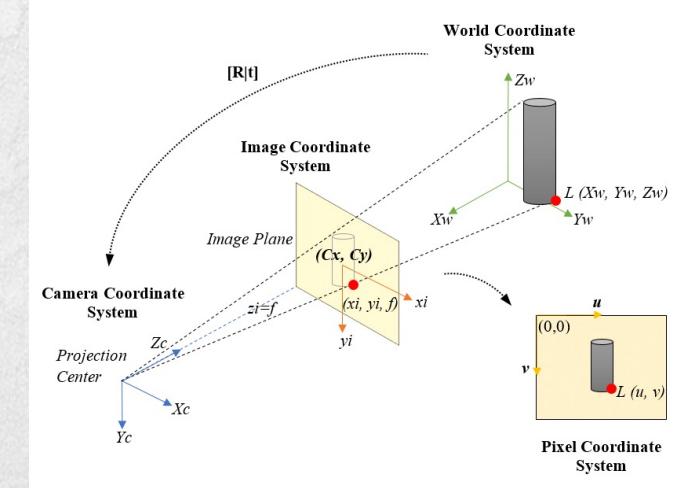
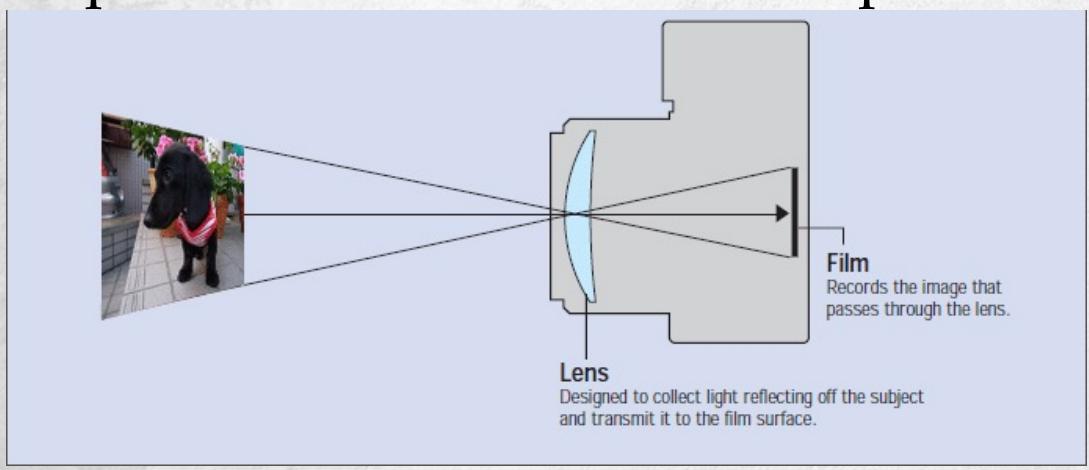
# 2D Perception



MISSISSIPPI STATE  
UNIVERSITY™

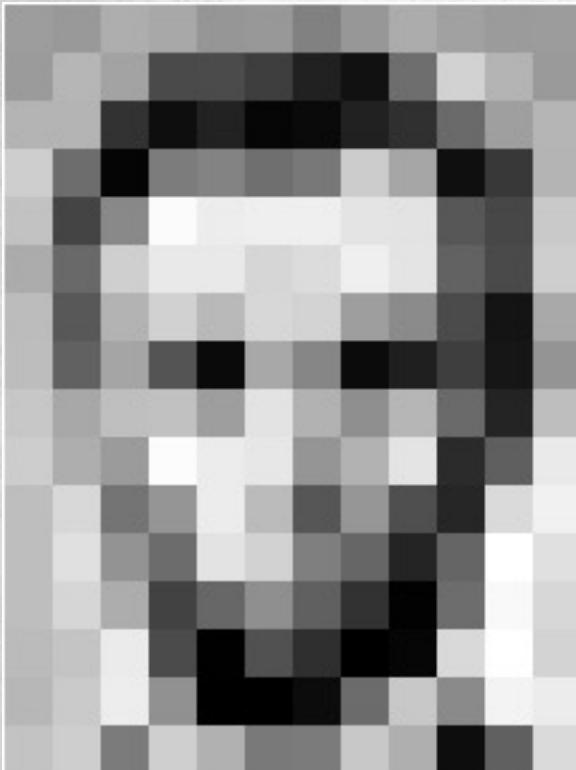
# Camera

- Uses Charge-Coupled Device (CCD) sensors to measure light intensity from the environment
- Lenses focus the light entering the camera and the size of the aperture can be widened or narrowed to let more or less light into the camera
- Shutter mechanism determines the amount of time the photosensitive surface is exposed to the light.



MISSISSIPPI STATE  
UNIVERSITY™

# Image Data Representation



157	153	174	168	150	152	129	151	172	161	155	166
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	94	6	10	33	48	106	169	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	209	175	13	96	218

157	153	174	168	150	152	129	151	172	161	155	166
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	94	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	209	175	13	96	218



MISSISSIPPI STATE  
UNIVERSITY™

# Image Scaling



Scaled down  
to 30%



Original  
raster image

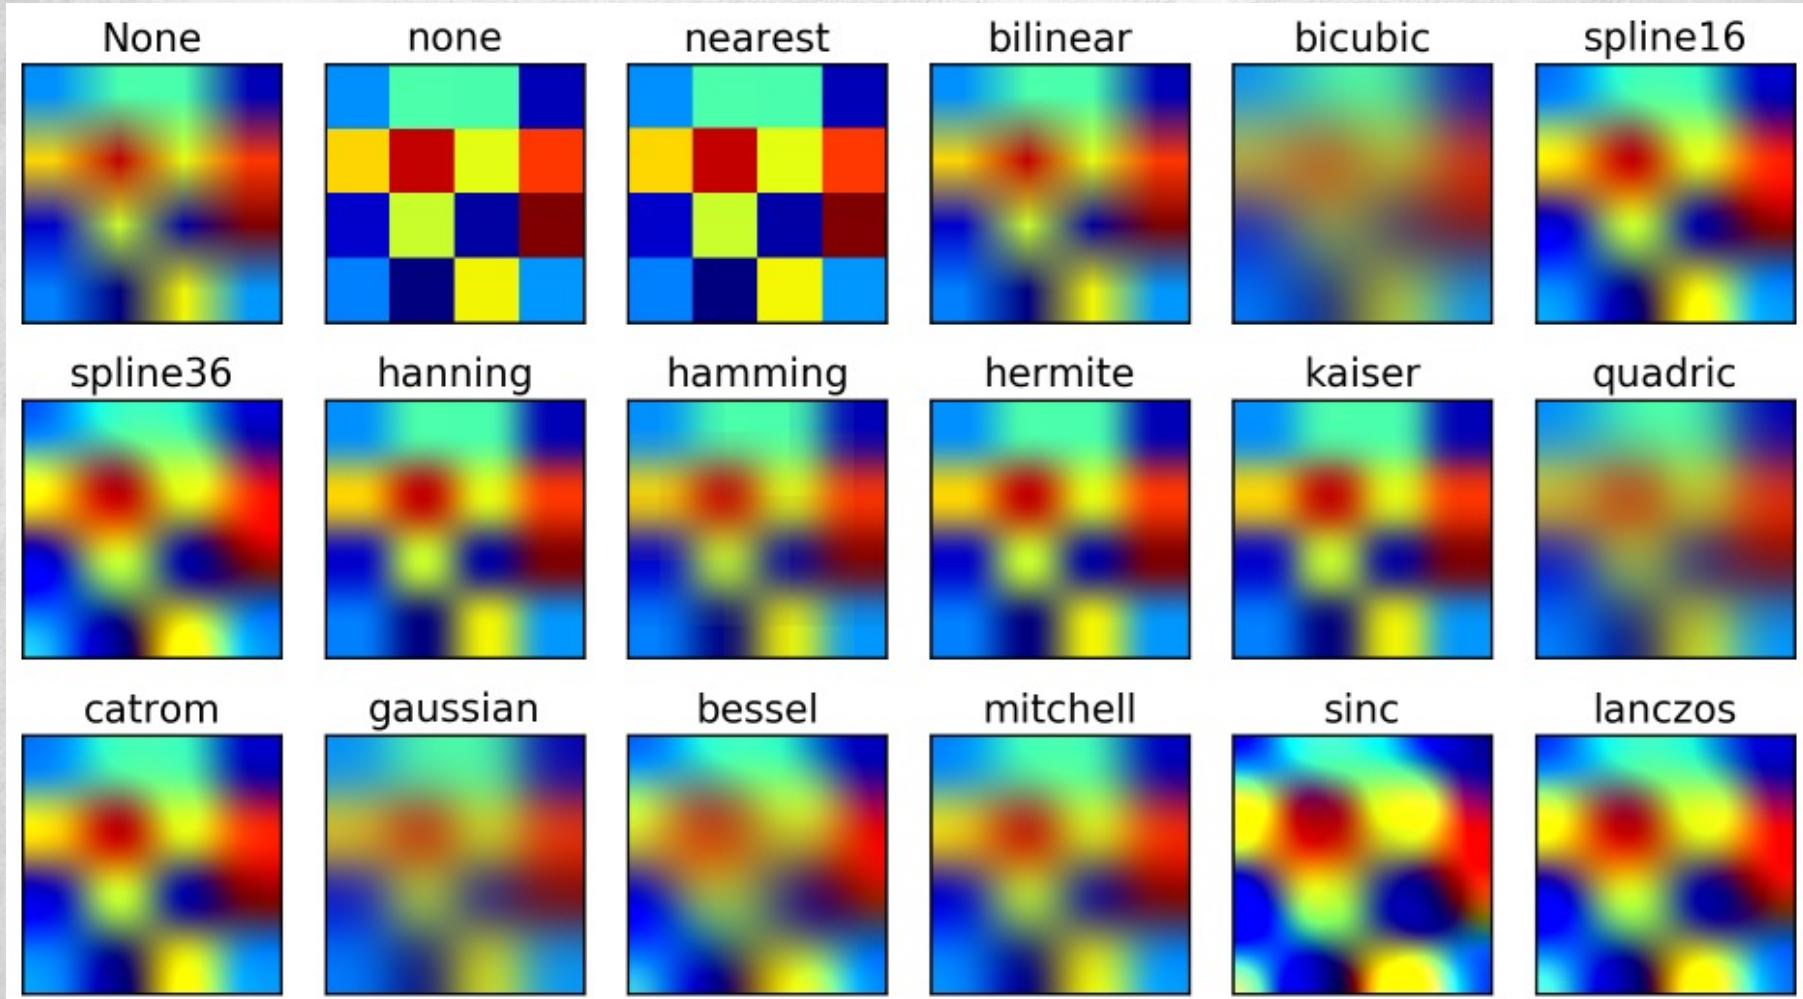


Scaled up  
to 300%

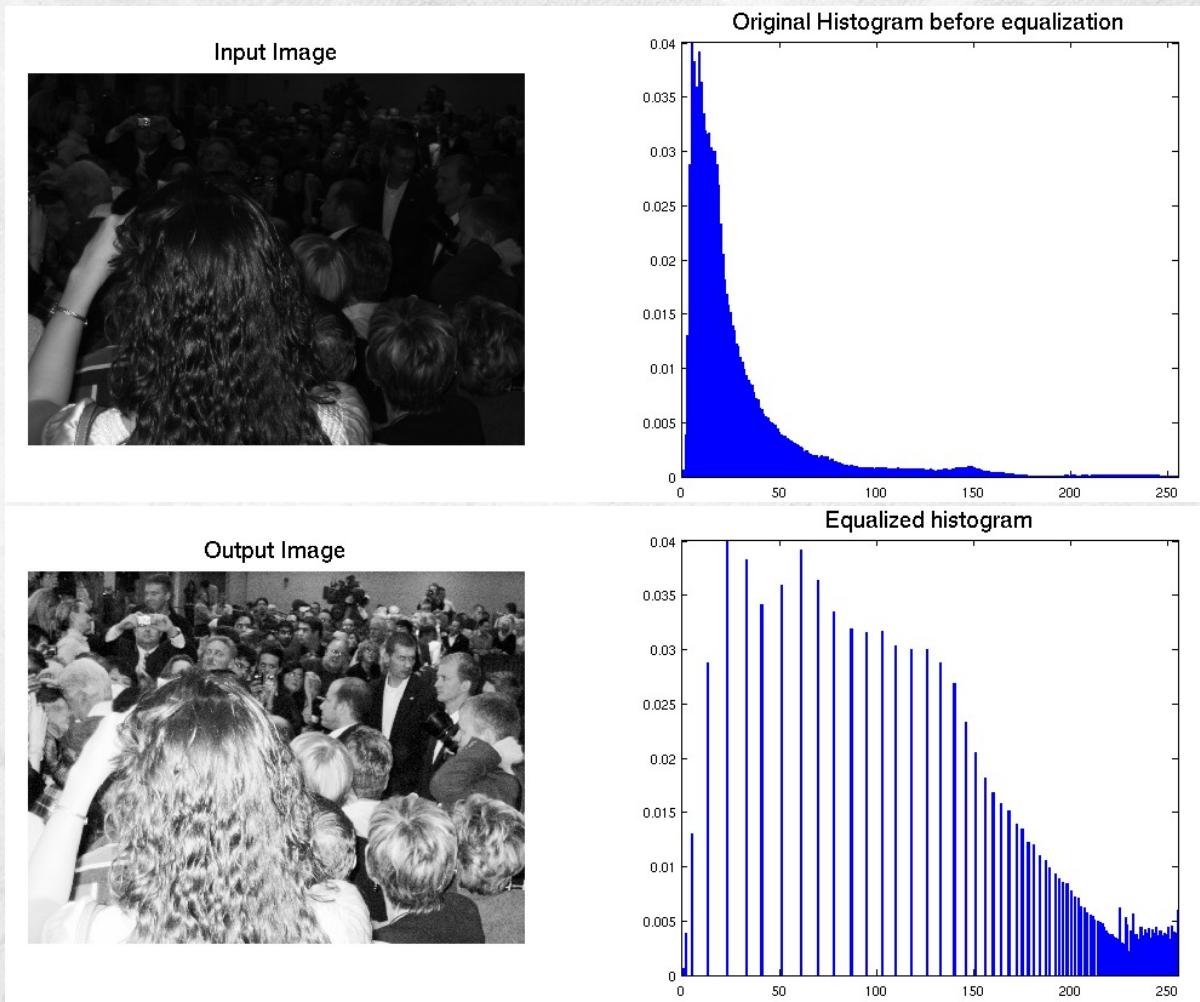


MISSISSIPPI STATE  
UNIVERSITY™

# Image Filtering

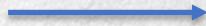


# Brightness Normalization



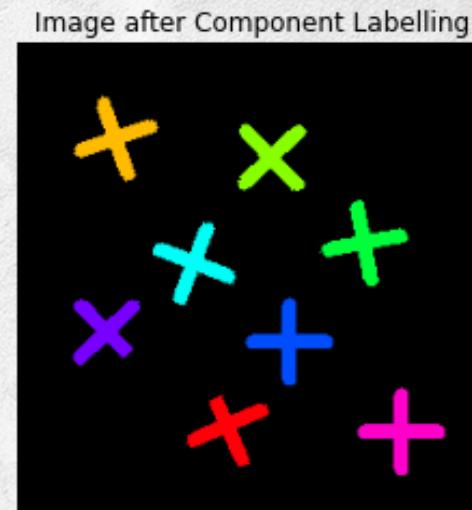
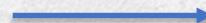
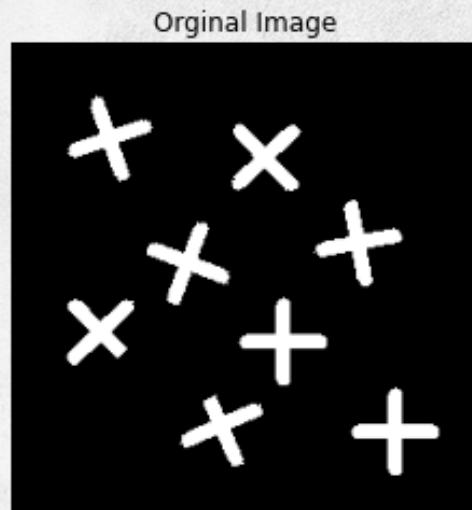
MISSISSIPPI STATE  
UNIVERSITY™

# Binary Thresholding

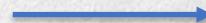


MISSISSIPPI STATE  
UNIVERSITY™

# Connected Components



1	1	0	0	1
0	1	0	1	0
0	1	0	0	0
0	0	1	1	1
0	0	0	0	1

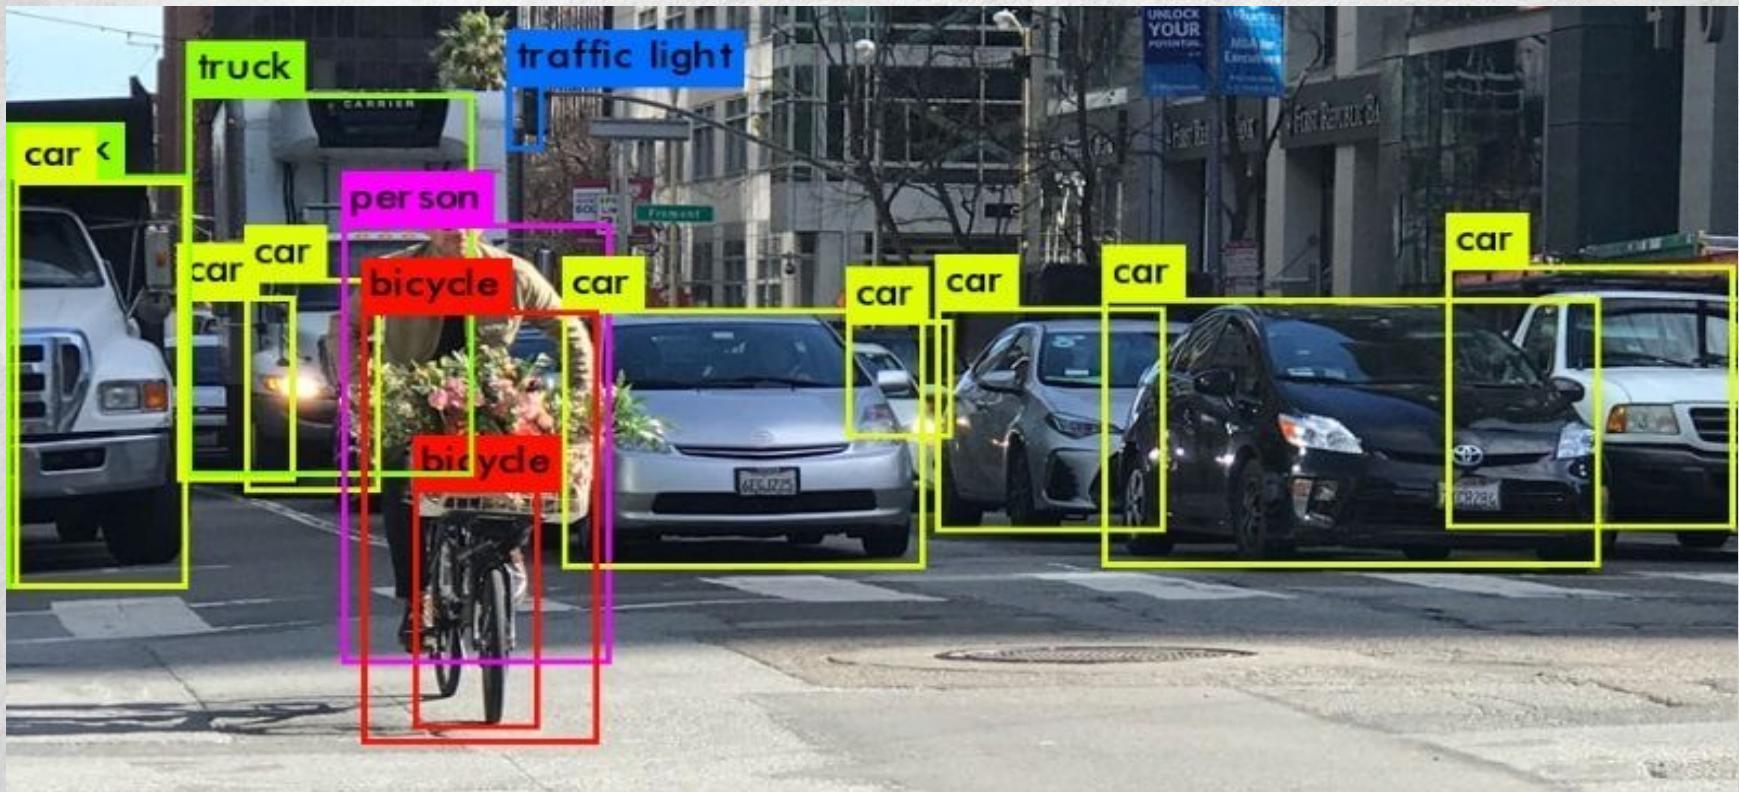


1	1	0	0	2
0	1	0	3	0
0	1	0	0	0
0	0	4	4	4
0	0	0	0	4

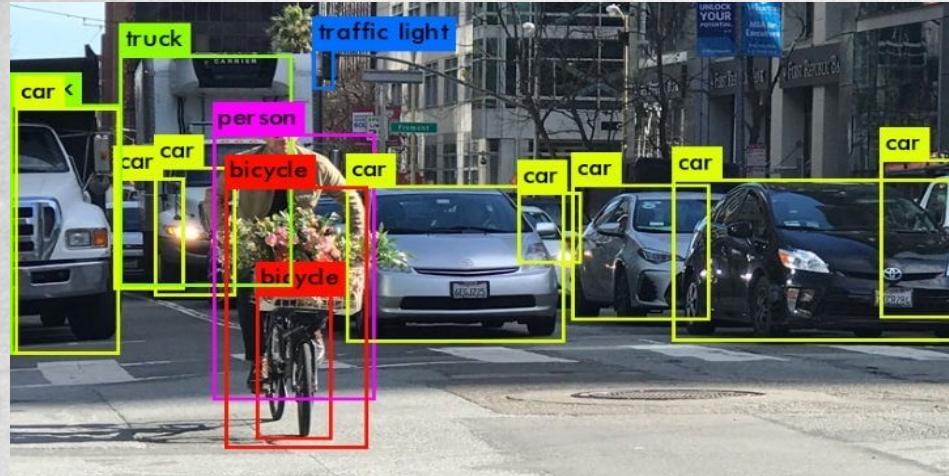


# Object Recognition

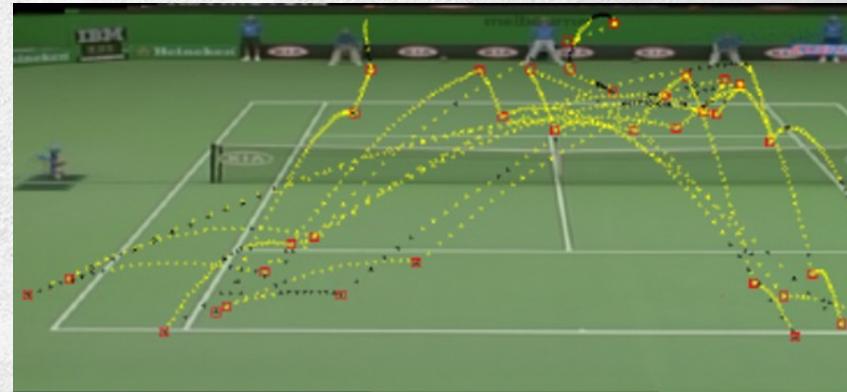
Computer vision technique that deals with detecting semantic objects of a certain class (e.g. humans, buildings, or cars) present in digital images and videos



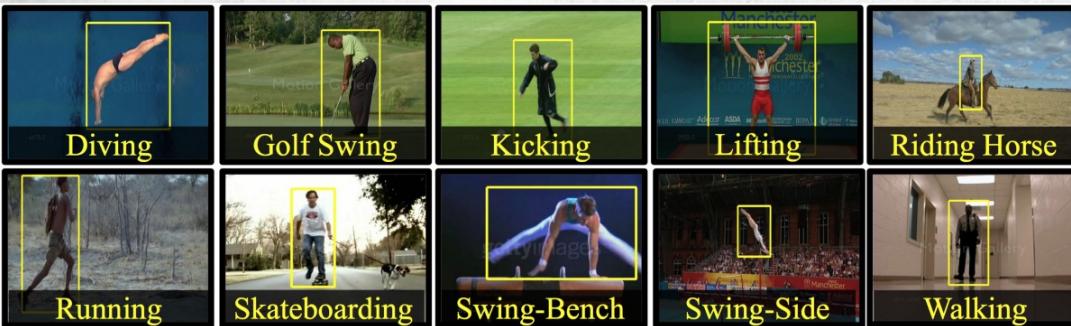
# Autonomous Driving



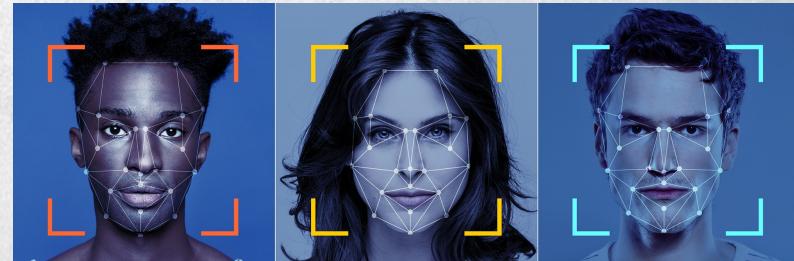
# Tracking



# Activity Recognition



# Face Recognition



MISSISSIPPI STATE  
UNIVERSITY™

# Different flavors of object recognition

Semantic Segmentation



No objects, just pixels

Classification + Localization



Single Object

Object Detection



Multiple Object

Instance Segmentation

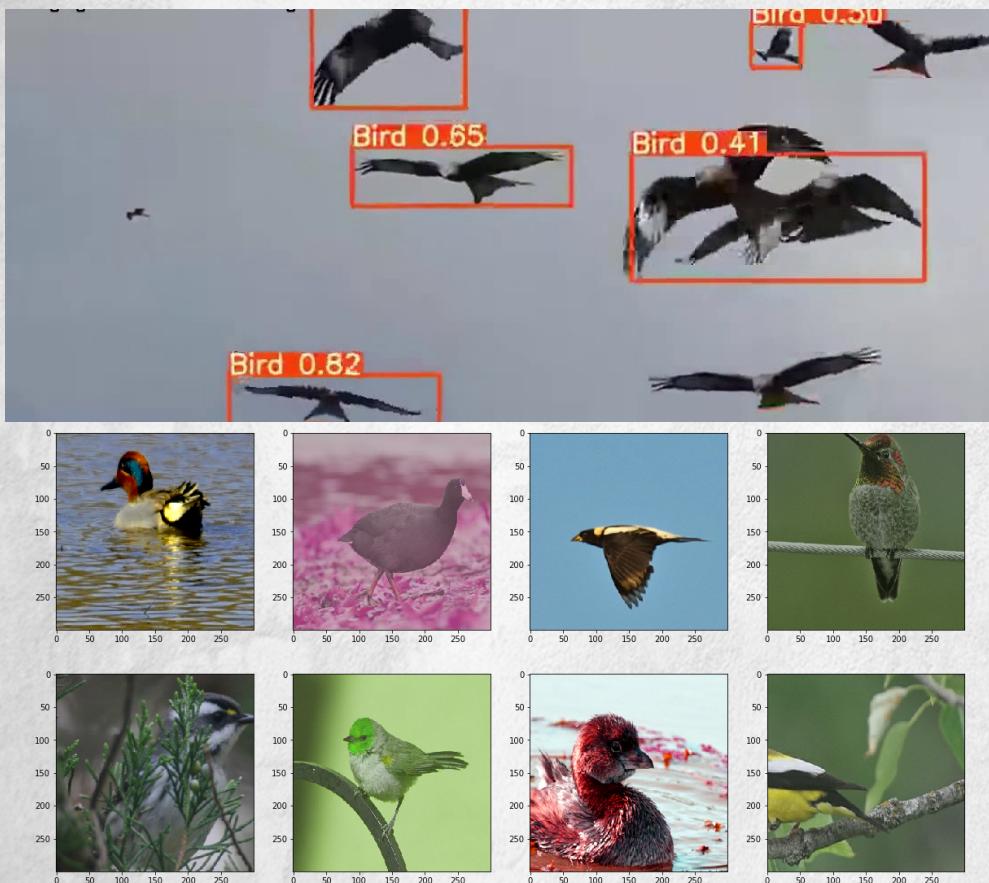


This image is CC0 public domain



MISSISSIPPI STATE  
UNIVERSITY™

# Real-world object recognition is difficult



WHEN A USER TAKES A PHOTO,  
THE APP SHOULD CHECK WHETHER  
THEY'RE IN A NATIONAL PARK...

SURE, EASY GIS LOOKUP.  
GIMME A FEW HOURS.

...AND CHECK WHETHER  
THE PHOTO IS OF A BIRD.

I'LL NEED A RESEARCH  
TEAM AND FIVE YEARS.



MISSISSIPPI STATE  
UNIVERSITY™

# Perception in Autonomous Driving

Computer Vision

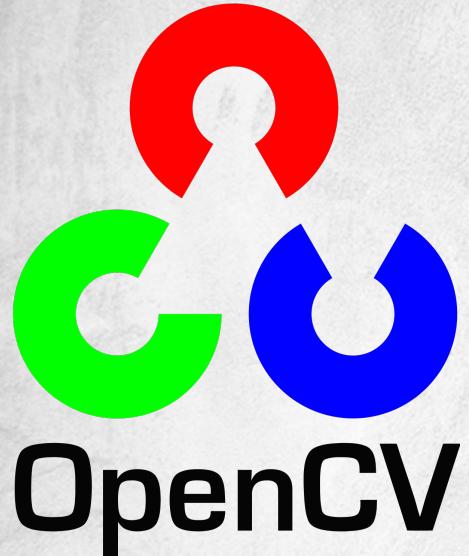


MISSISSIPPI STATE  
UNIVERSITY™

# Image Processing Software



MISSISSIPPI STATE  
UNIVERSITY™



## OpenCV

[https://docs.opencv.org/4.5.0/d6/d00/tutorial\\_py\\_root.html](https://docs.opencv.org/4.5.0/d6/d00/tutorial_py_root.html)



## *scikit-image*

[https://scikit-image.org/docs/dev/auto\\_examples/](https://scikit-image.org/docs/dev/auto_examples/)



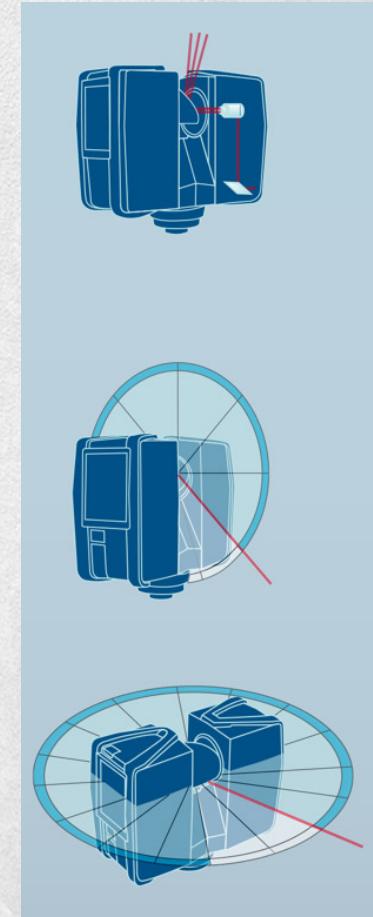
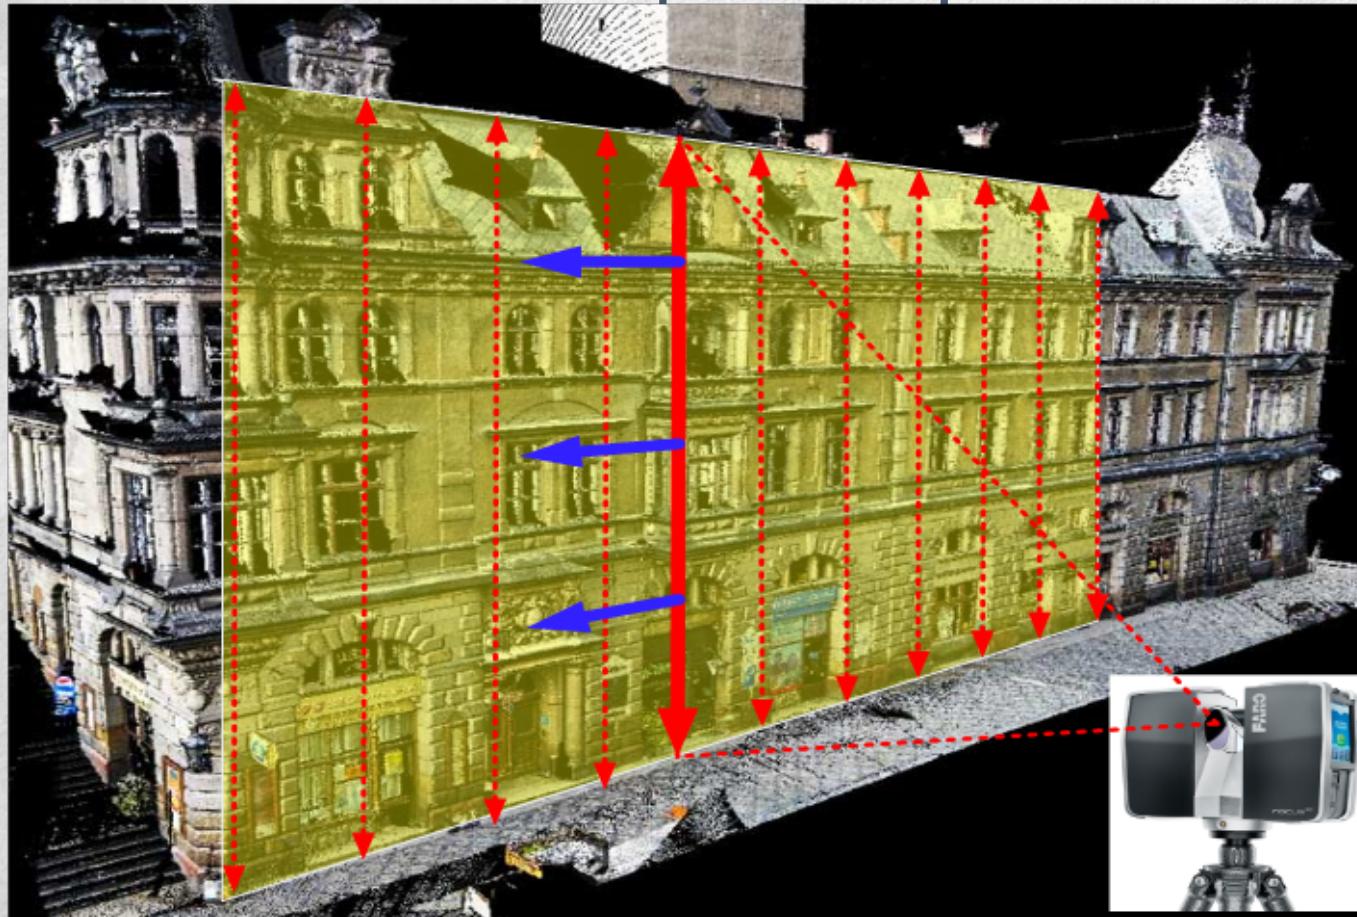
MISSISSIPPI STATE  
UNIVERSITY™

# 3D Perception



MISSISSIPPI STATE  
UNIVERSITY™

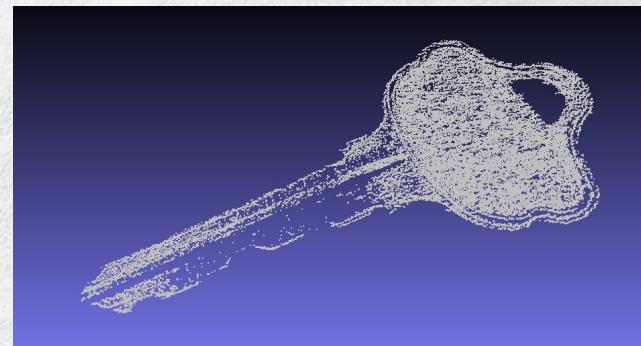
# LiDAR Scanning – working principles



MISSISSIPPI STATE  
UNIVERSITY™

# LiDAR scanning– point cloud

- Point cloud
  - product of a laser scanner
  - a set of points (millions)
  - point info: X, Y, Z, (I, R, G, B)
  - formats: fls, pts, xyz, las, txt...



MISSISSIPPI STATE  
UNIVERSITY™

# Point Cloud Data Representation

Open Ascii File

Filename: C:/Users/paolo/Desktop/NUVOLA.txt

Here are the first lines of this file. Choose an attribute for each column (one cloud at a time):

1	2	3
<input type="button" value="X coord. X"/>	<input type="button" value="Y coord. Y"/>	<input type="button" value="Z coord. Z"/>
589817.05	4884381.55	614.68
589816.37	4884382.87	613.98
589817.04	4884381.58	614.97
589816.64	4884382.56	614.03
589817.03	4884381.92	614.35
589816.70	4884382.19	614.33
589816.69	4884382.24	613.74
589816.70	4884382.20	613.99
589817.05	4884381.87	614.59
589816.37	4884382.50	614.31
589817.04	4884382.22	613.71
589817.07	4884381.88	614.94
589816.65	4884382.82	613.40
589816.97	4884382.86	613.70
589817.09	4884380.94	614.96

Separator  (ASCII code: 32) whitespace , ;

Skip lines 0  extract scalar field names from first line

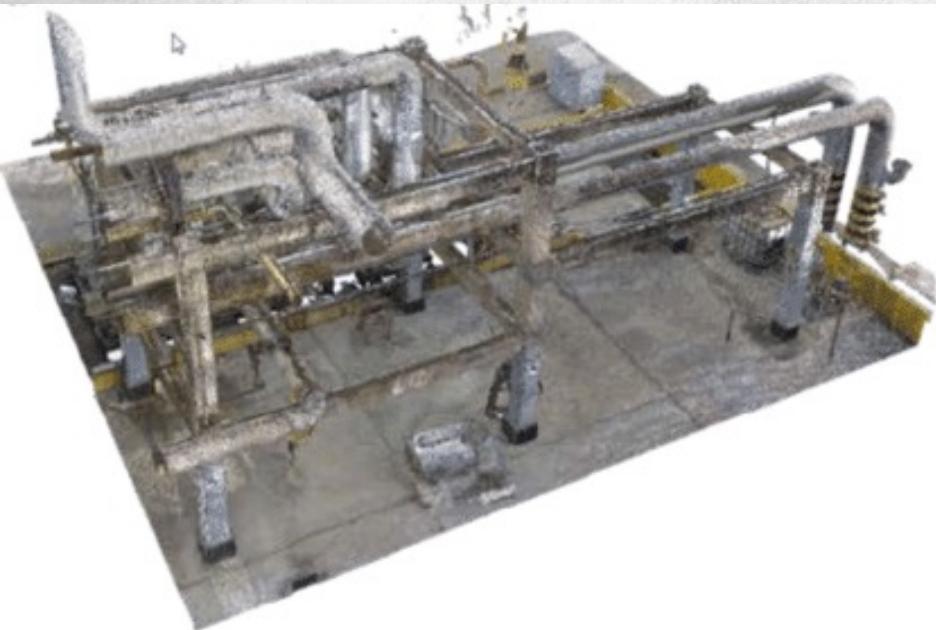
Max number of points per cloud 2000.00 Million

use comma as decimal character  Show labels in 2D

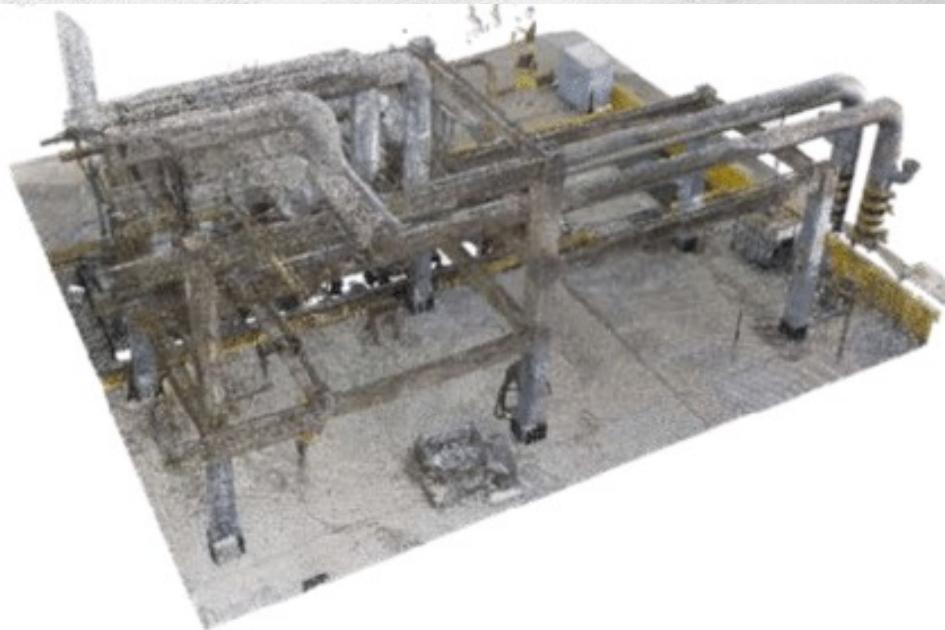


MISSISSIPPI STATE  
UNIVERSITY™

# Point Cloud Downsampling



(a) Original [3,467,072 points]

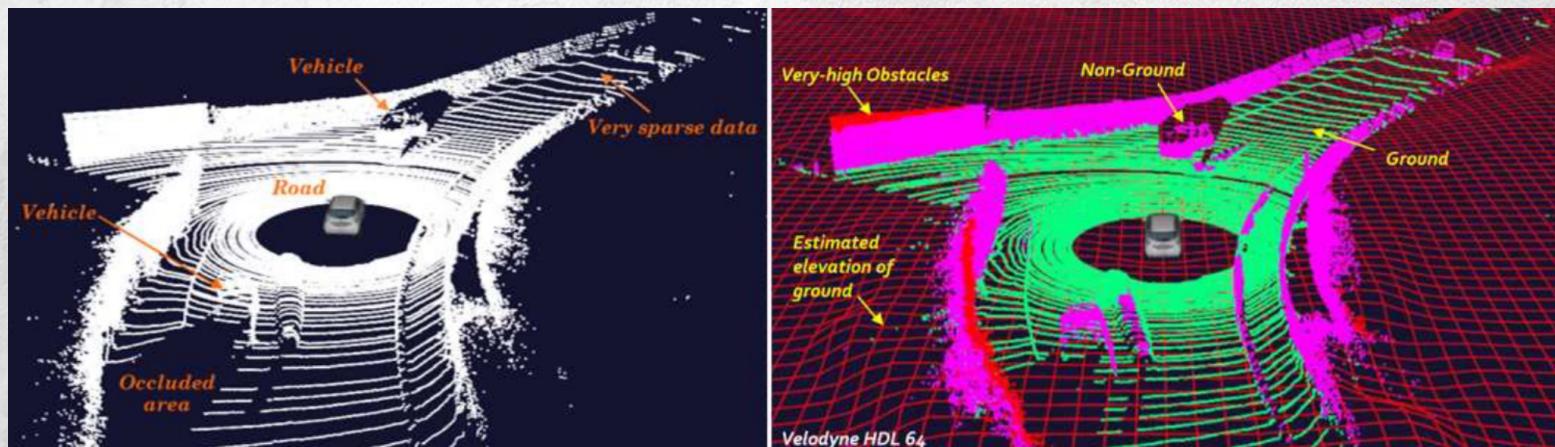
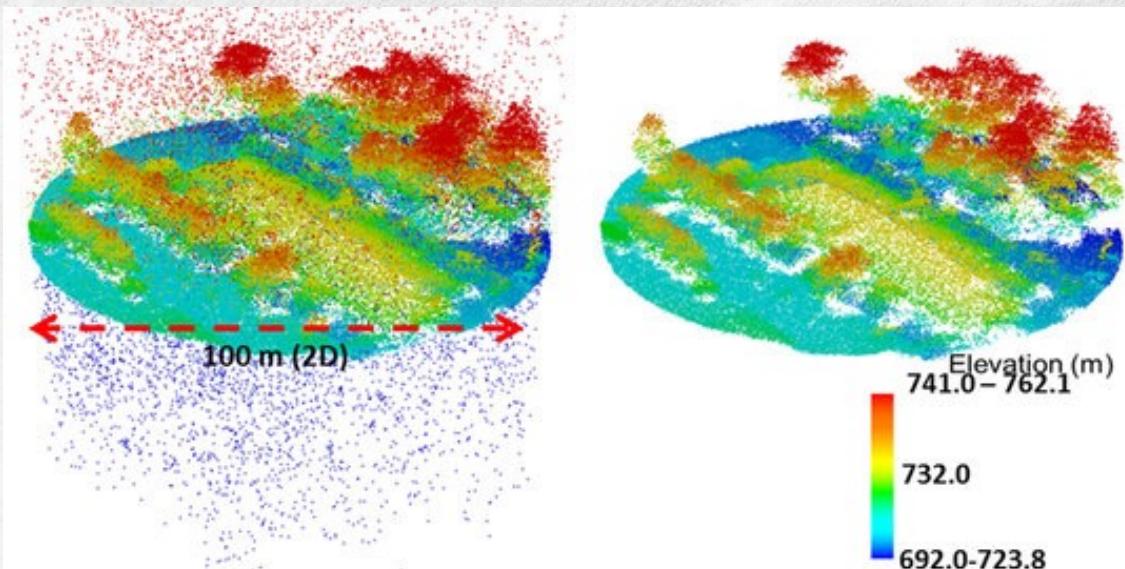


(b) Downsampled [322,797 points]



MISSISSIPPI STATE  
UNIVERSITY™

# Spatial Filtering

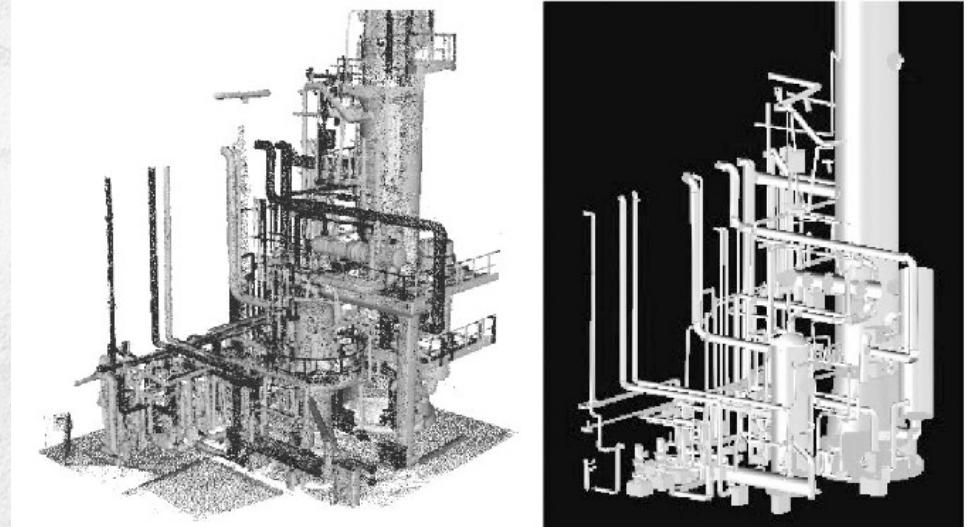


# Model-based Segmentation

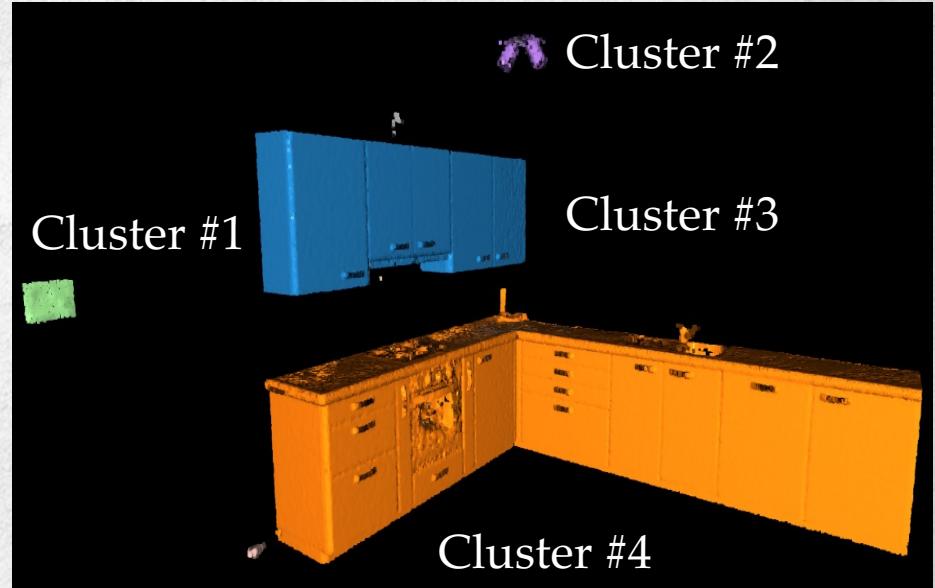
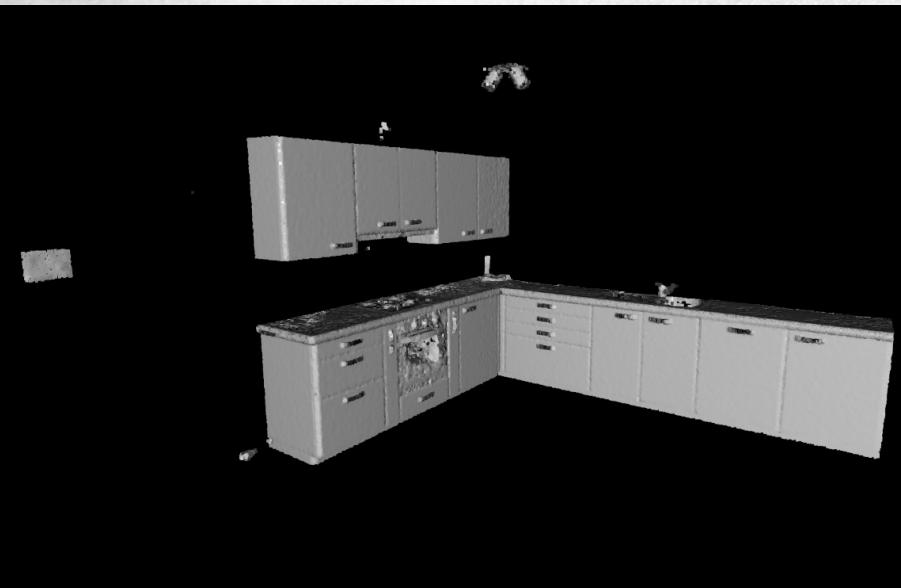
Plane segmentation



Cylinder segmentation



# Euclidean Clustering / Connected Components

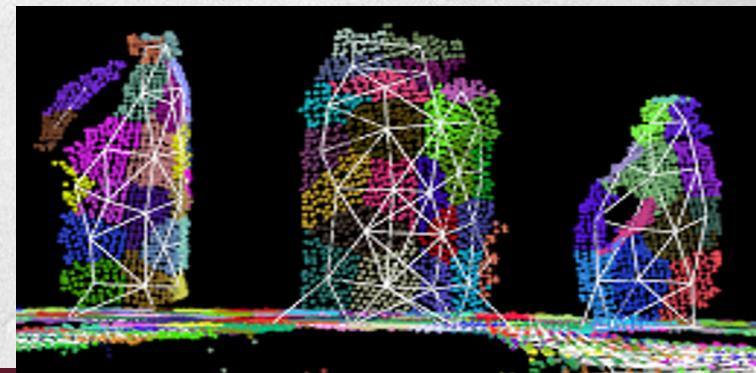


MISSISSIPPI STATE  
UNIVERSITY™

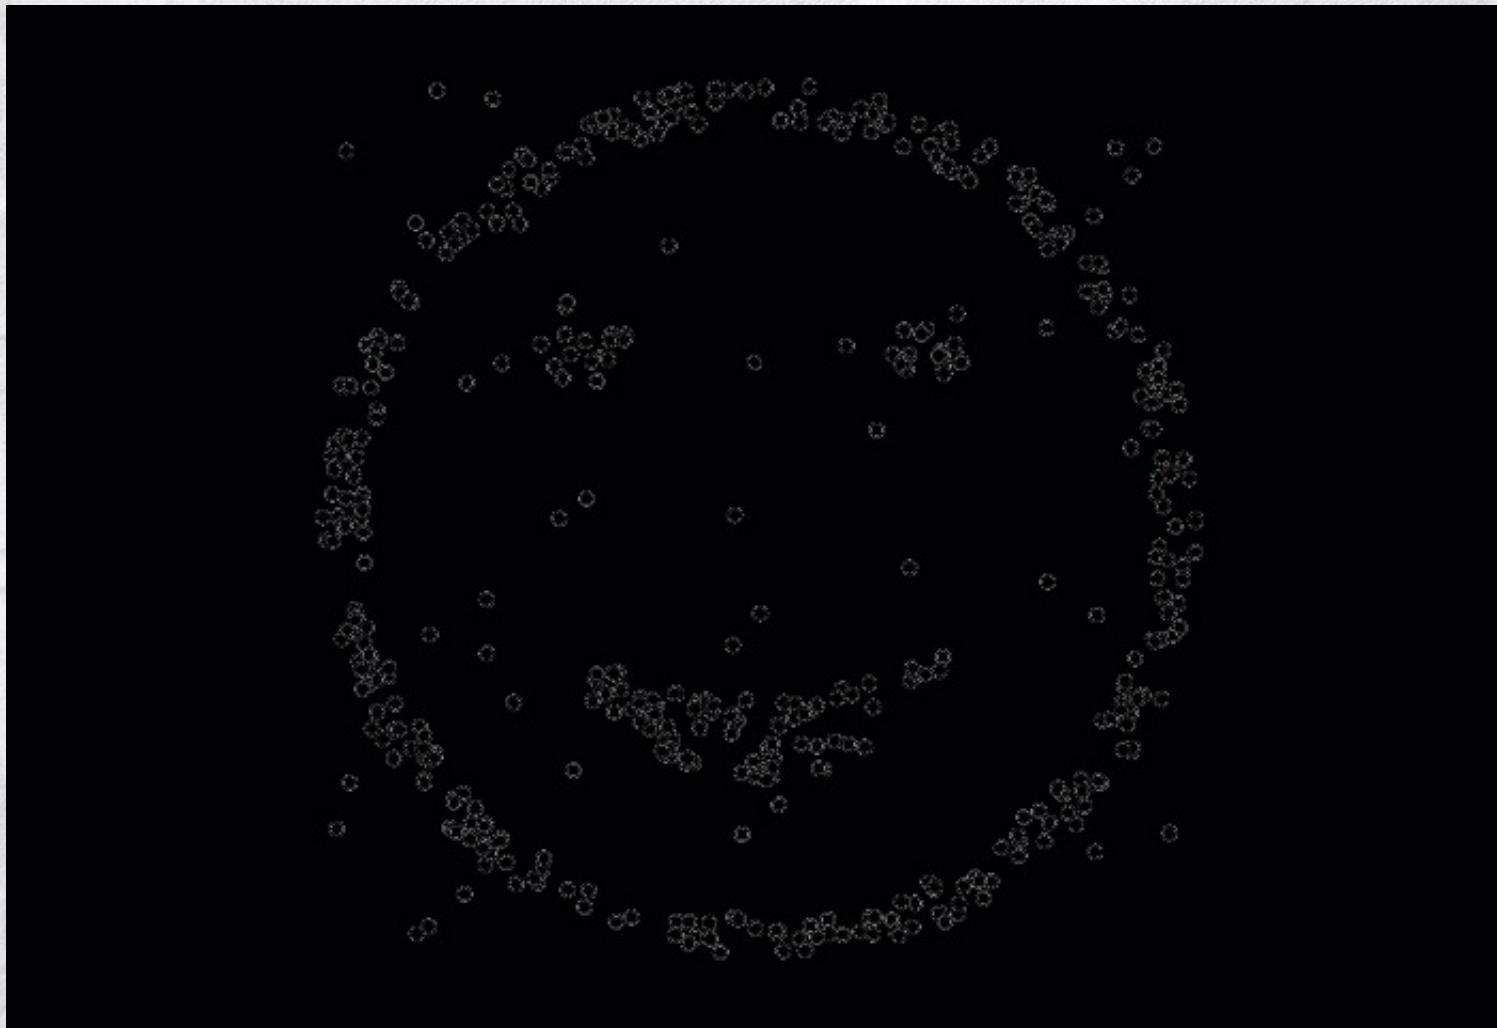
# Euclidean Clustering

## Point cloud clustering (Euclidean)

- Organize the point cloud into separate clusters corresponding to different objects
- Iterative process with the following steps:
  - Pick a seed point and initialize a cluster
  - Add all the neighboring points within a fixed distance threshold to the same cluster
  - Repeat for each adjacent point until no more neighboring points are found

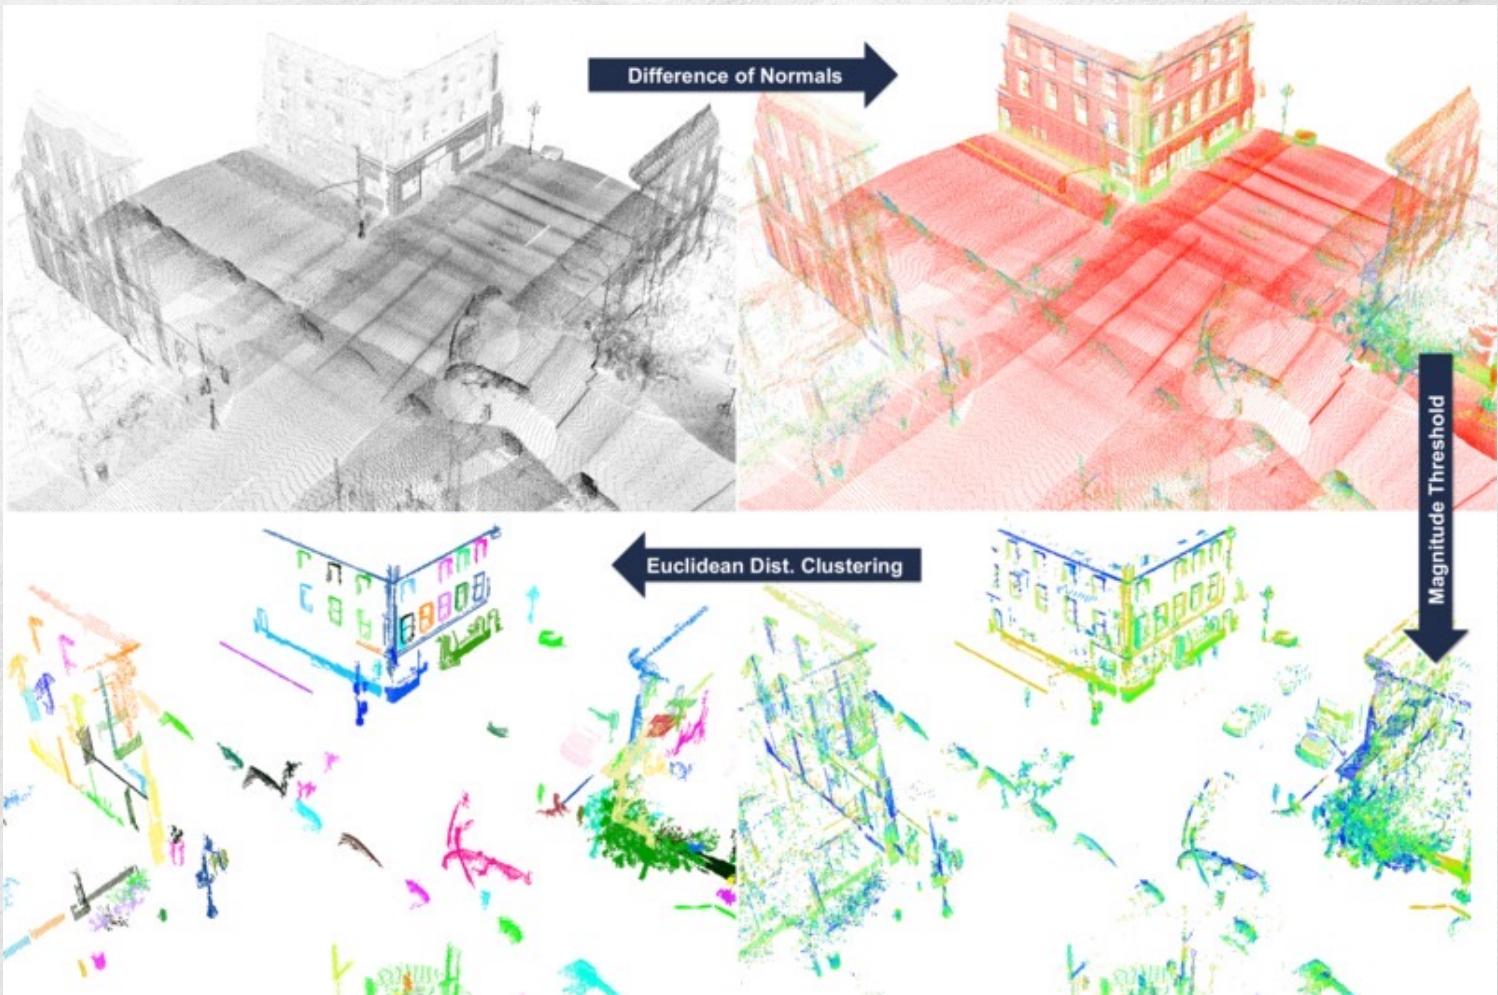


# Euclidean Clustering

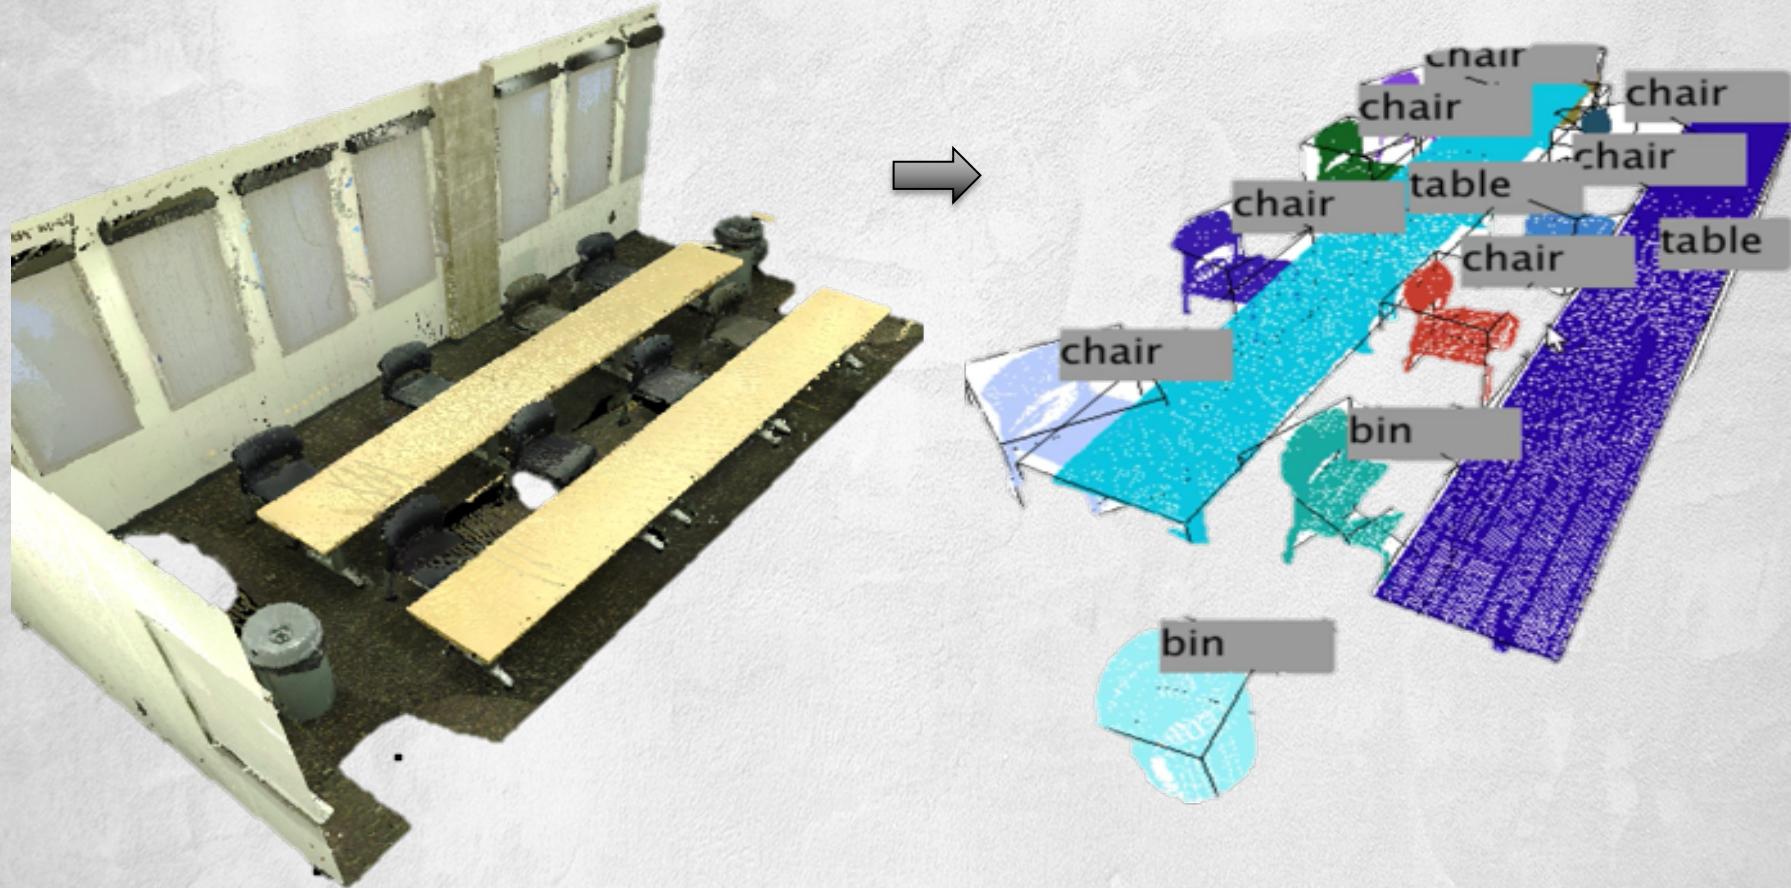


MISSISSIPPI STATE  
UNIVERSITY™

# Normal-based Segmentation



# Object Recognition



MISSISSIPPI STATE  
UNIVERSITY™

# Different flavors of object recognition

Classification



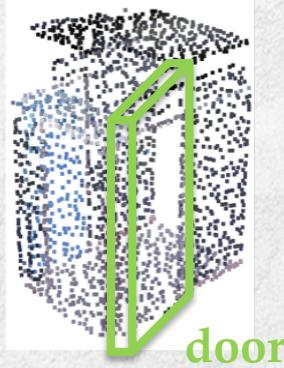
Input: Object  
Output: Object  
Label

Part Segmentation



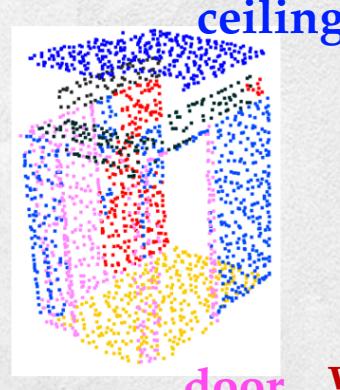
Input: Object  
Output: Point  
Labels

Object Detection



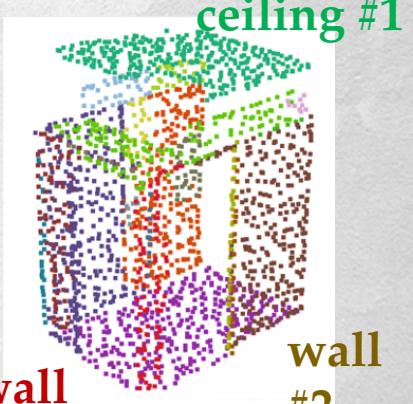
Input: Scene  
Output:  
Bounding Boxes

Semantic Segmentation



Input: Scene  
Output: Point  
Labels

Instance Segmentation



Input: Scene  
Output: Instance  
ID + Object Labels



# Point Cloud Processing Software

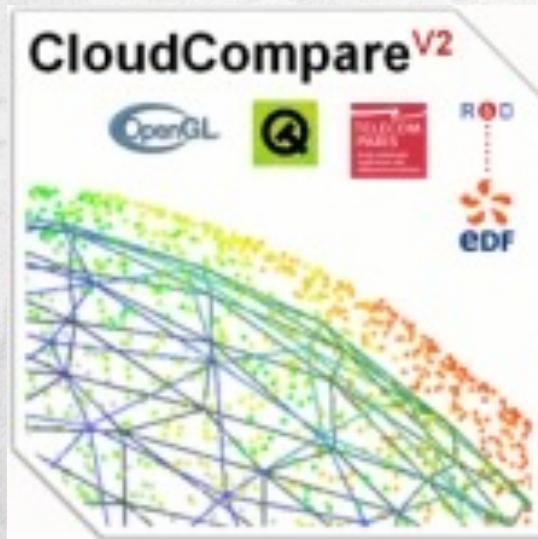


MISSISSIPPI STATE  
UNIVERSITY™



# Point Cloud Library

<https://pointclouds.org/>



# CloudCompare

<https://www.danielgm.net/cc/>



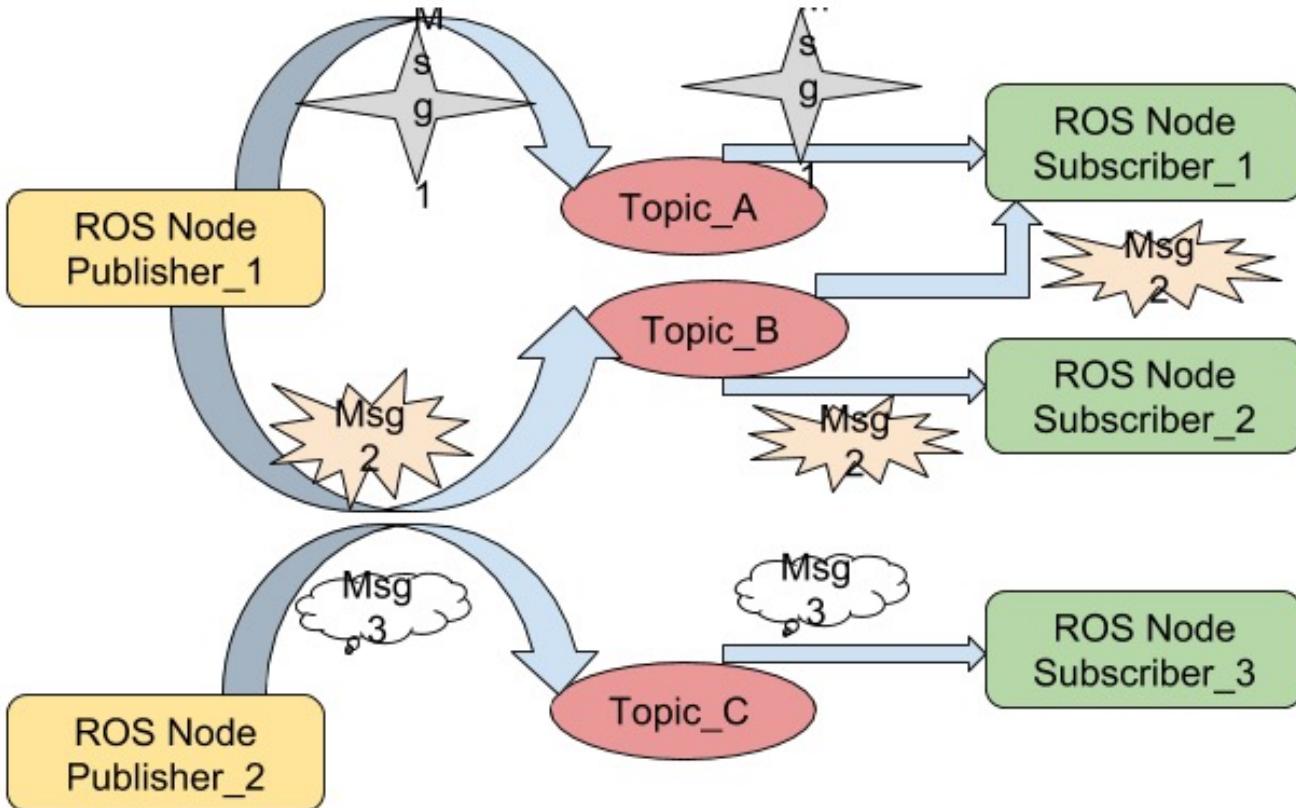
MISSISSIPPI STATE  
UNIVERSITY™

# Sensor Messages in ROS

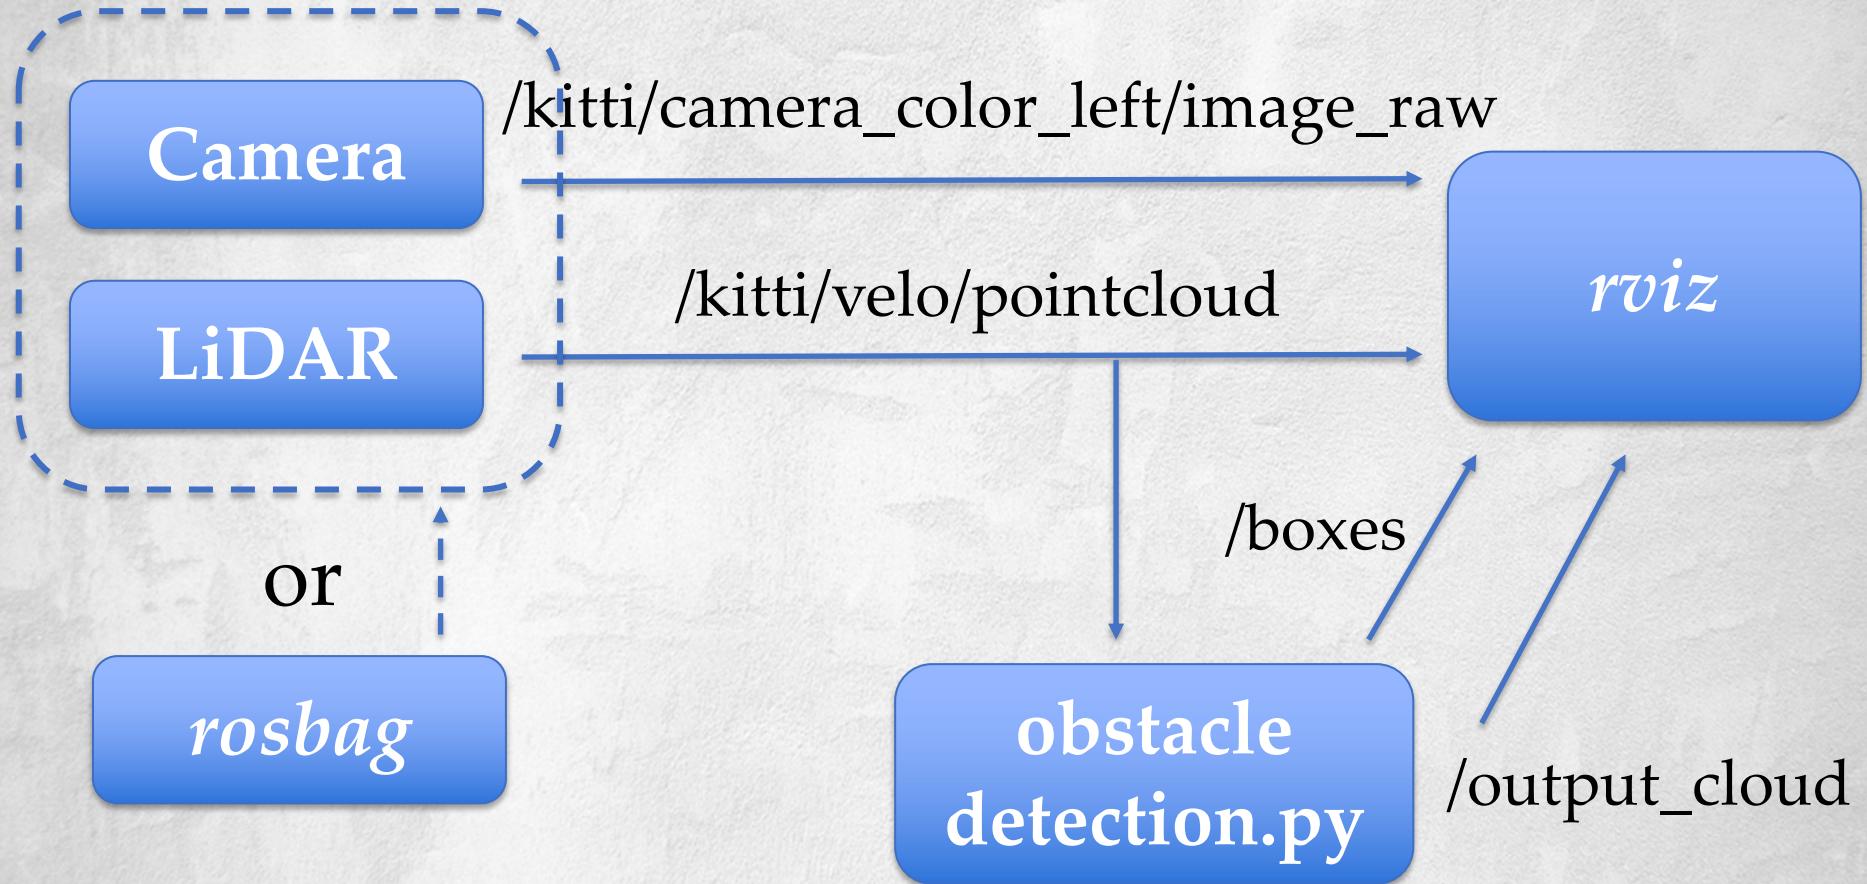


MISSISSIPPI STATE  
UNIVERSITY™

# *sensor\_msgs* API



# Example



# sensor\_msgs/Image Message

File: **sensor\_msgs/Image.msg**

## Raw Message Definition

```
# This message contains an uncompressed image
# (0, 0) is at top-left corner of image
#
Header header          # Header timestamp should be acquisition time of image
                        # Header frame_id should be optical frame of camera
                        # origin of frame should be optical center of camera
                        # +x should point to the right in the image
                        # +y should point down in the image
                        # +z should point into to plane of the image
                        # If the frame_id here and the frame_id of the CameraInfo
                        # message associated with the image conflict
                        # the behavior is undefined
uint32 height           # image height, that is, number of rows
uint32 width             # image width, that is, number of columns
#
# The legal values for encoding are in file src/image_encodings.cpp
# If you want to standardize a new string format, join
# ros-users@lists.sourceforge.net and send an email proposing a new encoding.
string encoding          # Encoding of pixels -- channel meaning, ordering, size
                        # taken from the list of strings in include/sensor_msgs/image_encodings.h
uint8 is_bigendian        # is this data big endian?
uint32 step               # Full row length in bytes
uint8[] data              # actual matrix data, size is (step * rows)
```



MISSISSIPPI STATE  
UNIVERSITY™

# sensor\_msgs/PointCloud2 Message

File: **sensor\_msgs/PointCloud2.msg**

## Raw Message Definition

```
# This message holds a collection of N-dimensional points, which may
# contain additional information such as normals, intensity, etc. The
# point data is stored as a binary blob, its layout described by the
# contents of the "fields" array.

# The point cloud data may be organized 2d (image-like) or 1d
# (unordered). Point clouds organized as 2d images may be produced by
# camera depth sensors such as stereo or time-of-flight.

# Time of sensor data acquisition, and the coordinate frame ID (for 3d
# points).
Header header

# 2D structure of the point cloud. If the cloud is unordered, height is
# 1 and width is the length of the point cloud.
uint32 height
uint32 width

# Describes the channels and their layout in the binary data blob.
PointField[] fields

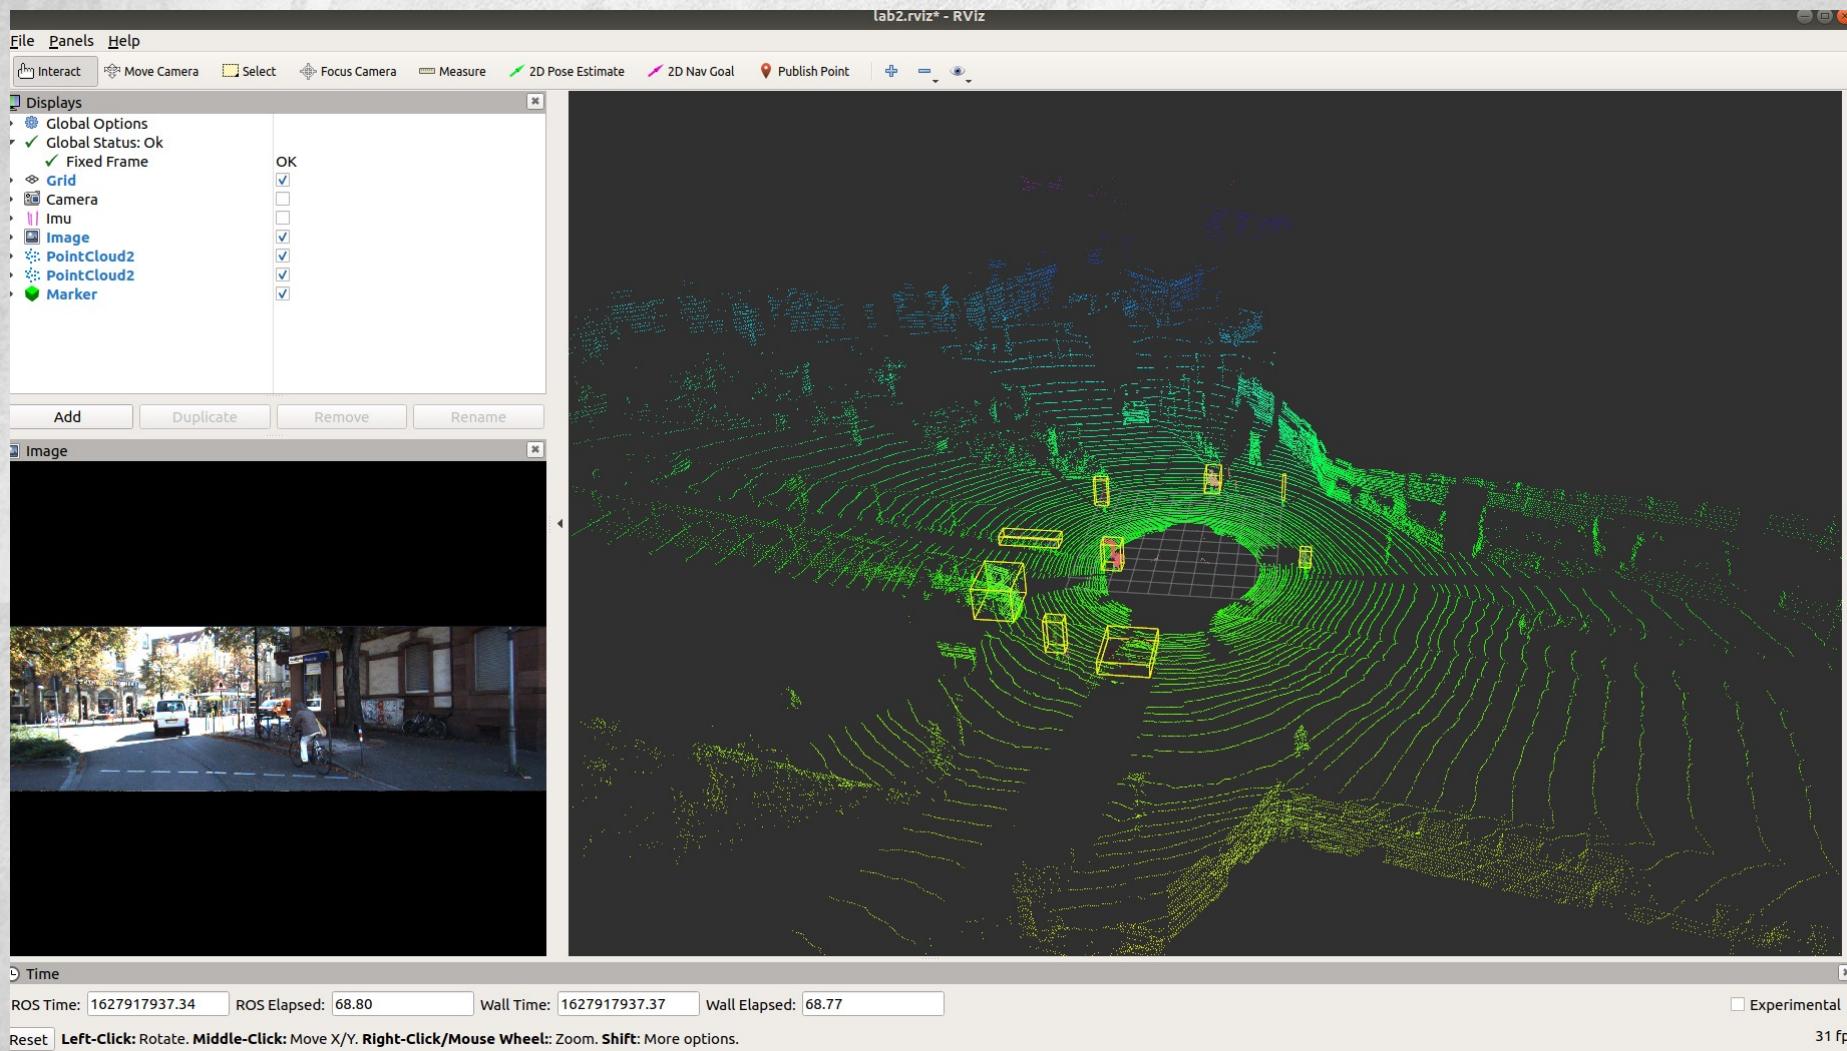
bool    is_bigendian # Is this data big endian?
uint32  point_step   # Length of a point in bytes
uint32  row_step     # Length of a row in bytes
uint8[] data         # Actual point data, size is (row_step*height)

bool is_dense        # True if there are no invalid points
```



MISSISSIPPI STATE  
UNIVERSITY™

# RVIZ



MISSISSIPPI STATE  
UNIVERSITY™

# Challenges in Perception



MISSISSIPPI STATE  
UNIVERSITY™



MISSISSIPPI STATE  
UNIVERSITY™

# References

1. [http://wiki.ros.org/sensor\\_msgs](http://wiki.ros.org/sensor_msgs)
2. [https://docs.opencv.org/4.5.0/d6/d00/tutorial\\_py\\_root.html](https://docs.opencv.org/4.5.0/d6/d00/tutorial_py_root.html)
3. [https://scikit-image.org/docs/dev/auto\\_examples/](https://scikit-image.org/docs/dev/auto_examples/)
4. <https://pointclouds.org/>
5. <https://www.danielgm.net/cc/>
6. <https://towardsdatascience.com/how-to-automate-3d-point-cloud-segmentation-and-clustering-with-python-343c9039e4f5>

