

AI Robotics

Particle Filter



MISSISSIPPI STATE
UNIVERSITY™

Overview



MISSISSIPPI STATE
UNIVERSITY™

Localization

Sensor Measurements:

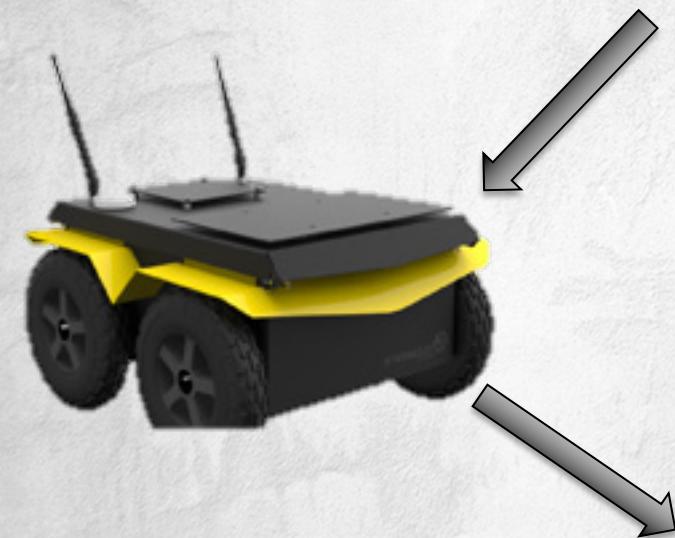
Robot State (or pose):
 $X_{0:t} = \{X_0, \dots, X_t\}$

$$Z_{0:t} = \{z_0, \dots, z_t\}$$

sense

Known

Unknown



act

Robot Controls:

$$u_{0:t} = \{u_0, u_1, \dots, u_t\}$$

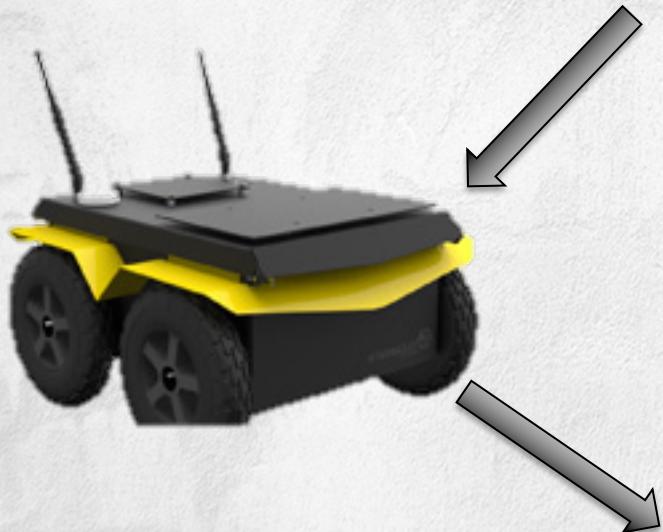


MISSISSIPPI STATE
UNIVERSITY™

Probabilistic State Estimation

$$P(x_t|u_1, z_1, u_2, z_2, \dots, u_t, z_t) = P(x_t|z_{1:t}, u_{1:t})$$

Robot State



sense

Sensor Model

$$P(z_t|x_t)$$

Motion Model

$$P(x_t|x_{t-1}, u_t)$$



MISSISSIPPI STATE
UNIVERSITY™

Types of Localization Algorithms

Localization

Non-probabilistic

Direct Sensing

Probabilistic

Bayesian Filtering

Dead Reckoning

GPS

Kalman Filter

Particle Filter



MISSISSIPPI STATE
UNIVERSITY™

Particle Filter



MISSISSIPPI STATE
UNIVERSITY™

Bayes Filter Algorithm

$$Bel(x_t) = \eta P(z_t | x_t) \int P(x_t | u_t, x_{t-1}) Bel(x_{t-1}) dx_{t-1}$$

Algorithm **Bayes_filter**($Bel(x_{t-1})$, u_t , z_t):

Previous State Current control
 input Current sensor
 measurement

$$Bel'(x_t) = \int P(x_t | u_t, x_{t-1}) Bel(x_{t-1}) dx$$

Prediction Step: Motion Update

$$Bel(x_t) = \eta P(z_t | x_t) Bel'(x_t)$$

Correction Step: Sensor Update

return $Bel(x_t)$

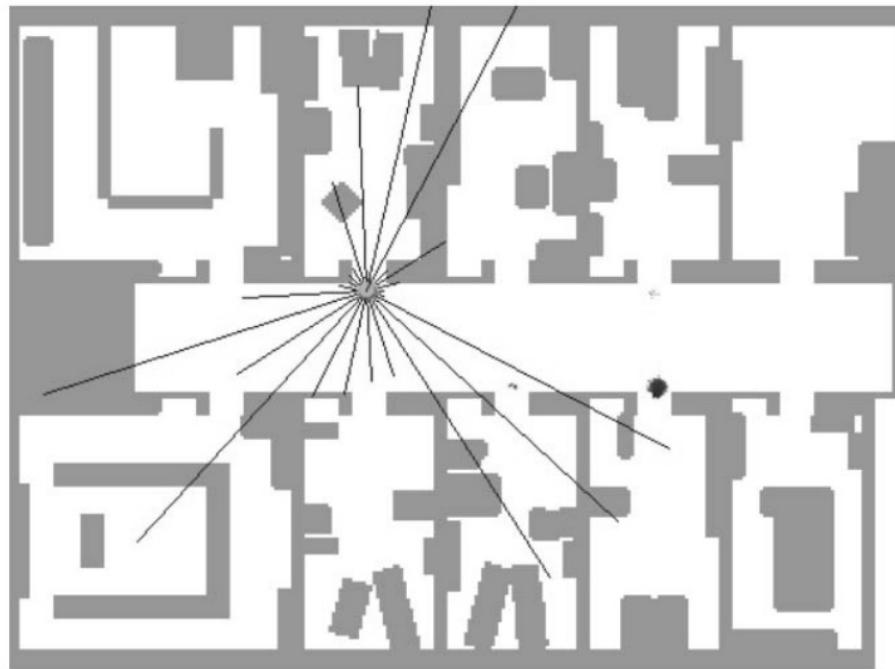
Current State



MISSISSIPPI STATE
UNIVERSITY™

The Problem with Kalman Filters in Robot Localization

- Kalman Filters only represent state variables as single Gaussians
- What if robot could be in one of two places?



Particle Filters (aka sequential Monte Carlo)



- Represents pdf as a set of samples (particles)
- Each particle contains one set of values for the state variables
- Good for non-Gaussian, multi-modal pdfs
- Find an approximate solution using a complex model (arbitrary pdf) rather than an exact solution using a simplified model (Gaussians)



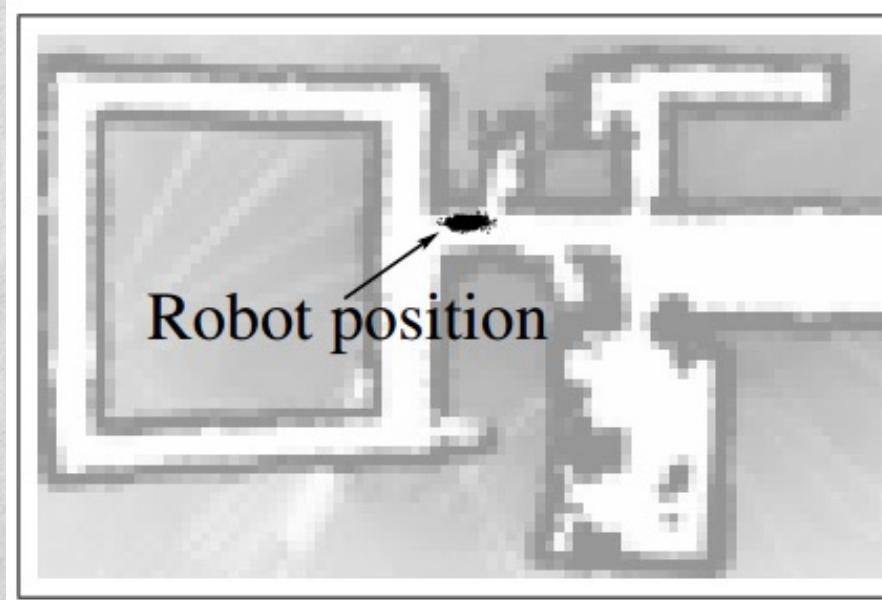
Particle Filters

- Main idea: Use samples to represent belief distribution
- Each sample will be called a particle
 - Each particle, or sample, corresponds to a potential pose of the robot
 - Intuitively, each particle is a “guess” at the actual pose
- Highly certain beliefs?
 - All particles very close together, or even the same
- Uncertain beliefs?
 - Particles distributed evenly around the space
- Set of particles approximates the distribution over poses
 - More particles = better approximation
 - Fewer = less precise approximation

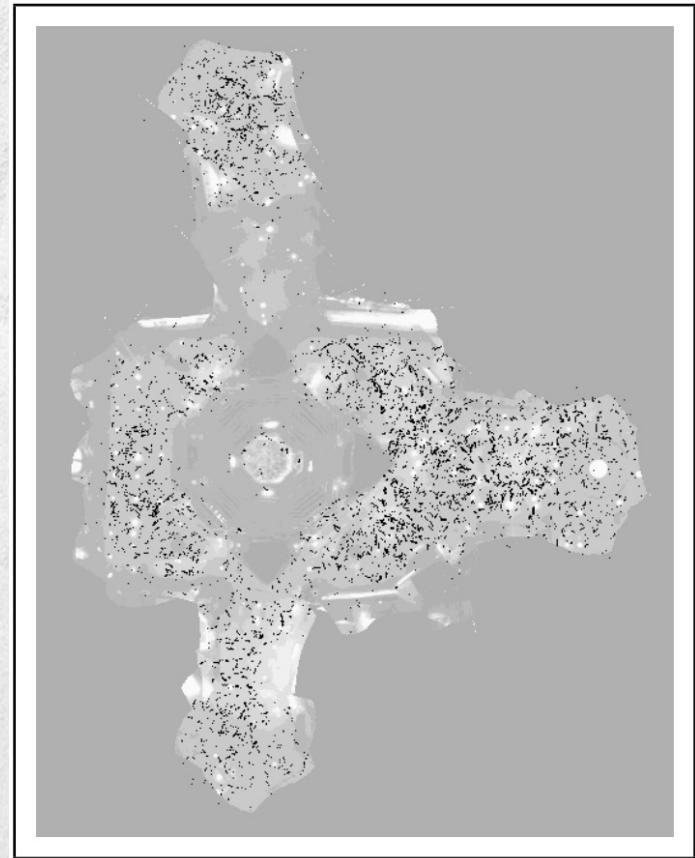


Example of Particle Distribution

Low Uncertainty



High Uncertainty

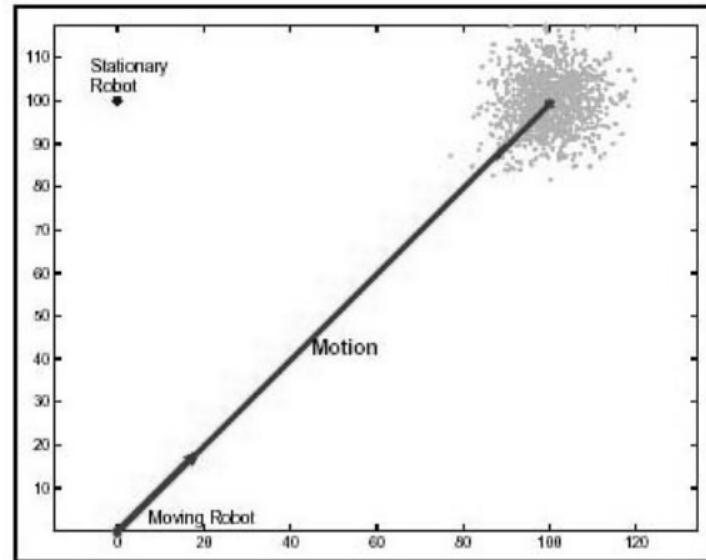


MISSISSIPPI STATE
UNIVERSITY™

1) Prediction: for each particle, sample and add random, noisy values from action model

Resulting proposal distribution ($q(x)$) approximates

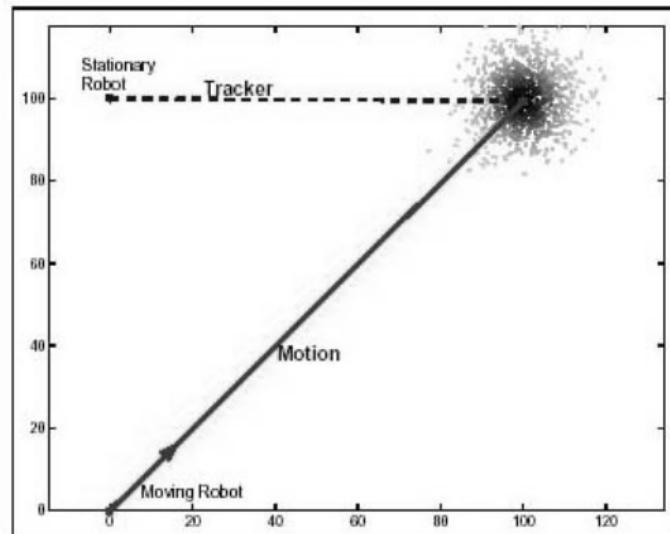
$$\int p(x_t | x_{t-1}, u_{t-1}) p(x_{t-1} | d_{0 \dots t-1}) dx_{t-1}$$



2) Update: each particle's weight is the likelihood of getting the sensor readings from that particle's hypothesis

The weight associated with each particle is

$w(x) = p(x)/q(x) = p(z_t | x_t)$, normalized so that all the weights sum to 1



3) Resample: new set of particles are chosen such that each particle survives in proportion to its weight



Resulting distribution is $p(x)$:

$$p(x_t | d_{o \dots t}) = \eta \underbrace{p(z_t | x_t)}_{p(x) - \text{the posterior probability}} \underbrace{\int p(x_t | u_{t-1}, x_{t-1}) p(x_{t-1} | d_{o \dots t-1}) dx_{t-1}}_{w(x) - \text{the importance weights}} \underbrace{p(x_{t-1} | d_{o \dots t-1})}_{q(x) - \text{the prior probability}}$$



Particle Filter: Steps (Simplified)

1. Initialize particles
2. At each time step
 1. Use motion model to propagate each particle forward (randomly)
 - (This will spread out particles that were the same)
 2. Use measurement model to assign each particle a weight
 - (Probability of sensor reading given particle's pose)
 3. Resample particles with replacement, with probability proportional to weight

At any point in time, our set of particles represents our current (approximate) beliefs

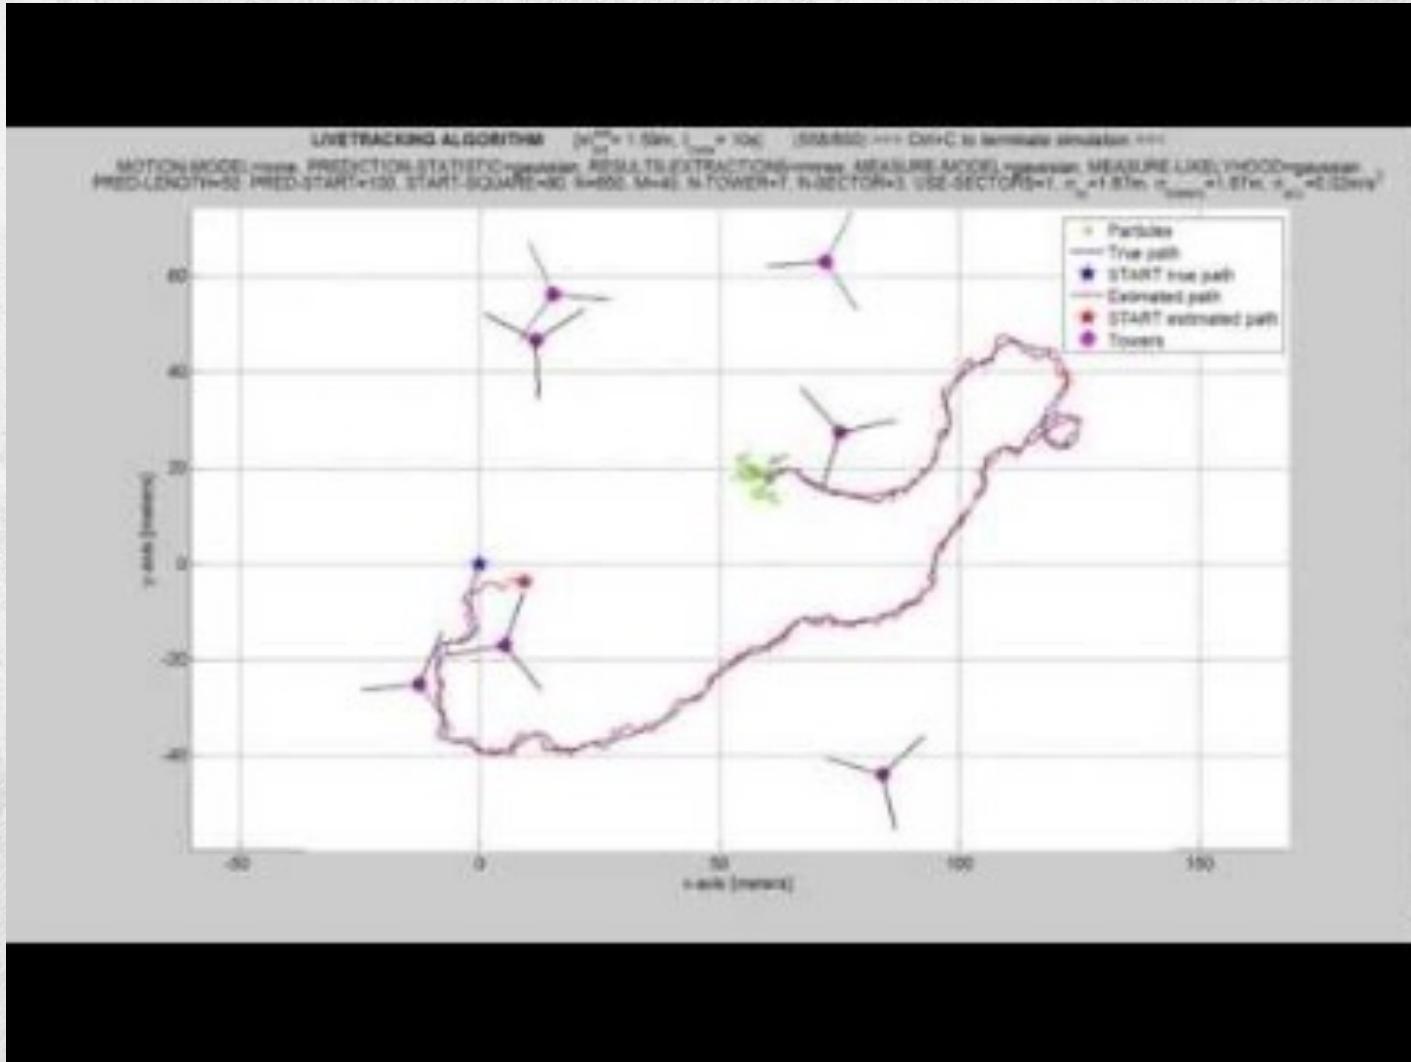


Particle Filter Algorithm

- Particle Filter(\mathcal{X}_{t-1} , u_t , z_t , m):
 - $\mathcal{X}'_t = \mathcal{X}_t = \emptyset$
 - for $m = 1$ to M do
 - $x_t^{[m]} = \text{sample_motion_model}(u_t, x_{t-1}^{[m]})$
 - $w_t^{[m]} = \text{measurement_model}(z_t, x_t^{[m]}, m,)$
 - $\mathcal{X}'_t = \mathcal{X}'_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$
 - endfor
 - for $m = 1$ to M do
 - draw i with probability $\propto w_i^{[m]}$
 - Assume that gives particle j
 - add copy of $x_t^{[j]}$ to \mathcal{X}_t
 - endfor
 - return \mathcal{X}_t



Particle Filter Demo

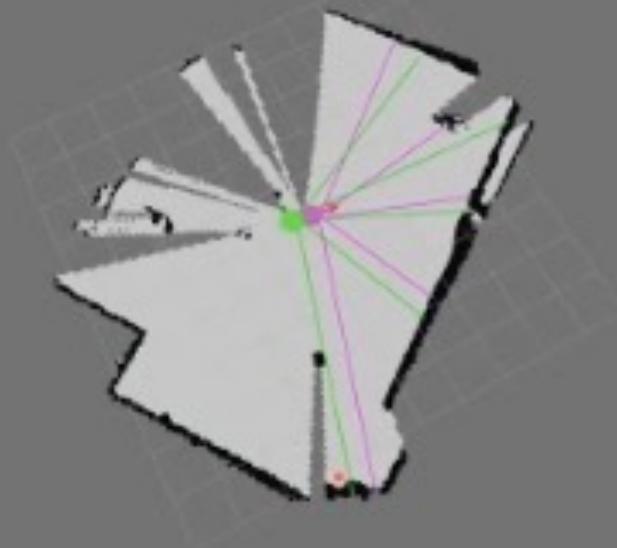


MISSISSIPPI STATE
UNIVERSITY™

Particle Filter Demo

Computing Particle Weights

Recall, we have the 'real' laser scan which the car observed (shown in green). Note we don't know where the green dot is but we do know the range measurements.



MISSISSIPPI STATE
UNIVERSITY™

References

1. https://web.mit.edu/16.412j/www/html/Advanced%20lectures/Slides/Hsaio_plinval_miller_ParticleFiltersPrint.pdf
2. <https://youtu.be/aUkBa1zMKv4>
3. <https://www.youtube.com/watch?v=HuWNhlpF4Ps>

