

AI Robotics

Path Planning



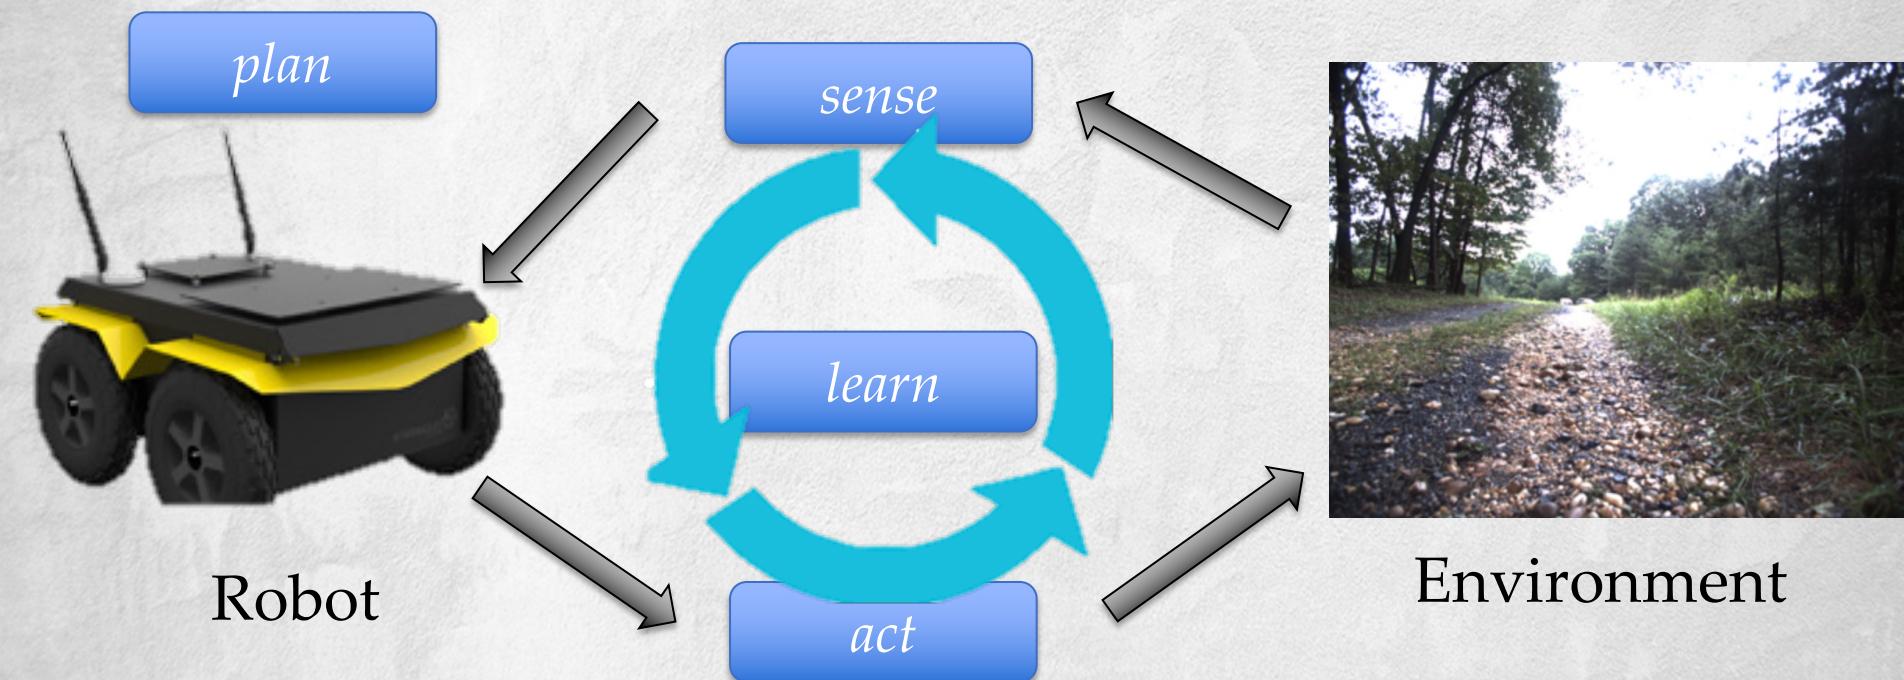
MISSISSIPPI STATE
UNIVERSITY™

Overview



MISSISSIPPI STATE
UNIVERSITY™

Robot Model



MISSISSIPPI STATE
UNIVERSITY™

Robot Model



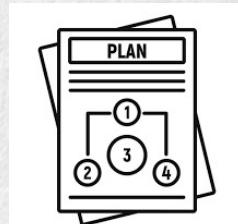
sensing ✓



localization ✓



mapping ✓



planning ?



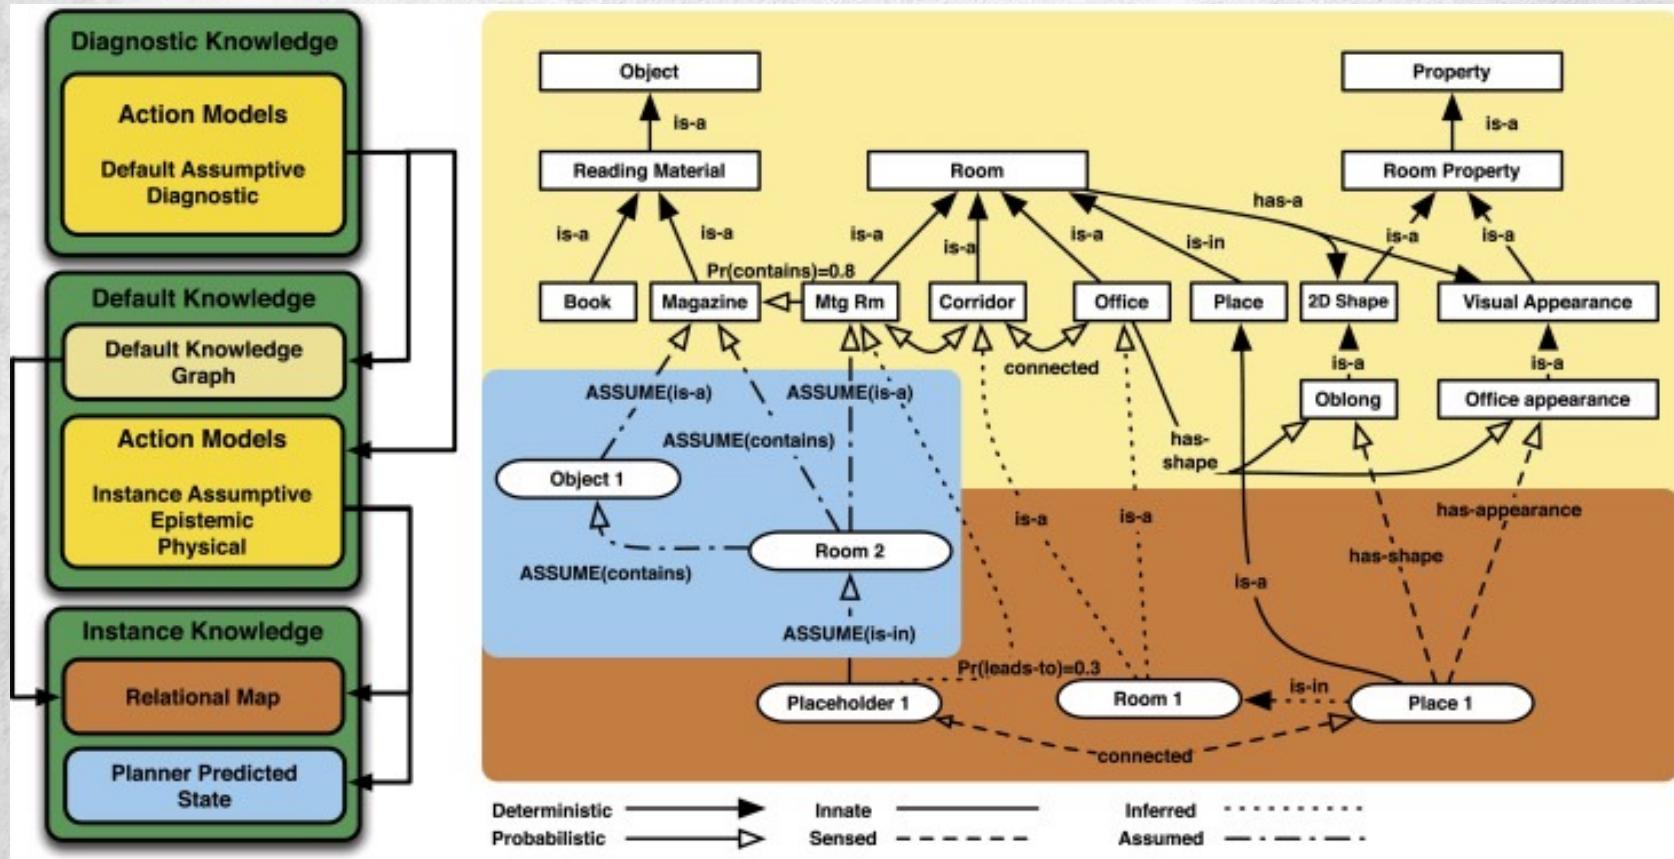
MISSISSIPPI STATE
UNIVERSITY™

Planning

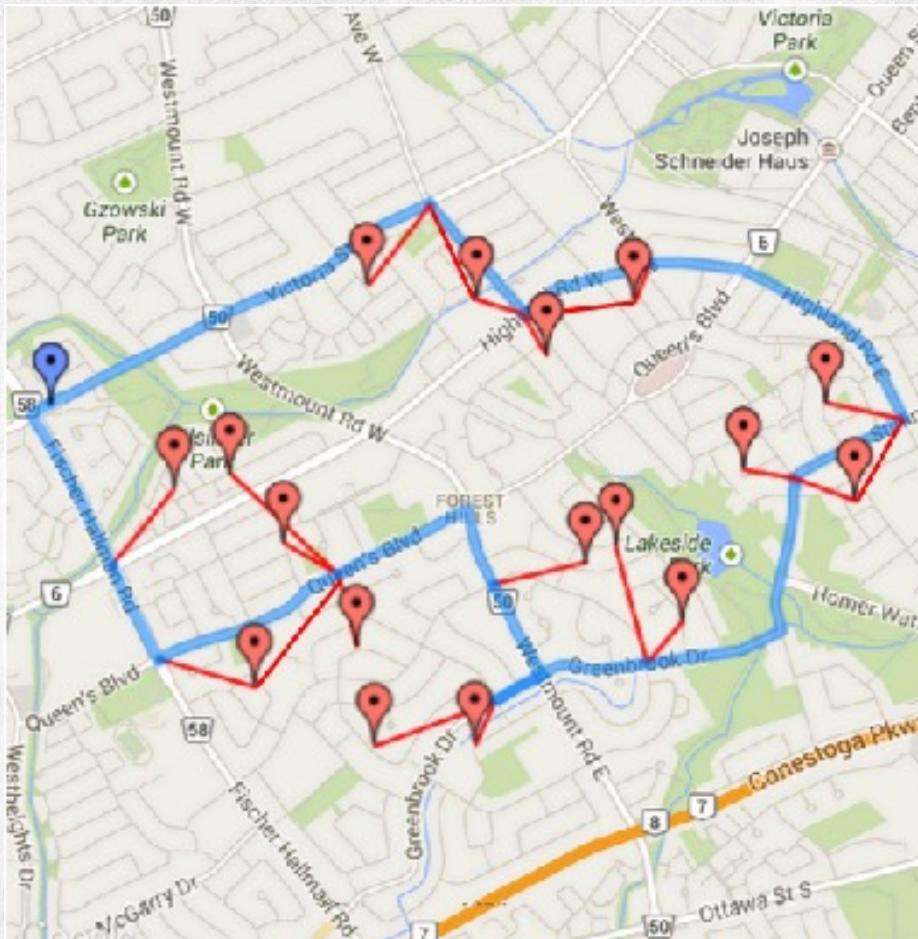
- Planning consists of any high level decision-making or problem-solving that the robot has to perform
- What might the planning portion of a robot need to know/do?
 - What tasks do I have to perform (Mission Planning)
 - What is the best way to get there? (Path planning)
 - How can I get in the correct configuration? (Motion planning)



Mission Planning

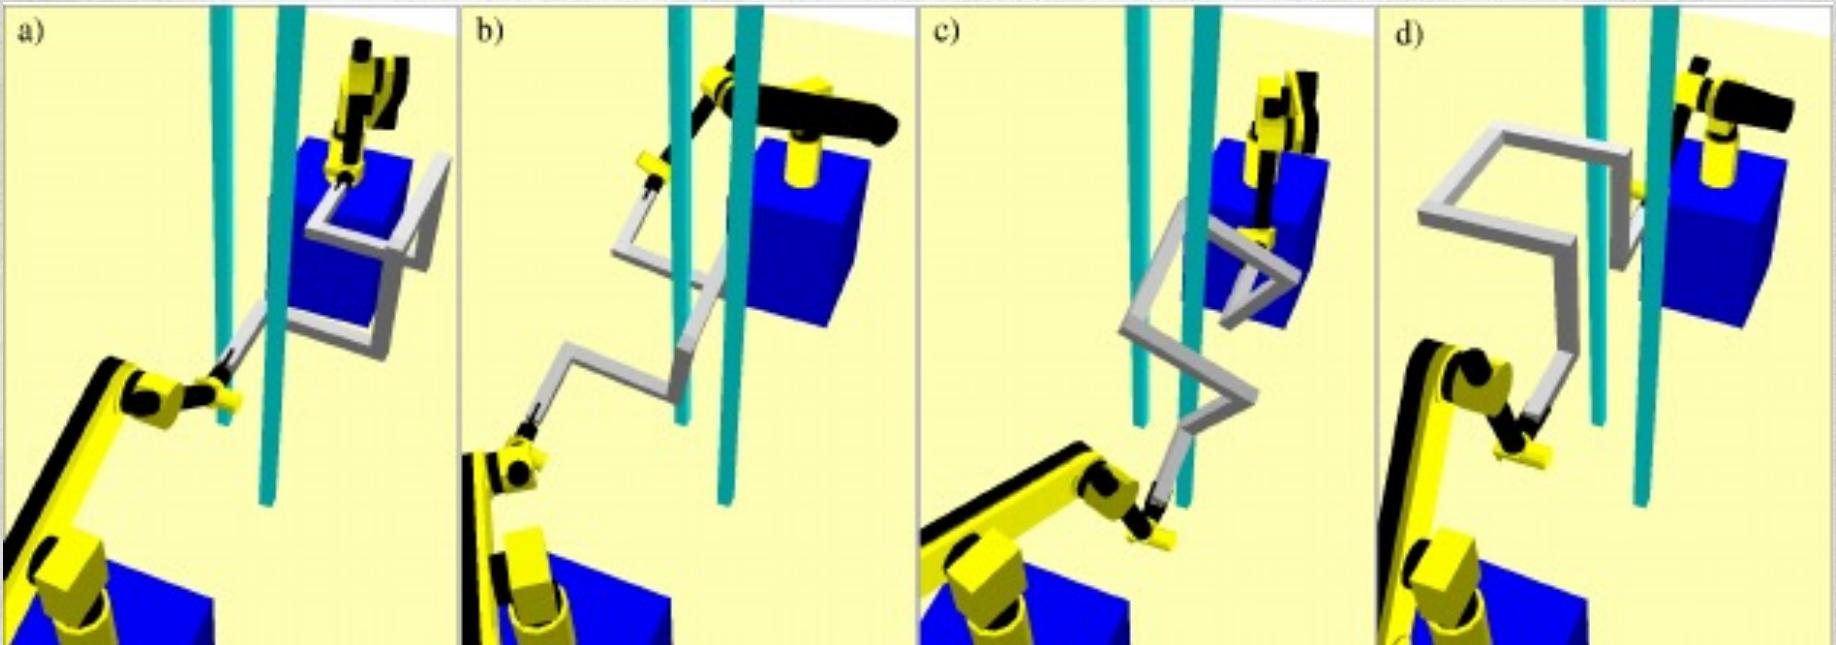


Path Planning



MISSISSIPPI STATE
UNIVERSITY™

Motion Planning



MISSISSIPPI STATE
UNIVERSITY™

Planning

- Mission planning
 - Which tasks to perform and what order to perform them to achieve a specified goal
 - Tasks may have preconditions and effects
- Path-planning
 - Focus is on robot's (x, y, θ) pose
 - What is it now?
 - What do we want it to be?
 - What sequence of poses will get me from where it is now to where it should be?
- Motion-planning
 - Focus also on robot's pose (called a configuration), but in general this can be a high-dimensional pose, consisting of many joint angles (think a robotic arm)



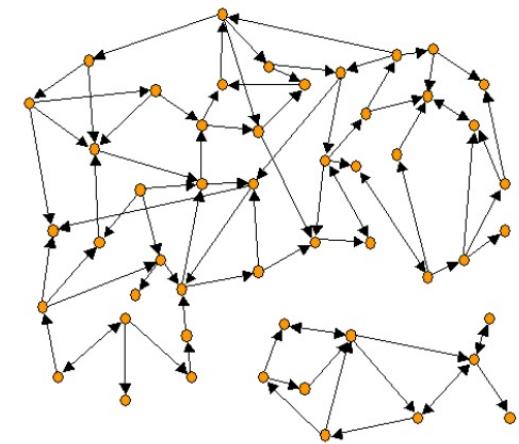
Problem Representation



MISSISSIPPI STATE
UNIVERSITY™

Search space

- Vertices – states
- Edges – connect neighboring states



- Planning algorithm needs to find a *solution* or a path from the initial state to the goal state
- The *cost* of a solution is the sum of the costs of each edge in the path
 - For our purposes, this will be the length of the path
- An *optimal solution* is a solution with the lowest cost of any solution

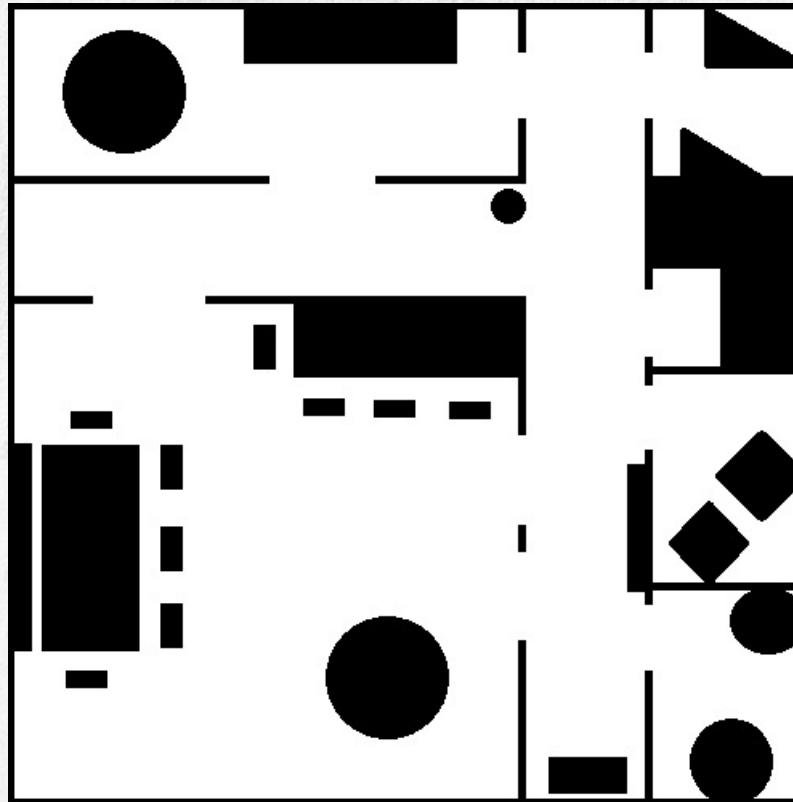


Search Space

- The **Search Space** for a planning problem is the set of all poses/configurations/states that we will planning in
- In robotics problems the search spaces are typically *continuous*. This means that there is an uncountably infinite number of possible states
- The planning algorithms we will discuss today require a *discrete* state space to work with
- How can we come up with a discrete version of a continuous state space?



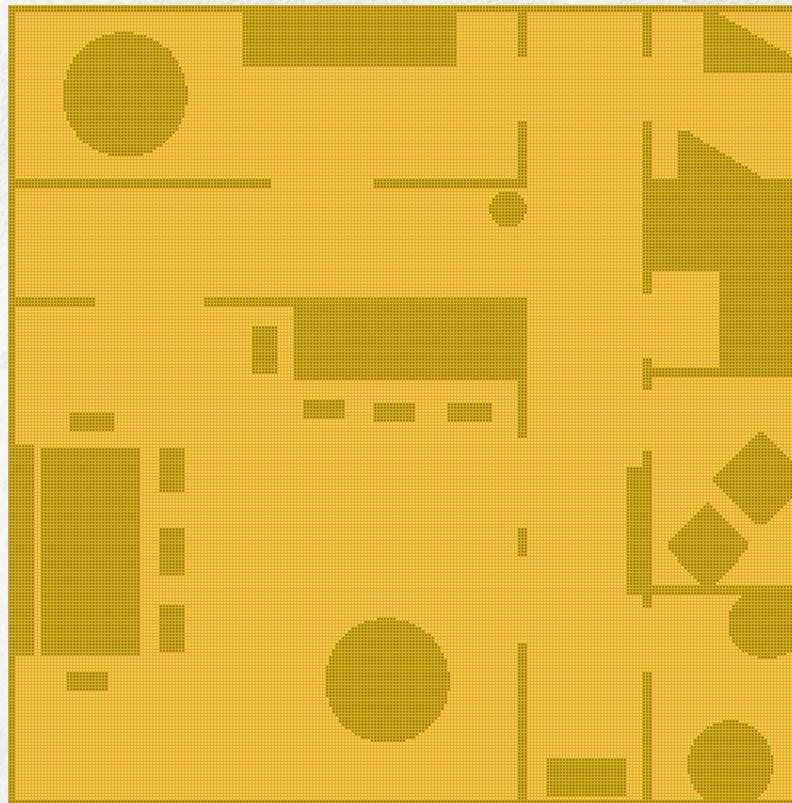
Original Map



MISSISSIPPI STATE
UNIVERSITY™

Grid-based discretization

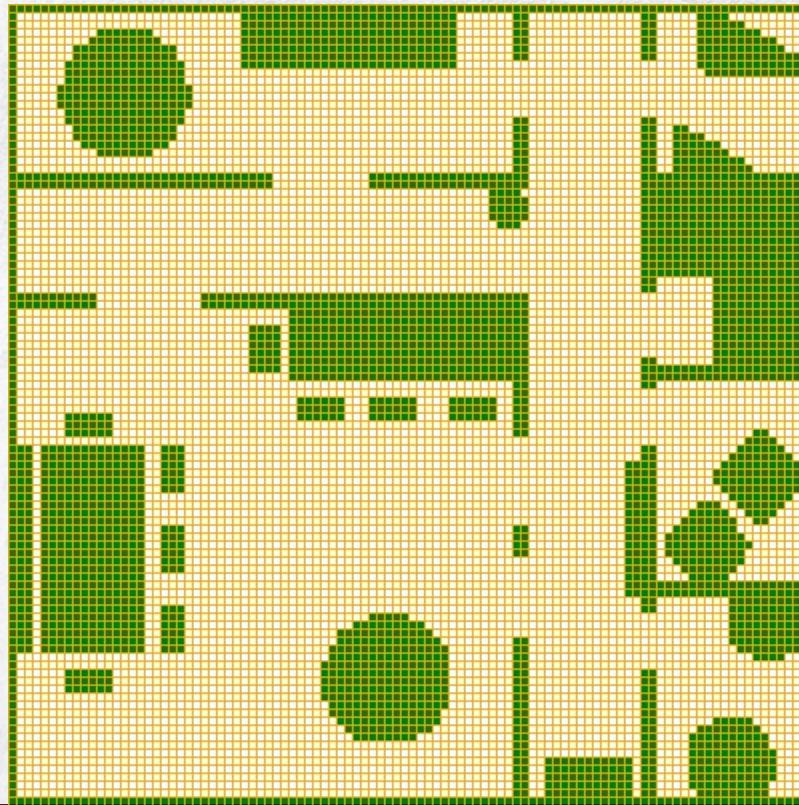
2 pix per cell =
61752 xy-cells



MISSISSIPPI STATE
UNIVERSITY™

Grid-based Discretization

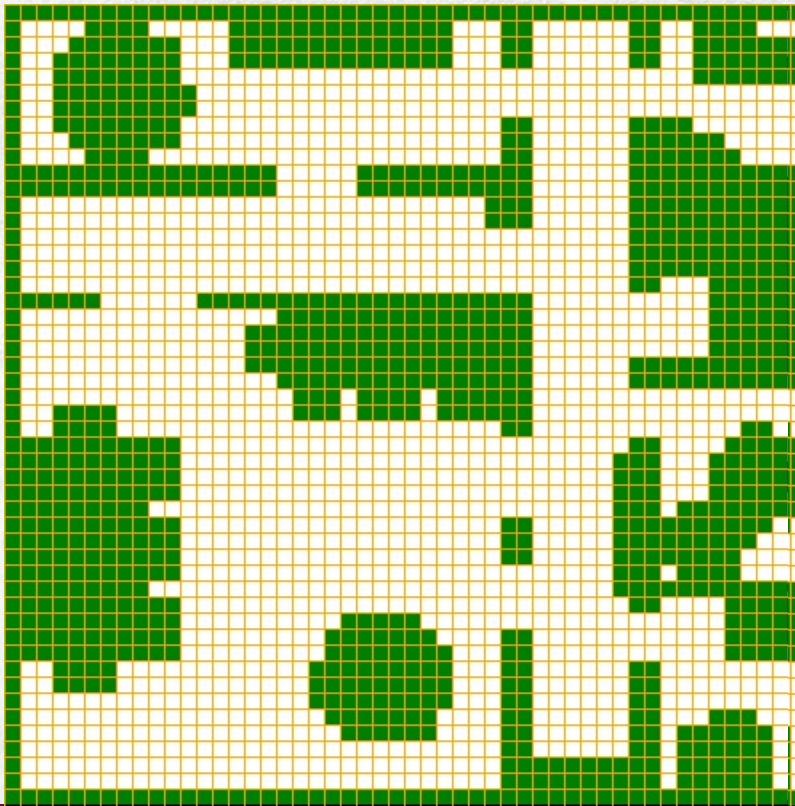
5 pix per cell =
9702 xy-cells



MISSISSIPPI STATE
UNIVERSITY™

Grid-based Discretization

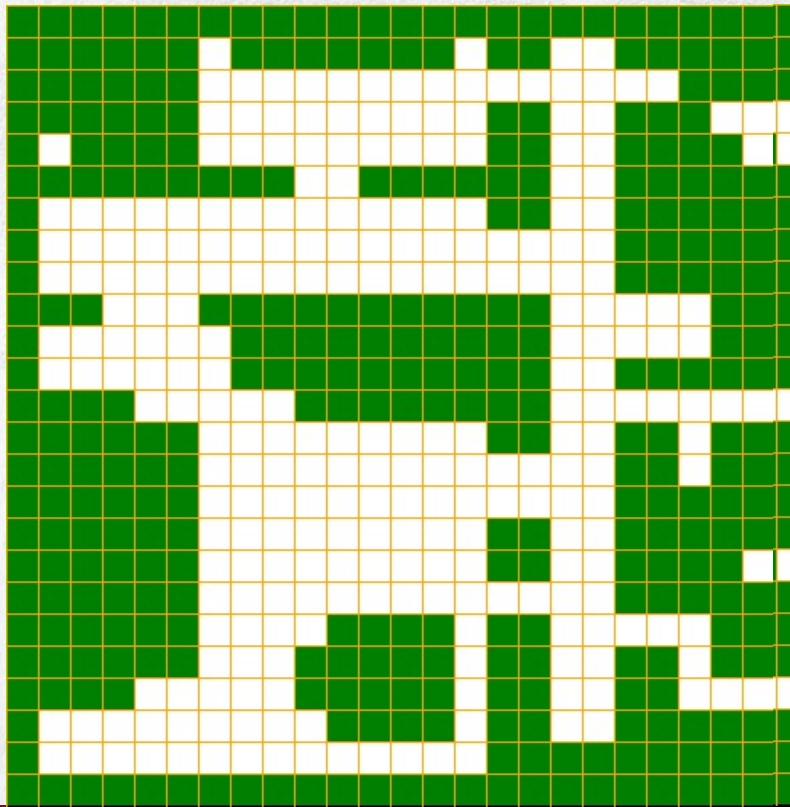
10 pix per cell =
2352 xy-cells



MISSISSIPPI STATE
UNIVERSITY™

Grid-based Discretization

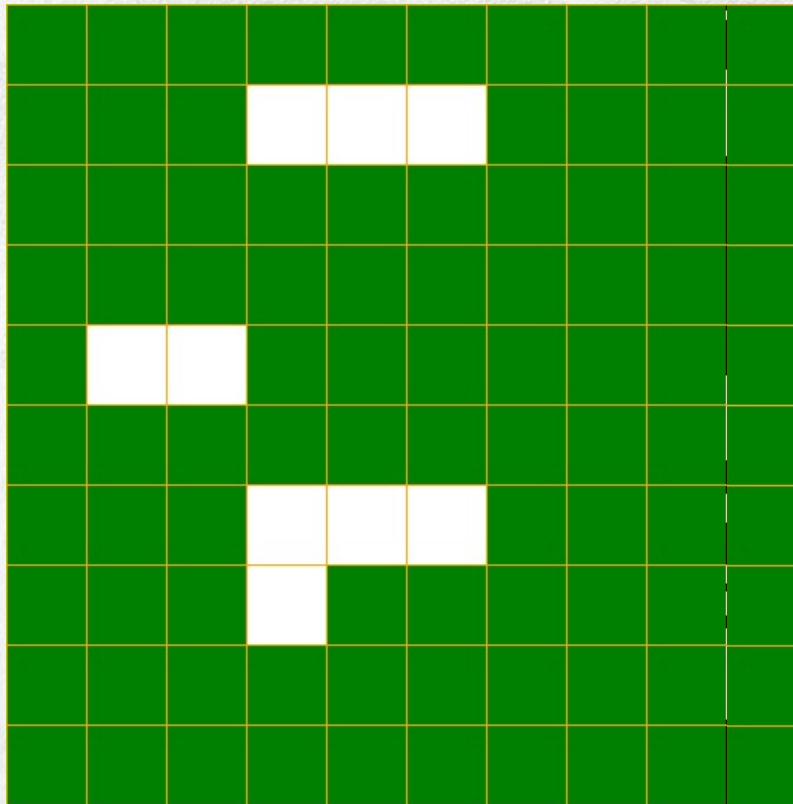
20 pix per cell =
552 xy-cells



MISSISSIPPI STATE
UNIVERSITY™

Grid-based Discretization

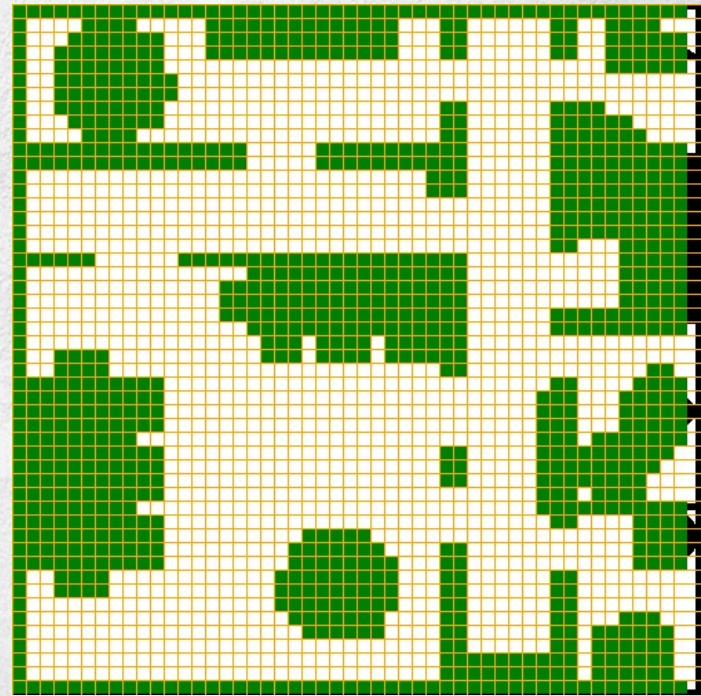
50 pix per cell =
72 xy-cells



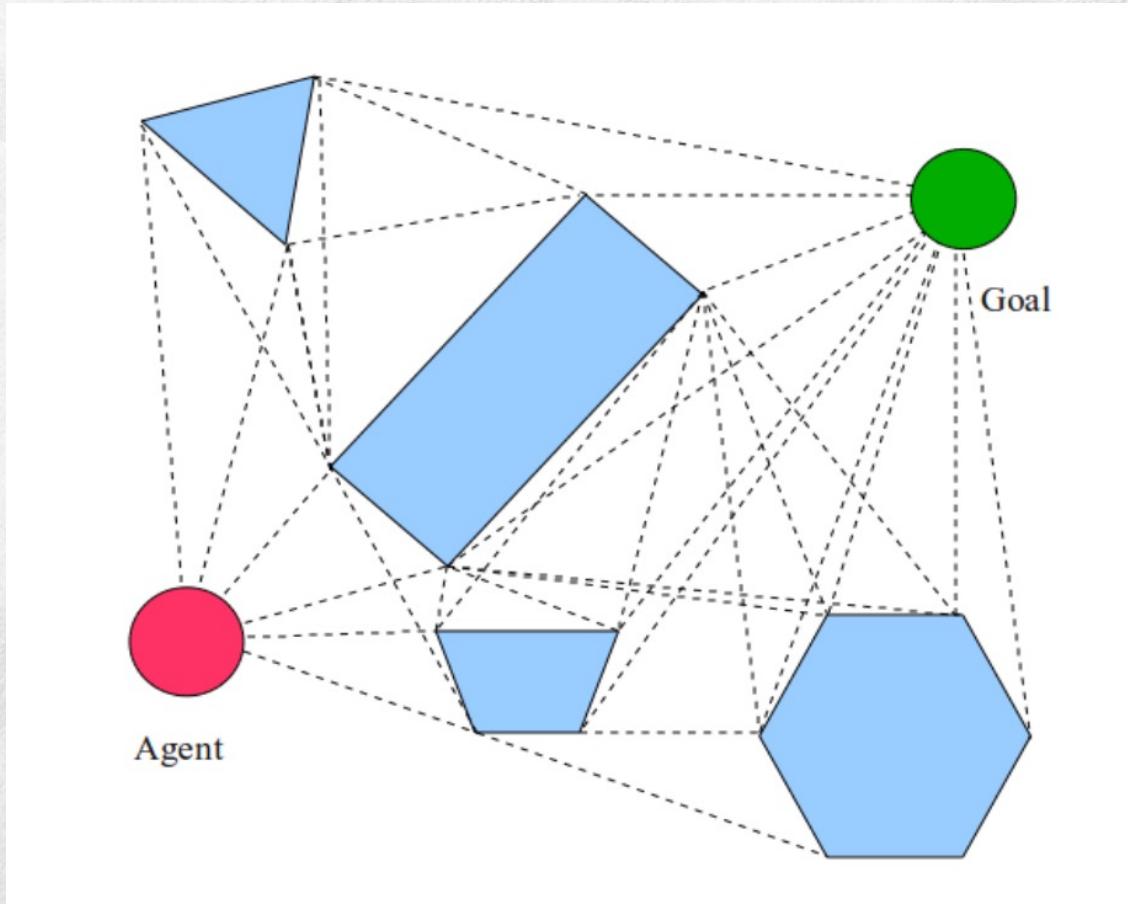
MISSISSIPPI STATE
UNIVERSITY™

Grid-based discretization

- Advantages
 - Easy to compute
 - Conceptually straightforward
- Disadvantages
 - Doesn't scale well
 - Ignores features of environment
 - Might make possible paths impossible if resolution is too large

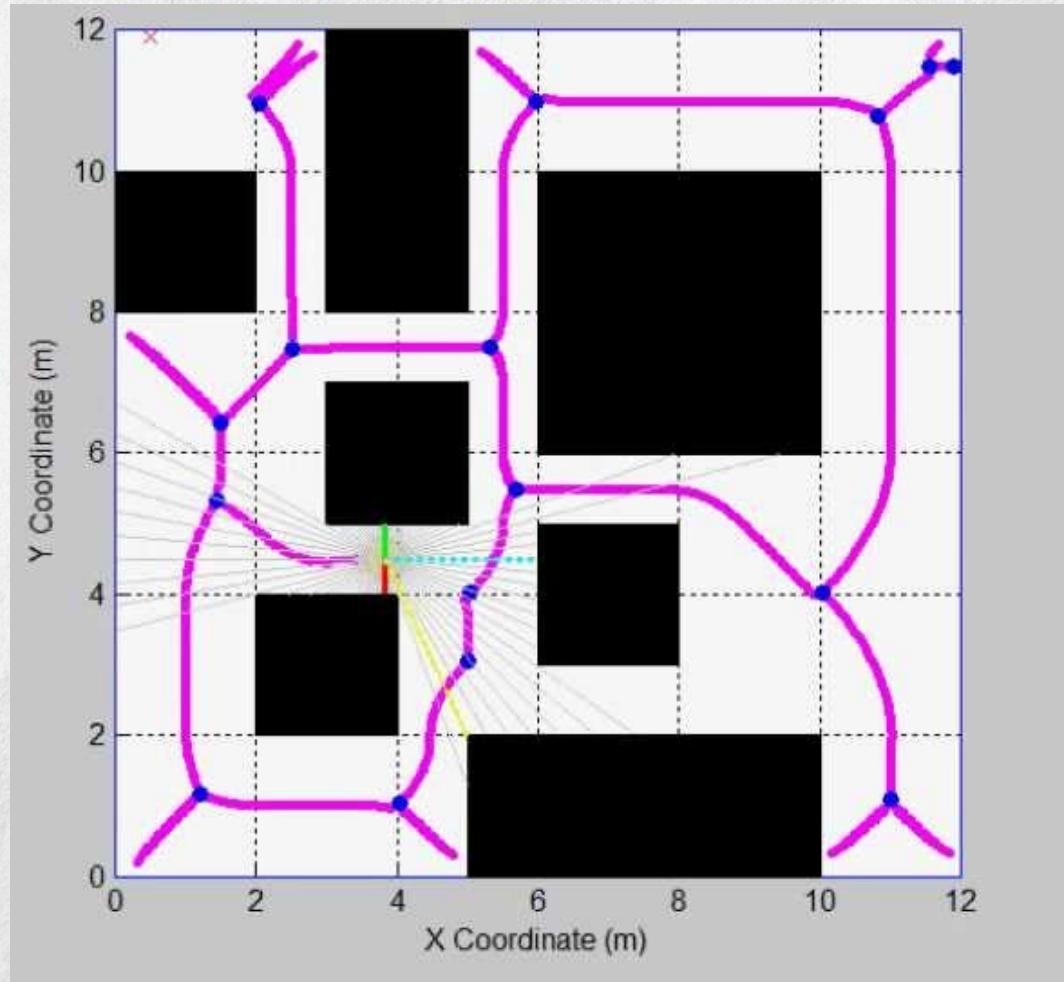


Visibility Graph



MISSISSIPPI STATE
UNIVERSITY™

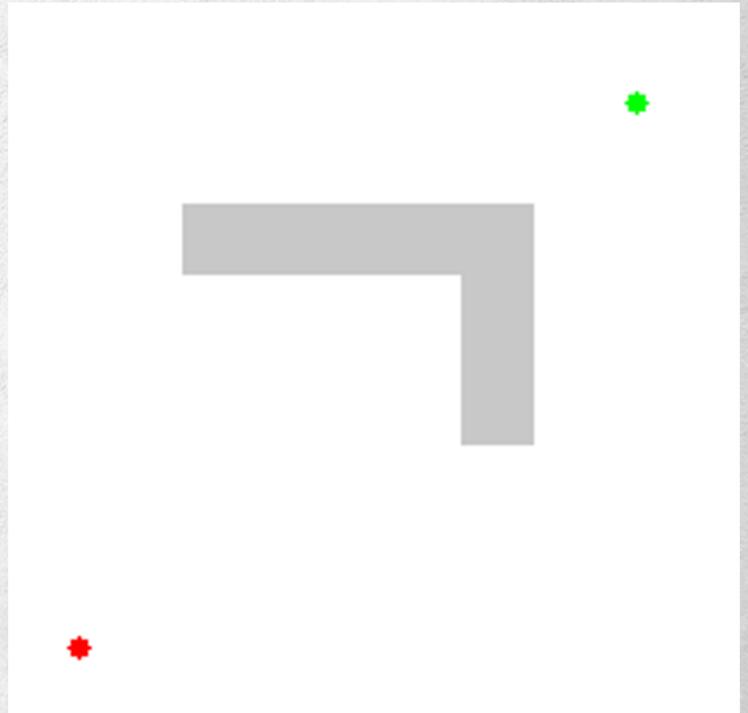
Voronoi Diagram



MISSISSIPPI STATE
UNIVERSITY™

Search-Based Planning

1. Make a set of points
 - Each point = one possible robot state
2. Make connections between neighboring points
 - Forming a graph
3. Remove any connections that go through obstacles
4. Search for a path through the graph



Search Algorithms



MISSISSIPPI STATE
UNIVERSITY™

Search Space and Planning Algorithms

- So we have our discretized search space
- How do our algorithms interact with it?
- A search space provides:
 - An initial state
 - For any state, a set of neighboring states, that can be easily accessed from that state
 - i.e. neighbor grid cells
 - This will be impacted by how you decide to do it
 - All 8 neighbors on grid? Only up-down/left-right?
 - The cost for getting from any state to a neighbor
 - Test to see if a state is a goal (goal-test)



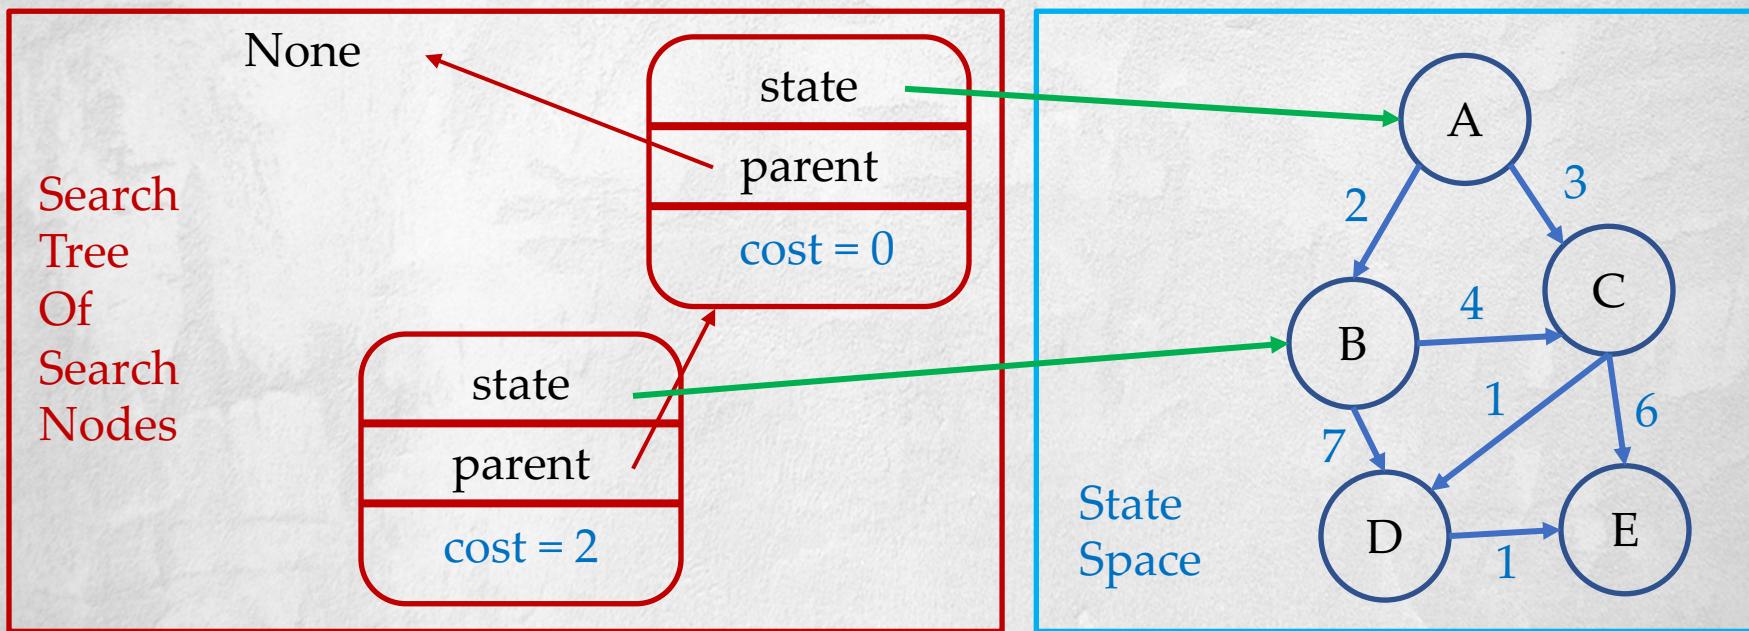
Search Algorithms

- The algorithms we discuss today will all work in a similar way
- They will each build up a tree of *search nodes*, where each search node corresponds to a path in the search space
- The algorithms differ in how they decide which search node, or partial path, to further explore at each time step
- Search will stop once a partial path has reached the goal state

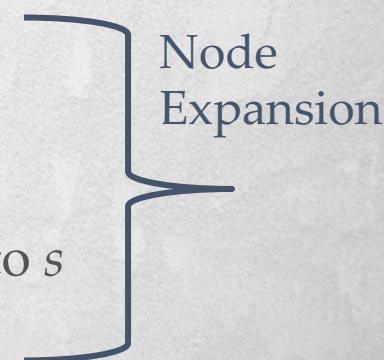


Search Nodes

- These are data structures that have:
 - A search state they are pointing to
 - A pointer to their parent search node (tracing these from a given search node to the initial search node gives us the partial path represented by the search node)
 - The cost of the path from the initial node to this search node

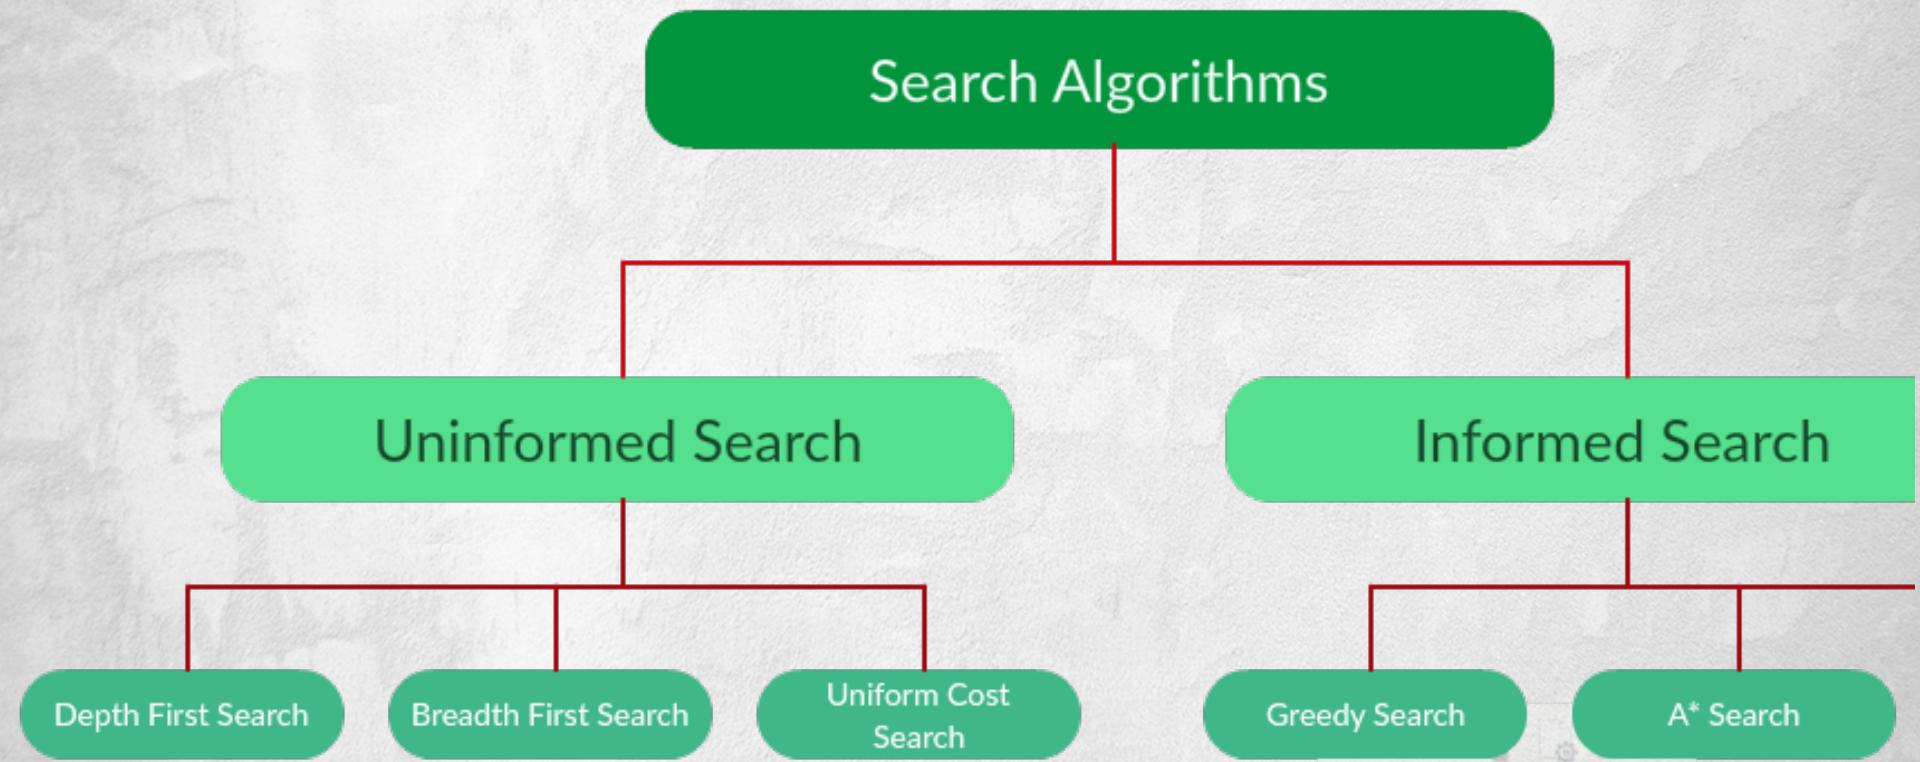


Basic Search Algorithm

- Create START node **Node generation**
 - State is initial state
 - Cost is 0
 - Parent pointer is None
 - Insert START into FRINGE (queue which stores search nodes)
 - While FRINGE is not empty:
 - Get *node* from FRINGE according to strategy
 - Is *node.state* a goal? If Yes, then we are done (return *node*)
 - If No, then for each neighboring state *s* of *node.state*:
 - Create search node *n* with **Node generation**
 - State = *s*
 - Cost = *node.cost* + cost of getting from *node.state* to *s*
 - Parent pointer = *node*
- 
- Node Expansion



Search Algorithm Categories

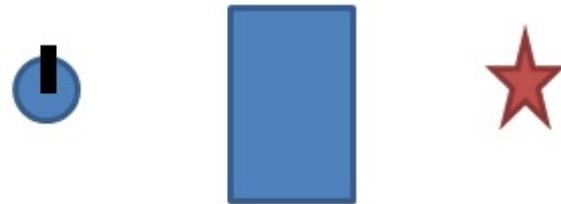


Uninformed Search Algorithms

- Uninformed = Blind = no information about each state (other than goal test)
- Depth-First Search
 - FRINGE is a LIFO queue / stack
 - Node that was added most recently is the next to explore
- Breadth-First Search
 - FRINGE is a FIFO queue
 - Node that has been in the longest is the next to explore
- Uniform-Cost Search
 - FRINGE is a priority queue, sorted by cost
 - Like Breadth-First, but takes cost into account
 - Always expand the node with the lowest cost
- Which of these will always find a solution? Always find an optimal solution?

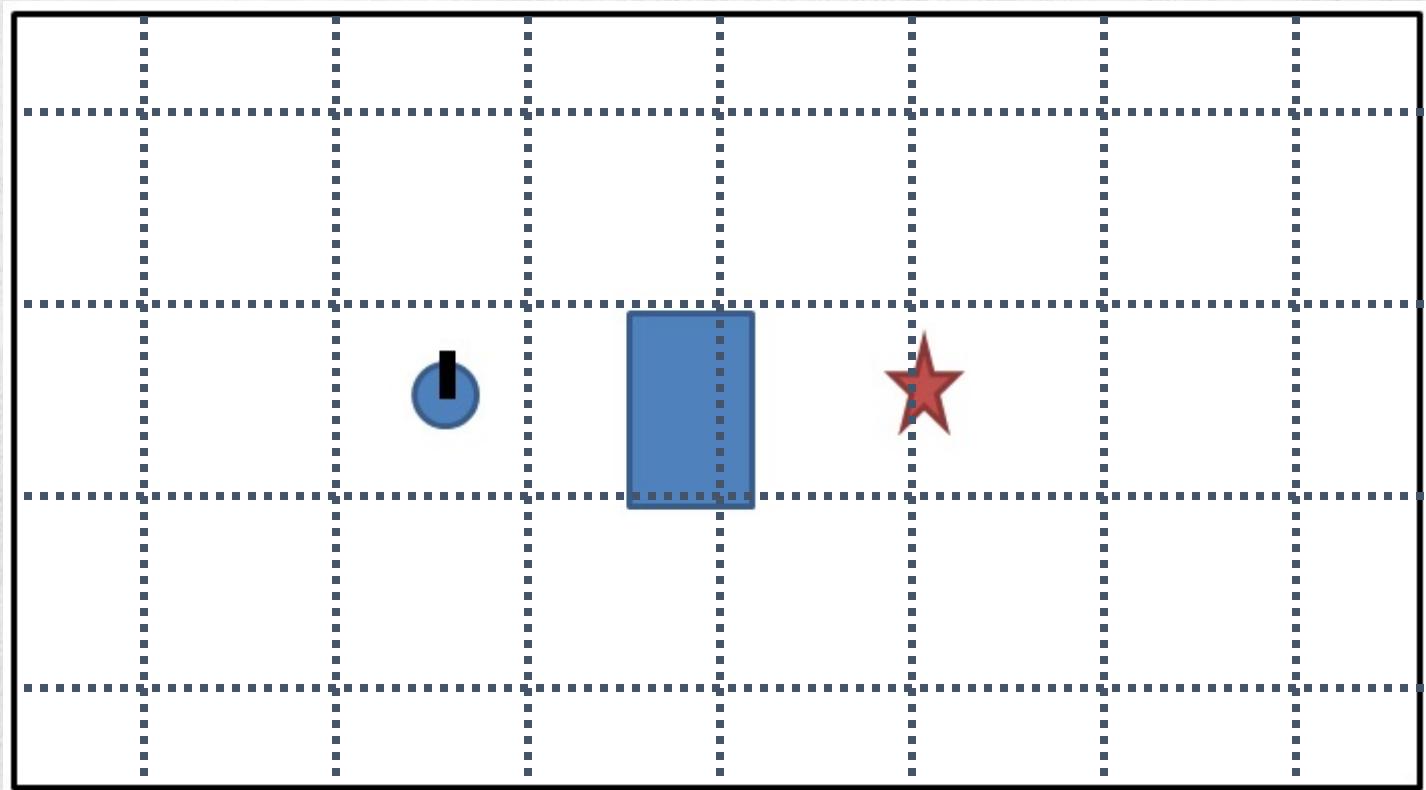


Let's work some examples



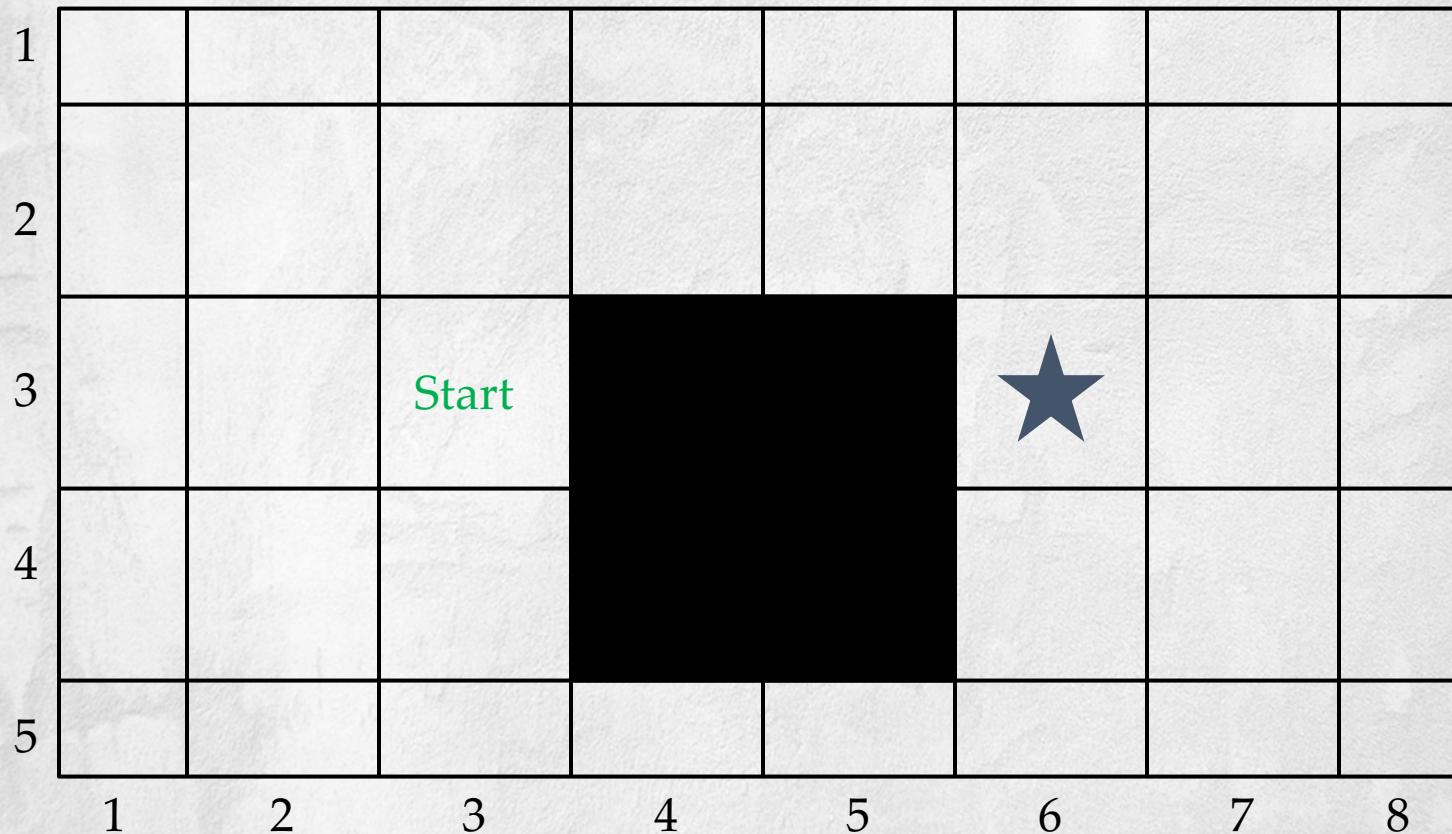
MISSISSIPPI STATE
UNIVERSITY™

Discretize the search space



MISSISSIPPI STATE
UNIVERSITY™

Discretized search space



Start = (3,3) Goal = (6,3)

Neighbors = N,E,S,W (in that order)



MISSISSIPPI STATE
UNIVERSITY™

Depth-First Search

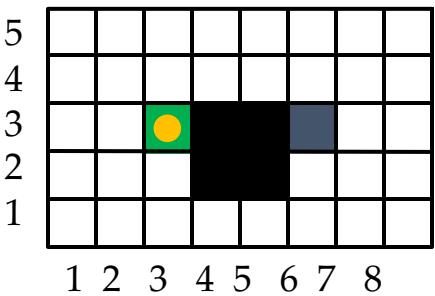


MISSISSIPPI STATE
UNIVERSITY™

Depth First Search

```
procedure DFS_iterative( $G$ ,  $v$ ) is
    let  $S$  be a stack
     $S.push(v)$ 
    while  $S$  is not empty do
         $v = S.pop()$ 
        if  $v$  is not labeled as discovered then
            label  $v$  as discovered
            for all edges from  $v$  to  $w$  in  $G.adjacentEdges(v)$  do
                 $S.push(w)$ 
```

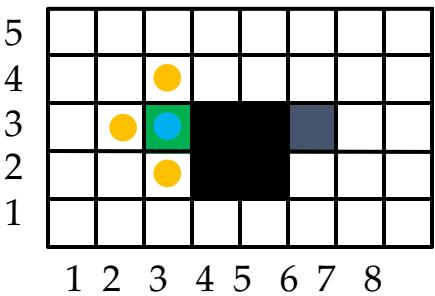




Search Tree: Depth-first Search (no duplicates in fringe)

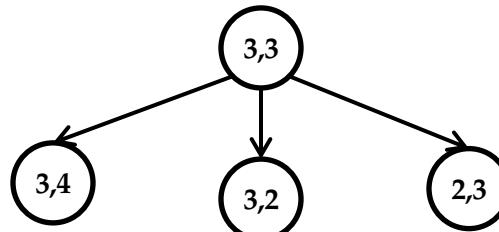


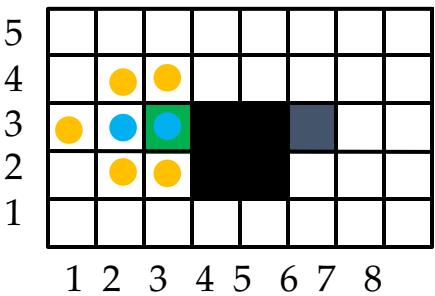
Neighbors = N,E,S,W (in that order)



Neighbors = N,E,S,W (in that order)

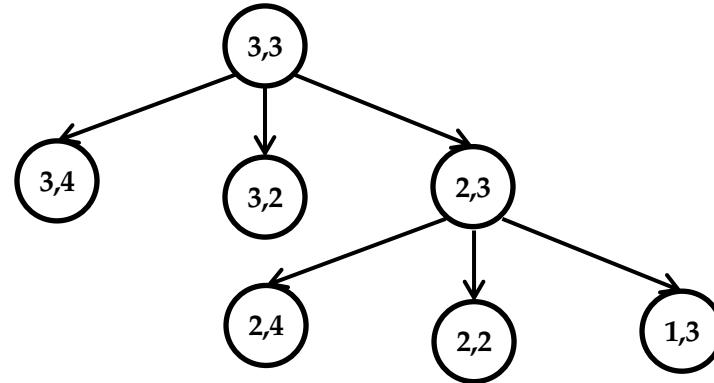
Search Tree: Depth-first Search (no duplicates in fringe)

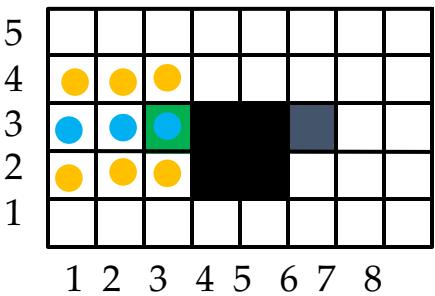




Neighbors = N,E,S,W (in that order)

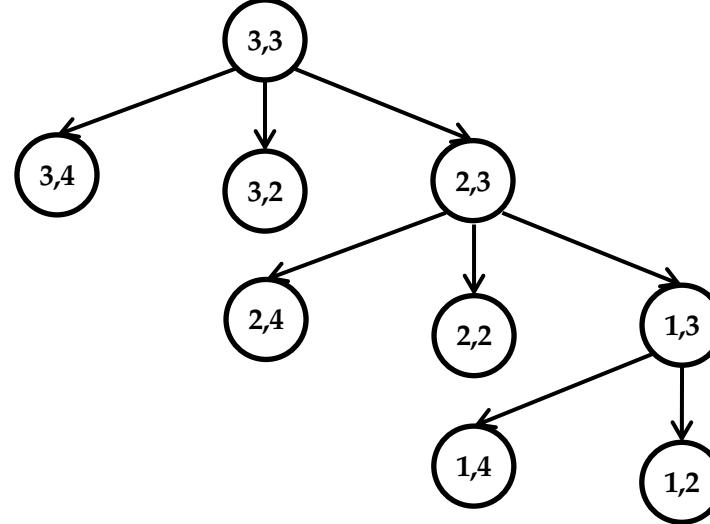
Search Tree: Depth-first Search (no duplicates in fringe)

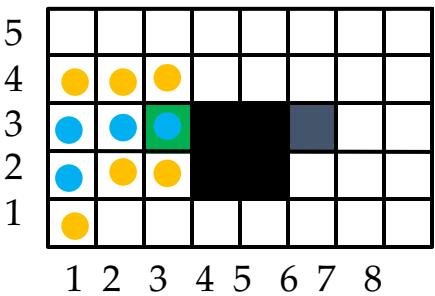




Neighbors = N,E,S,W (in that order)

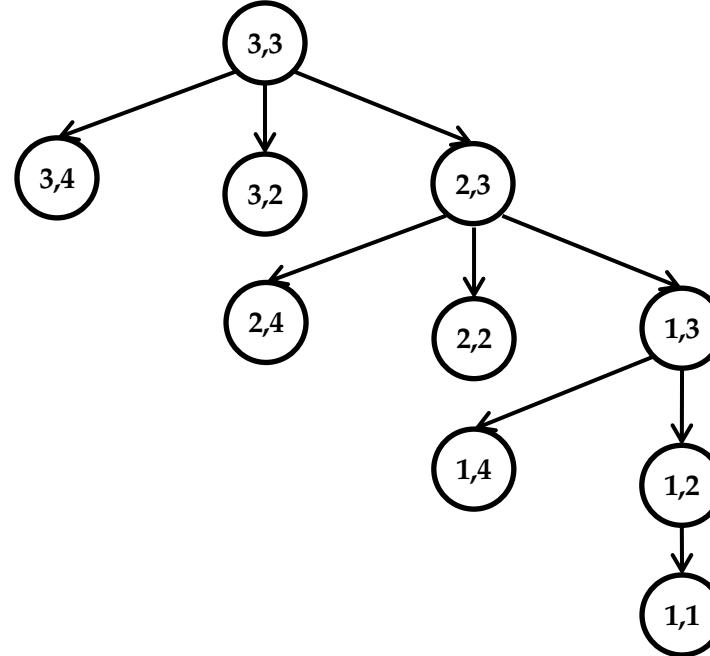
Search Tree: Depth-first Search (no duplicates in fringe)

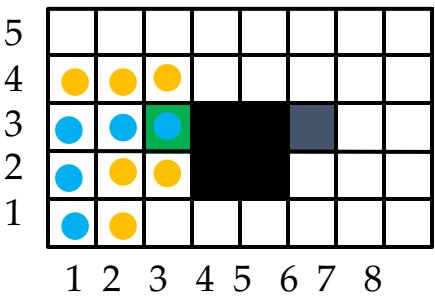




Neighbors = N,E,S,W (in that order)

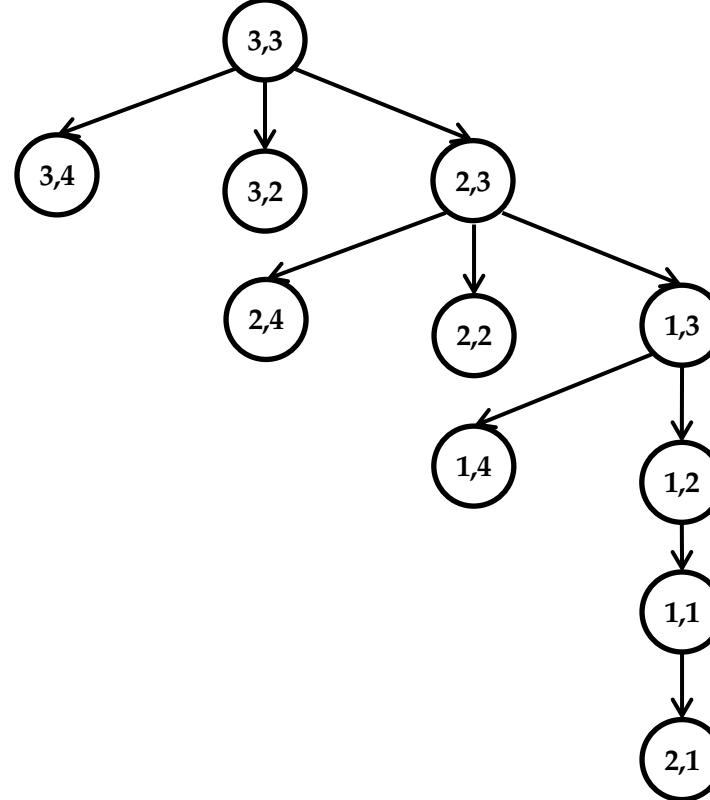
Search Tree: Depth-first Search (no duplicates in fringe)

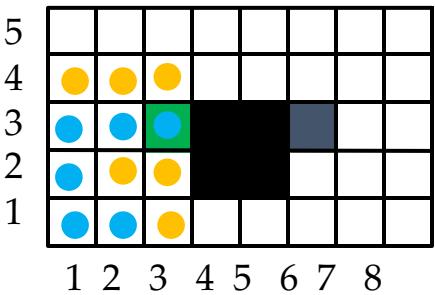




Neighbors = N,E,S,W (in that order)

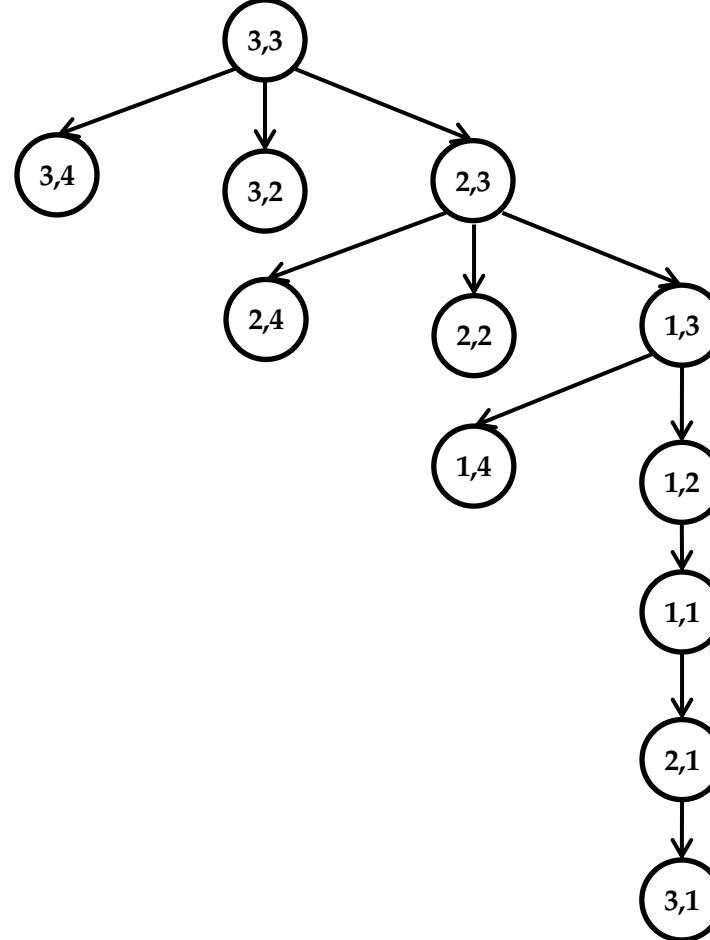
Search Tree: Depth-first Search (no duplicates in fringe)

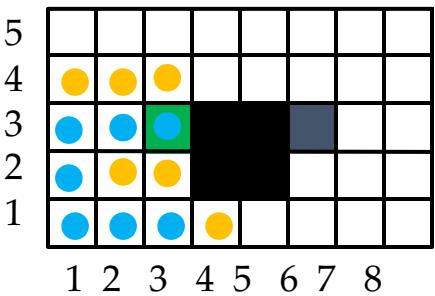




Neighbors = N,E,S,W (in that order)

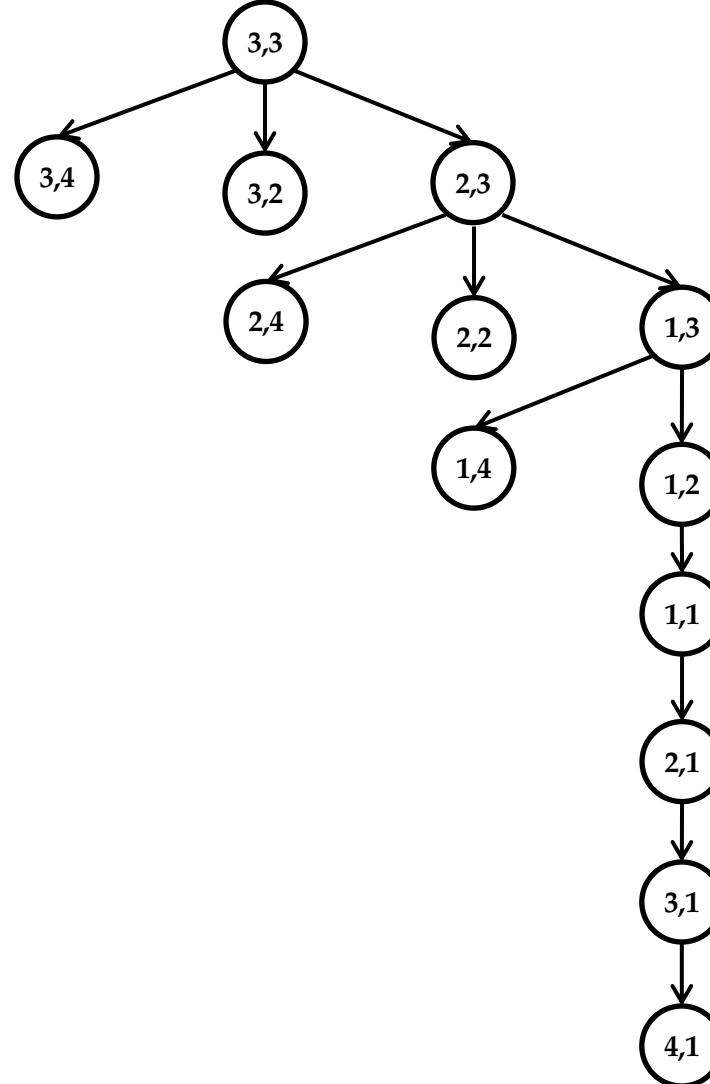
Search Tree: Depth-first Search (no duplicates in fringe)

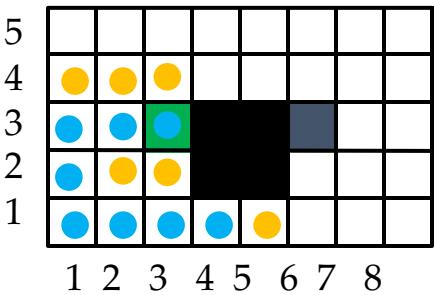




Neighbors = N,E,S,W (in that order)

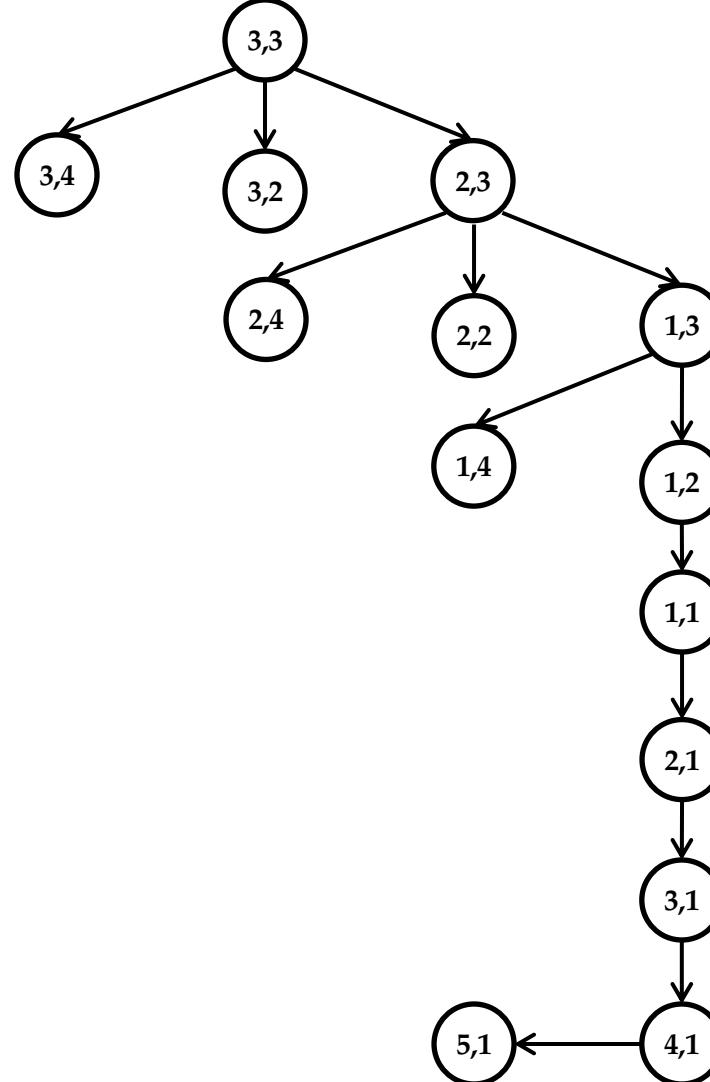
Search Tree: Depth-first Search (no duplicates in fringe)

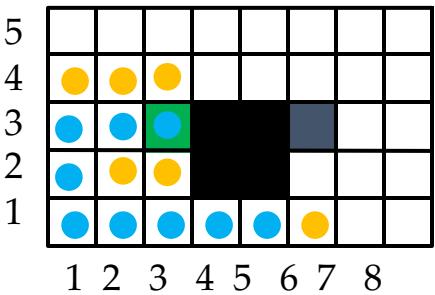




Neighbors = N,E,S,W (in that order)

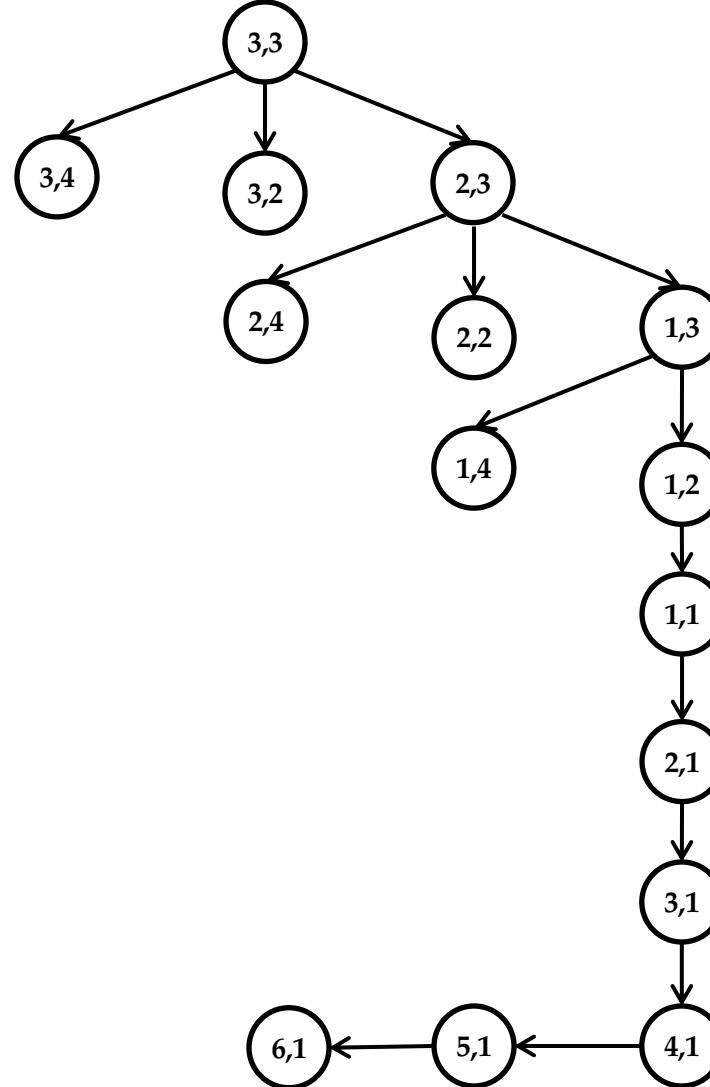
Search Tree: Depth-first Search (no duplicates in fringe)

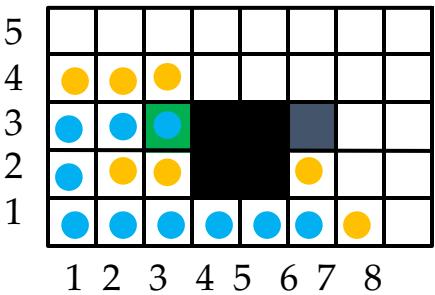




Neighbors = N,E,S,W (in that order)

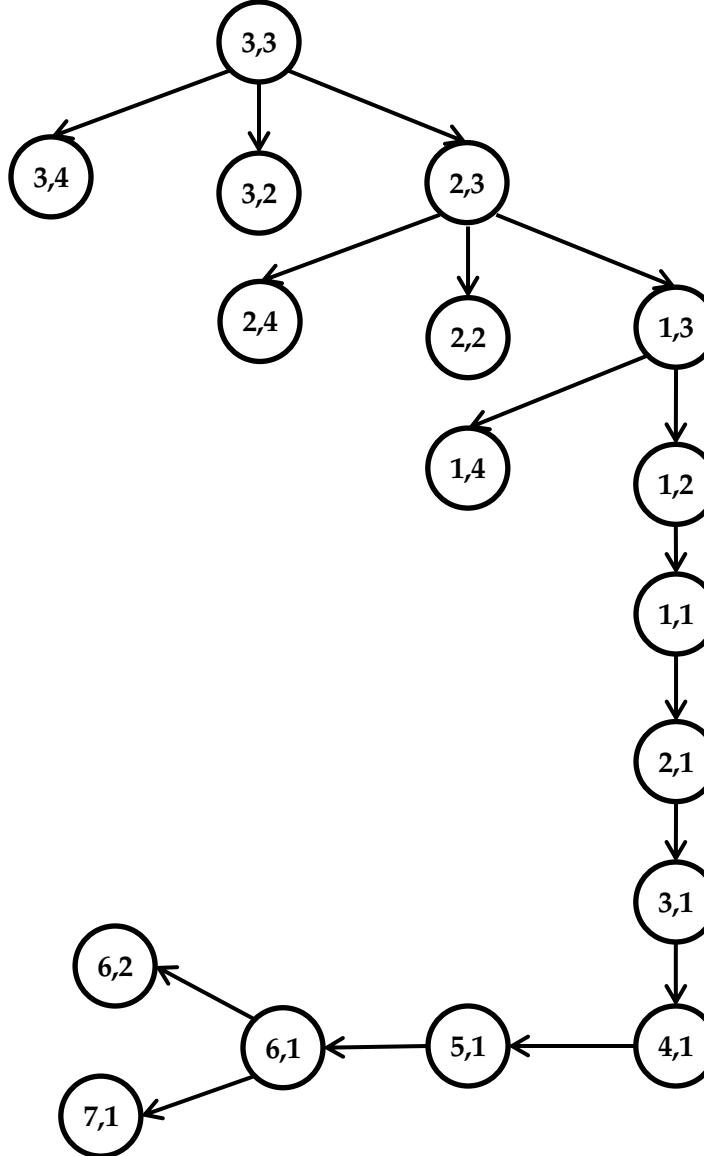
Search Tree: Depth-first Search (no duplicates in fringe)

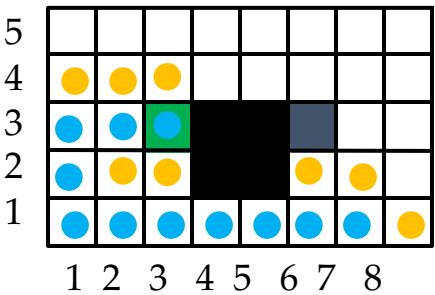




Neighbors = N,E,S,W (in that order)

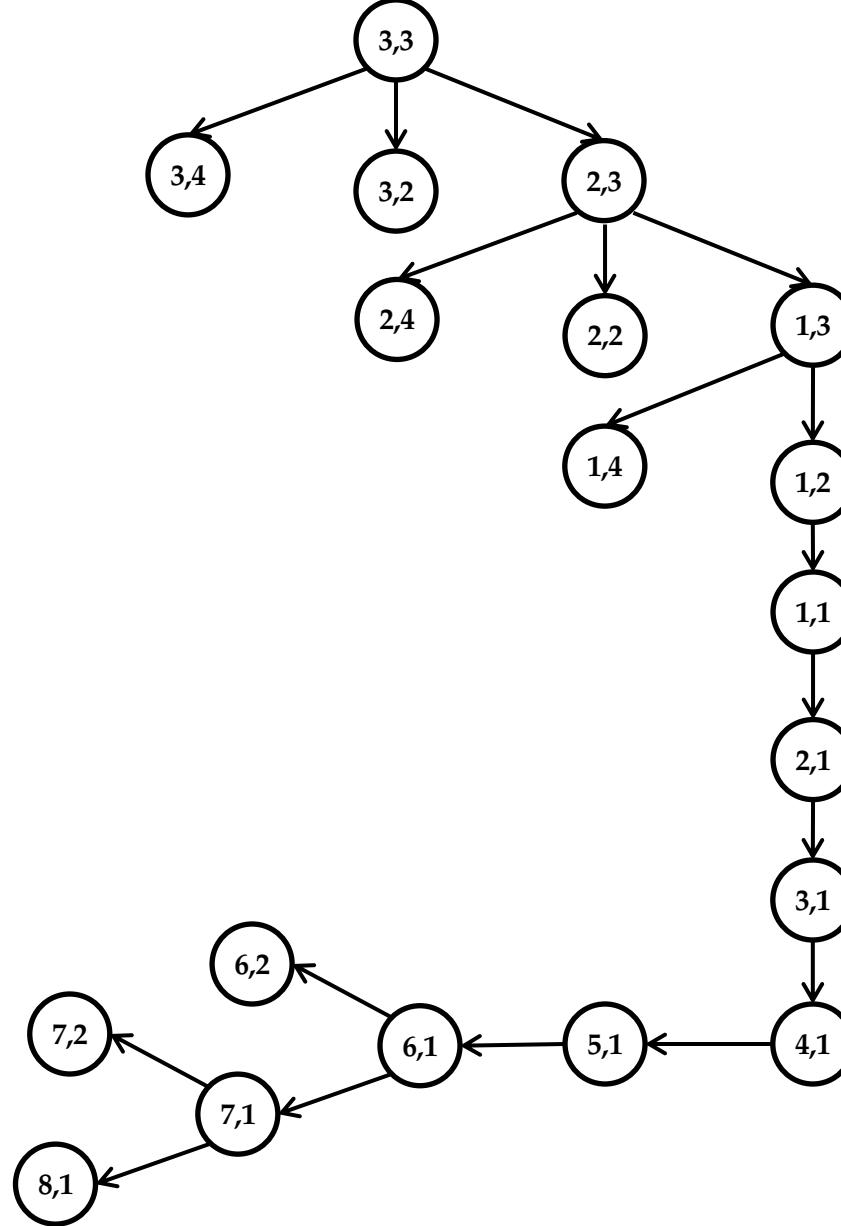
Search Tree: Depth-first Search (no duplicates in fringe)

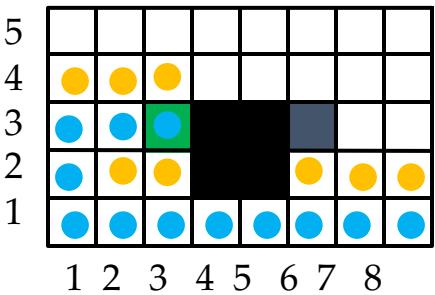




Neighbors = N,E,S,W (in that order)

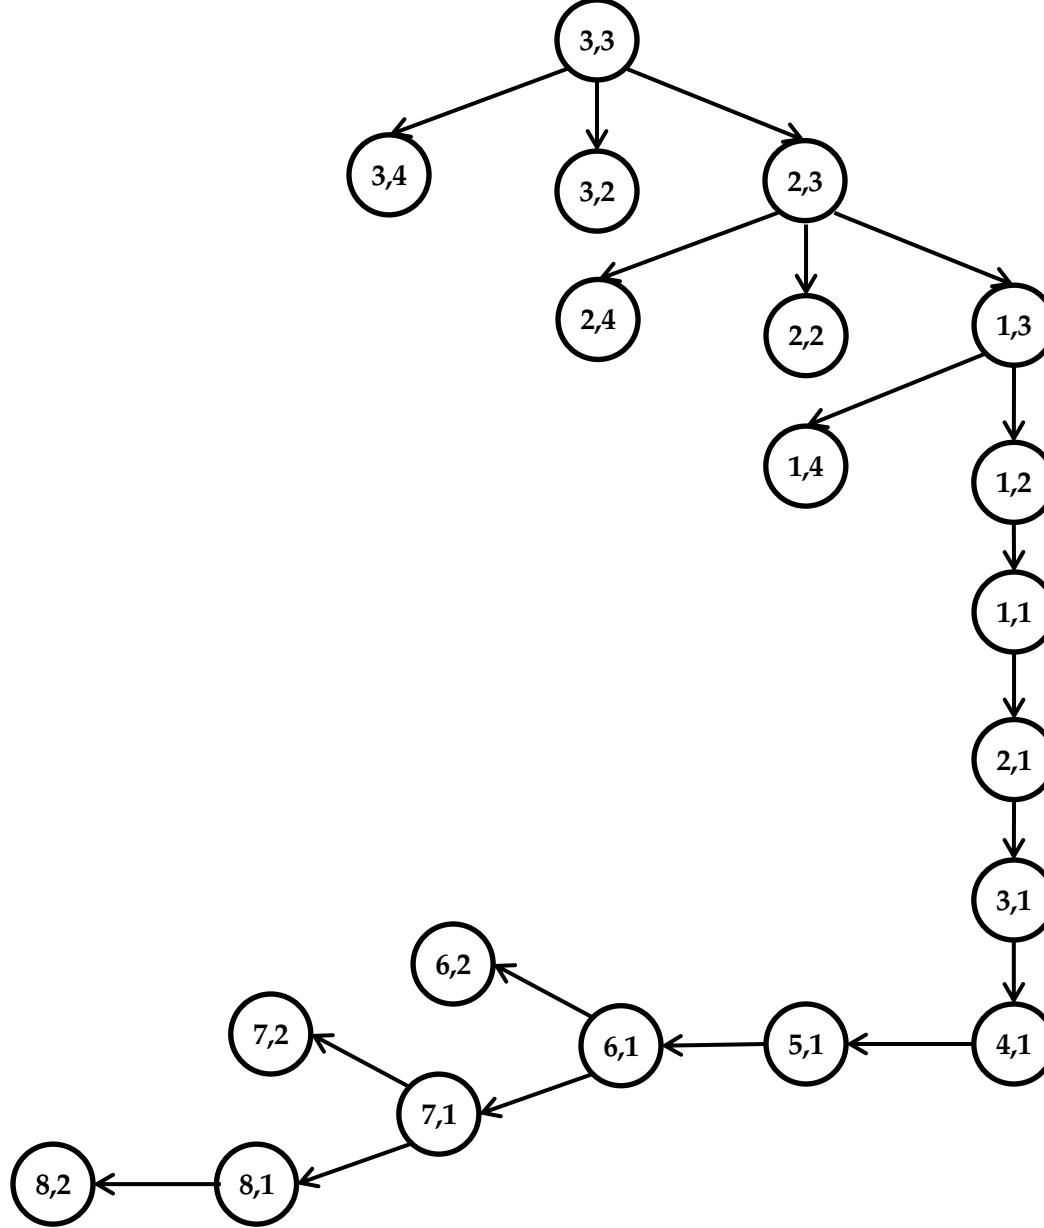
Search Tree: Depth-first Search (no duplicates in fringe)

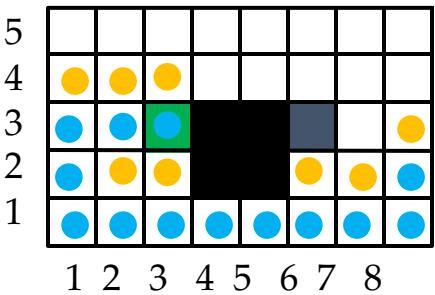




Neighbors = N,E,S,W (in that order)

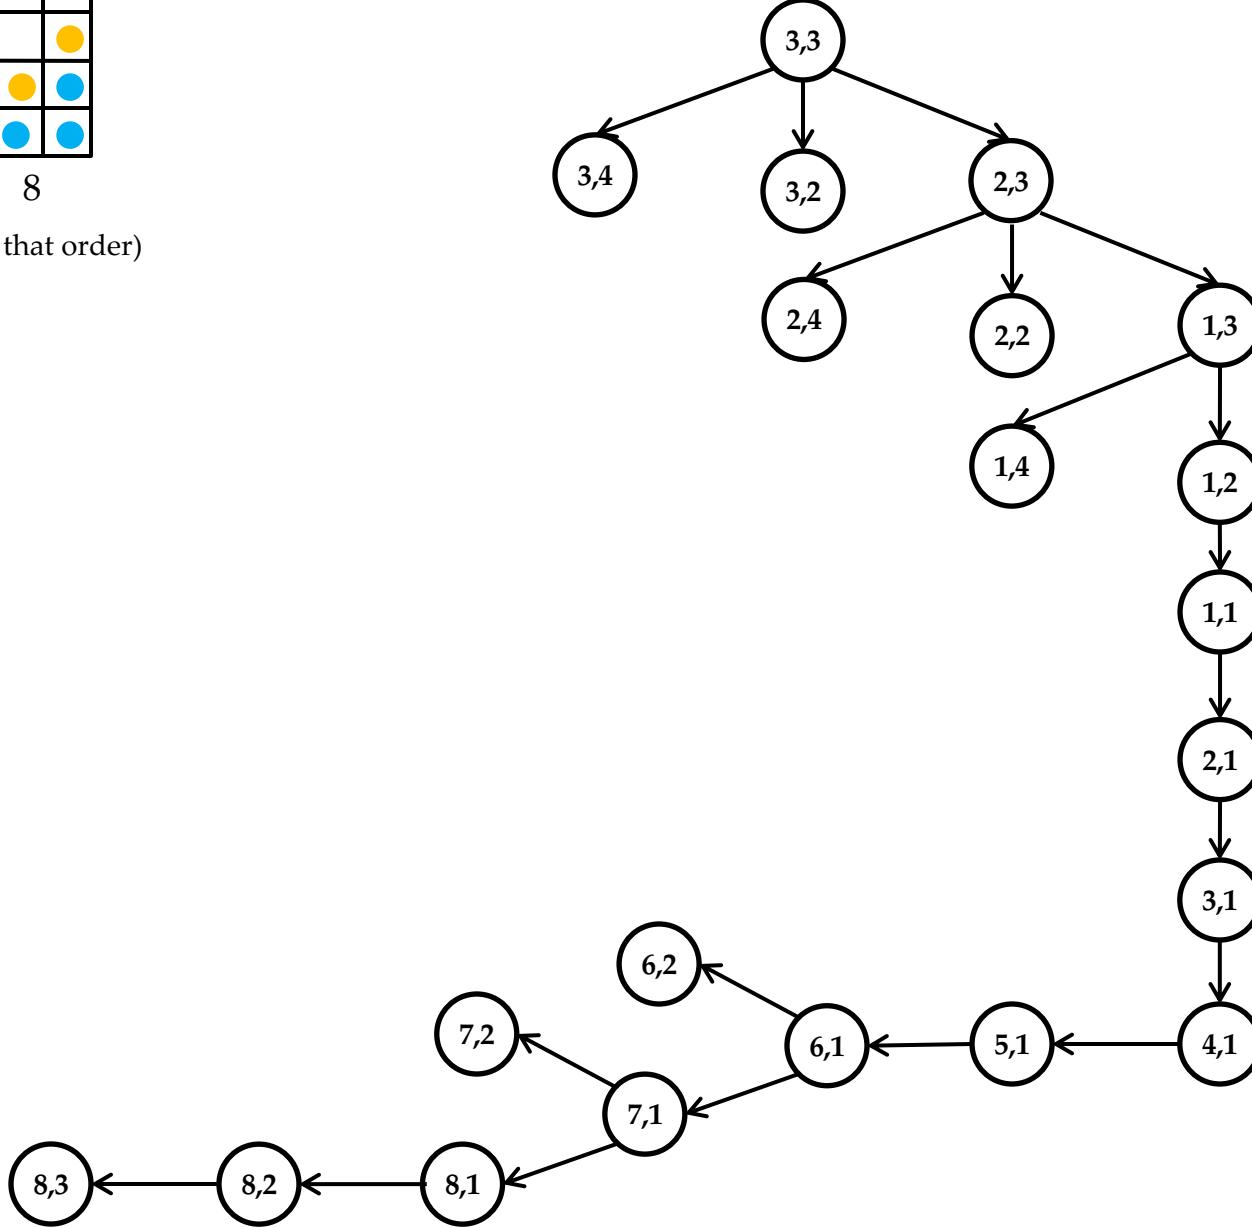
Search Tree: Depth-first Search (no duplicates in fringe)

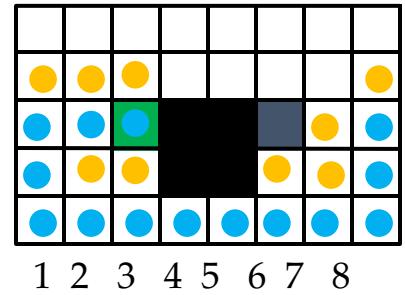




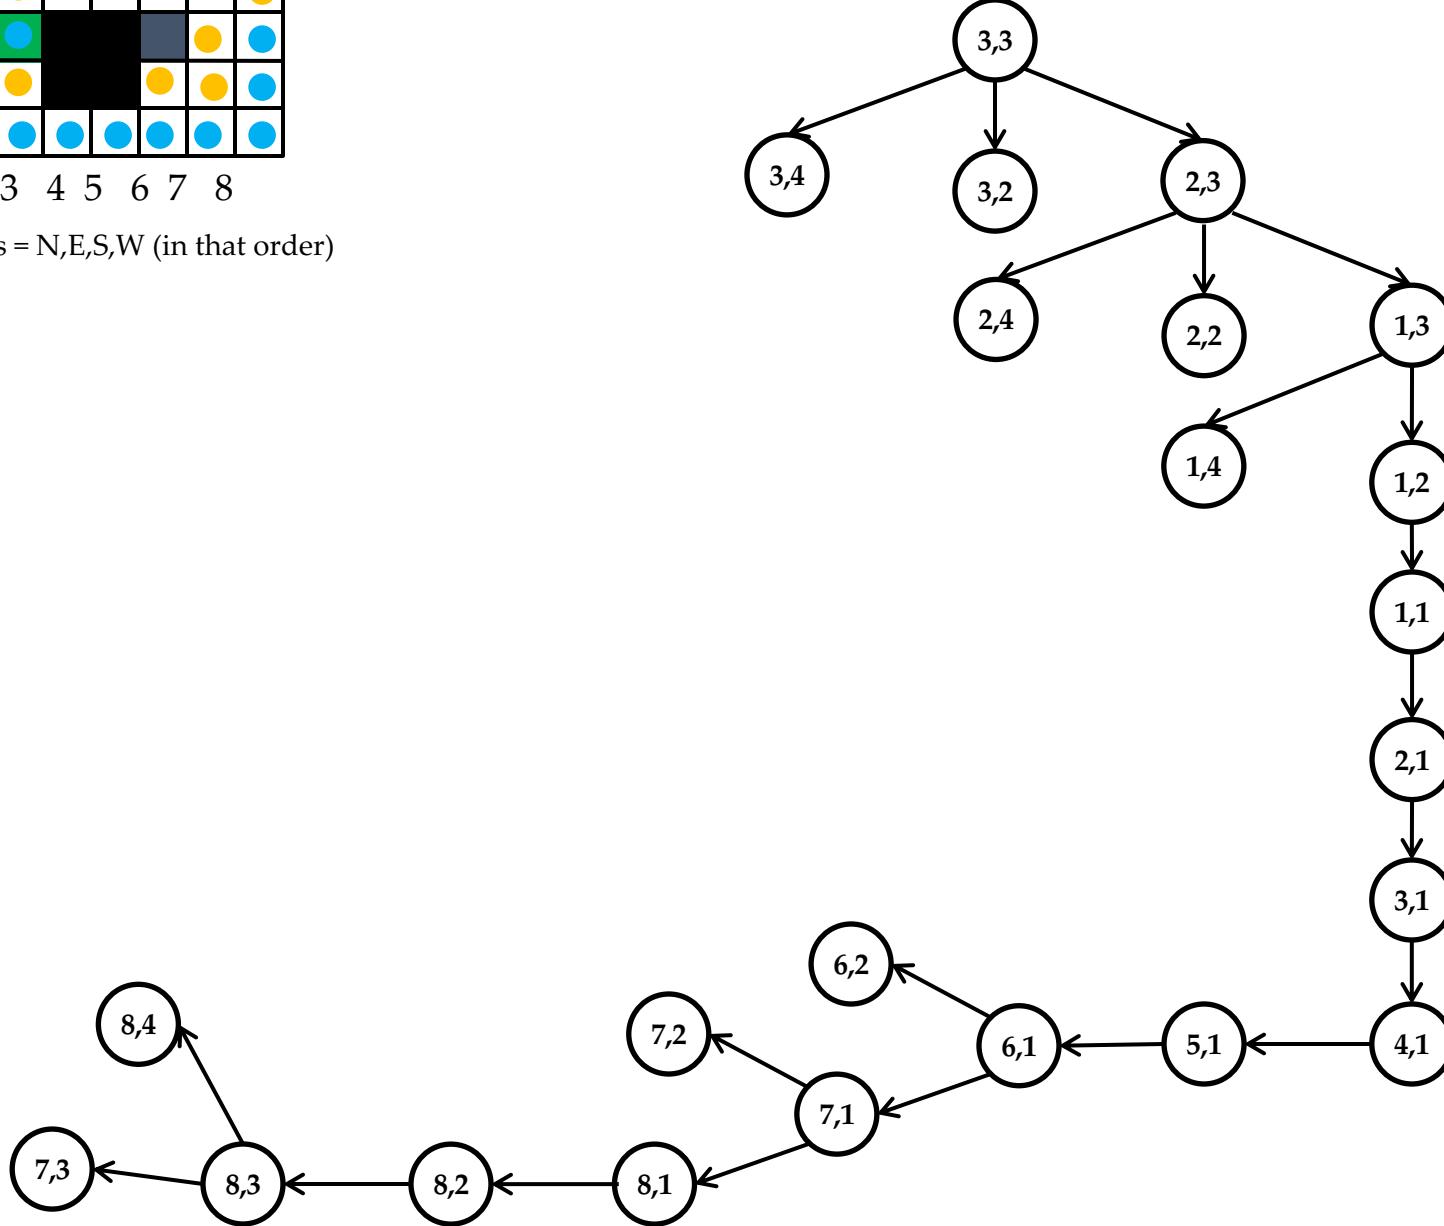
Neighbors = N,E,S,W (in that order)

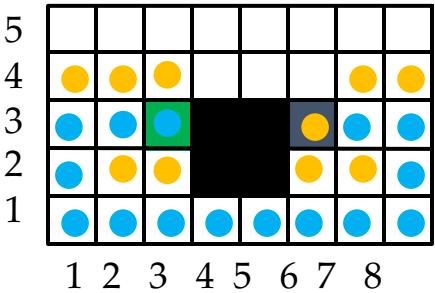
Search Tree: Depth-first Search (no duplicates in fringe)



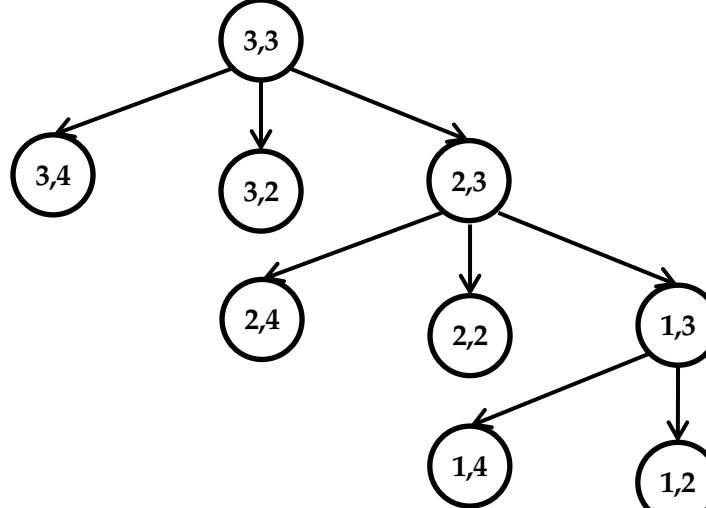


Search Tree: Depth-first Search (no duplicates in fringe)

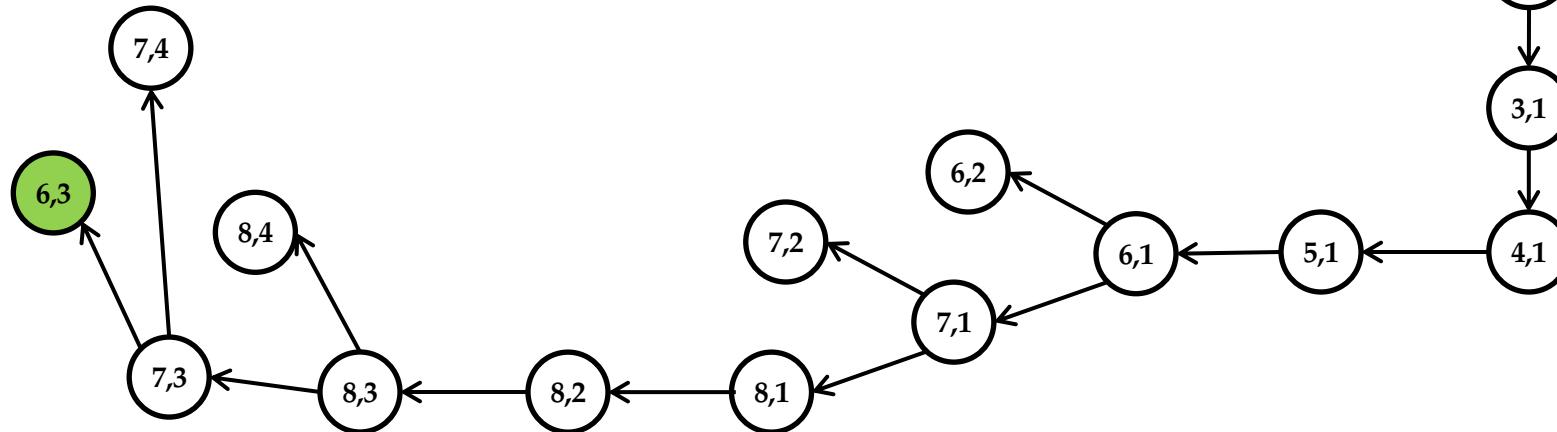


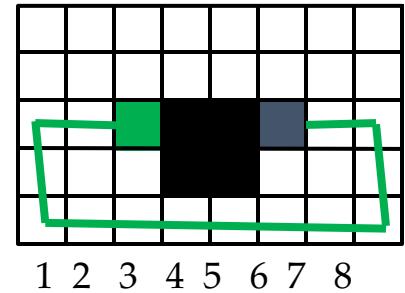


Search Tree: Depth-first Search (no duplicates in fringe)



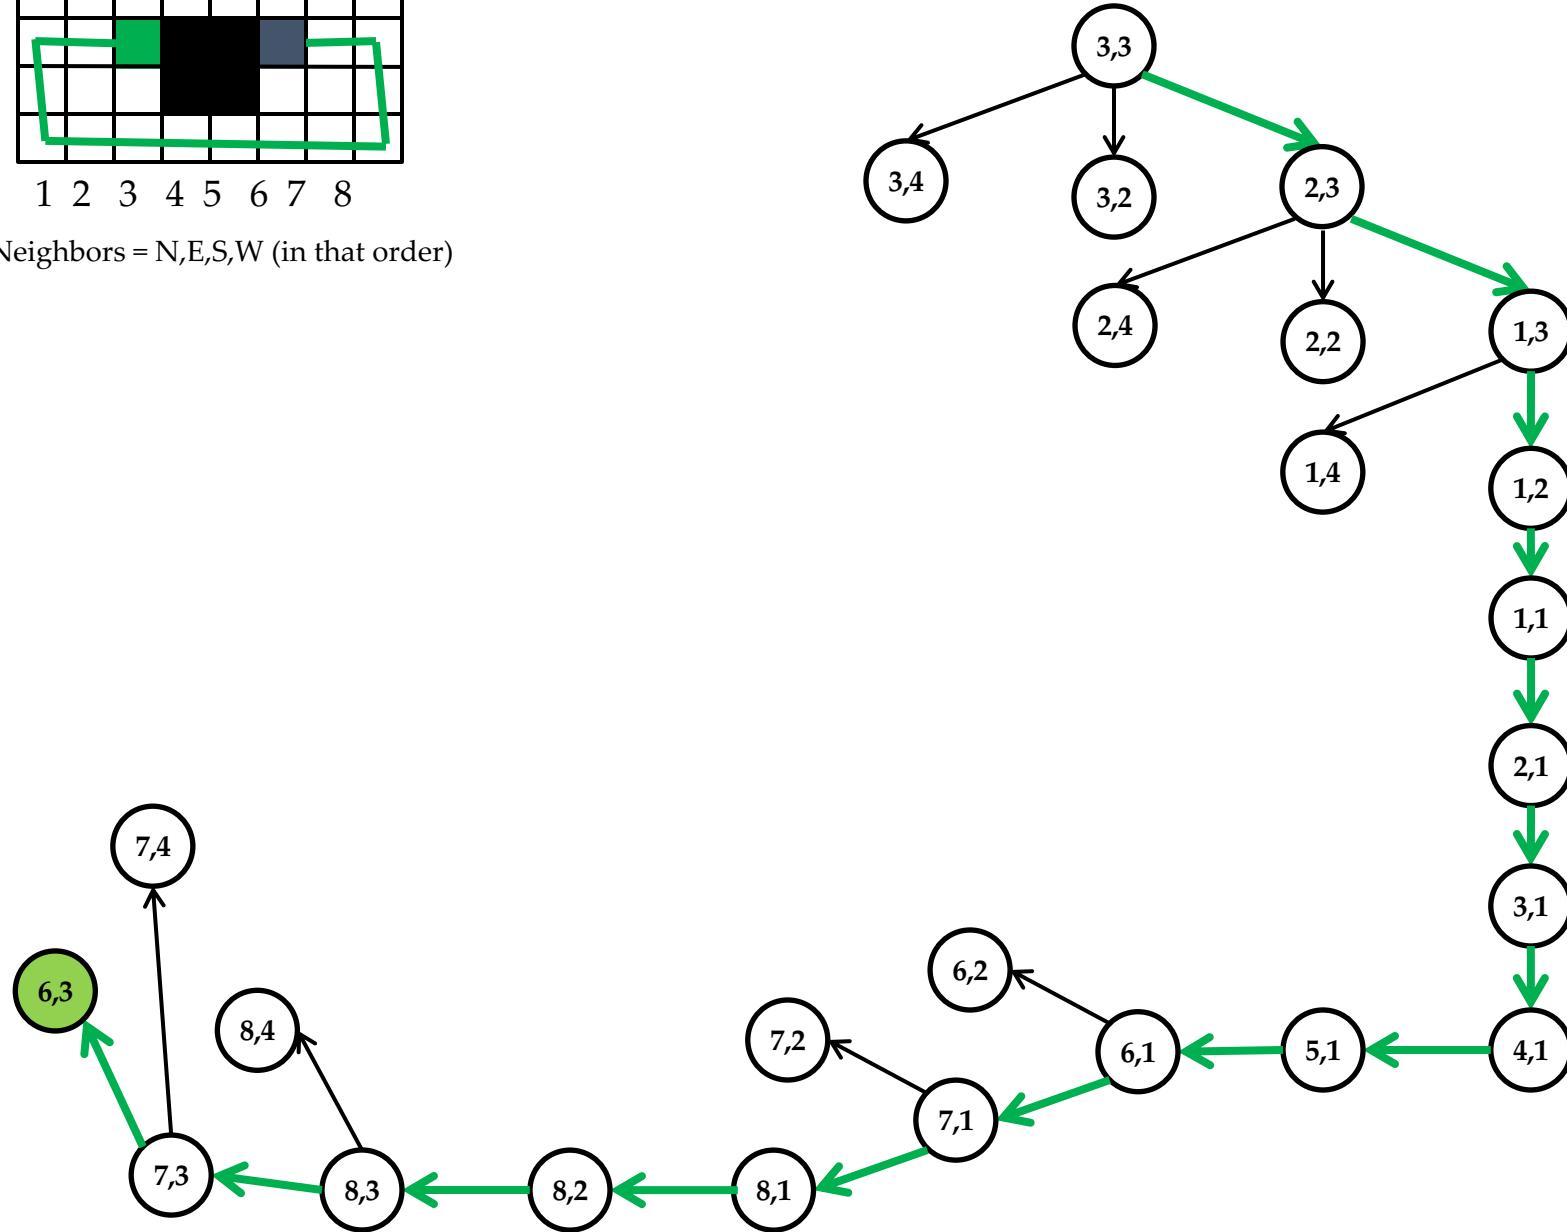
Neighbors = N,E,S,W (in that order)





Neighbors = N,E,S,W (in that order)

Search Tree: Depth-first Search (no duplicates in fringe)



Breadth-First Search

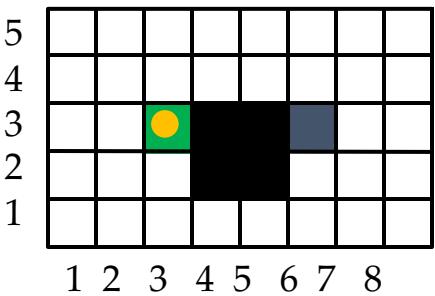


MISSISSIPPI STATE
UNIVERSITY™

Breadth First Search

```
procedure BFS( $G$ ,  $root$ ) is
    let  $Q$  be a queue
    label  $root$  as explored
     $Q.enqueue(root)$ 
    while  $Q$  is not empty do
         $v := Q.dequeue()$ 
        if  $v$  is the goal then
            return  $v$ 
        for all edges from  $v$  to  $w$  in  $G.adjacentEdges(v)$  do
            if  $w$  is not labeled as explored then
                label  $w$  as explored
                 $Q.enqueue(w)$ 
```

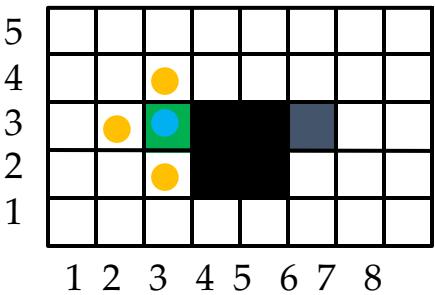




Search Tree: Breadth-first Search (no duplicates in fringe)

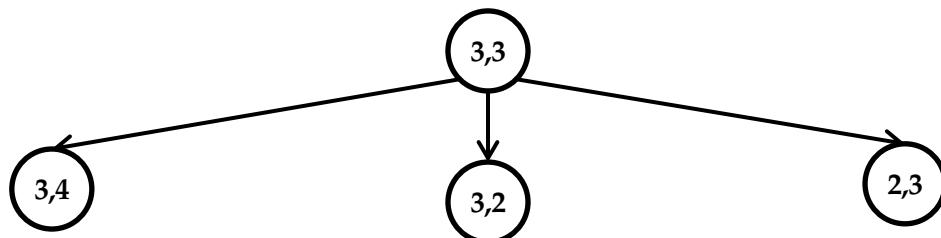


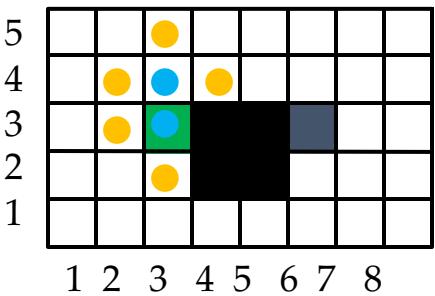
Neighbors = N,E,S,W (in that order)



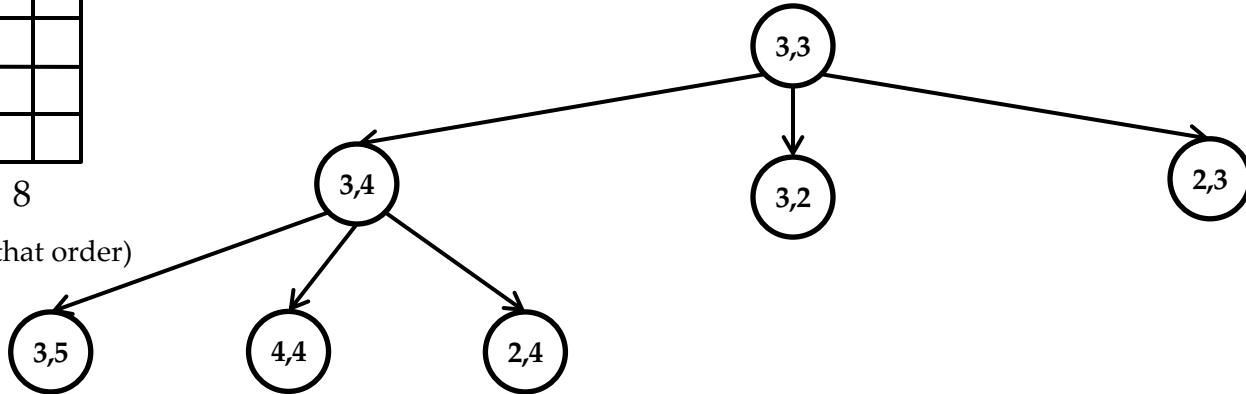
Neighbors = N,E,S,W (in that order)

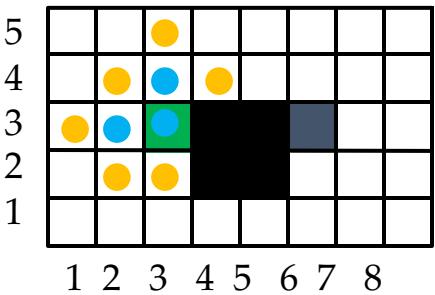
Search Tree: Breadth-first Search (no duplicates in fringe)





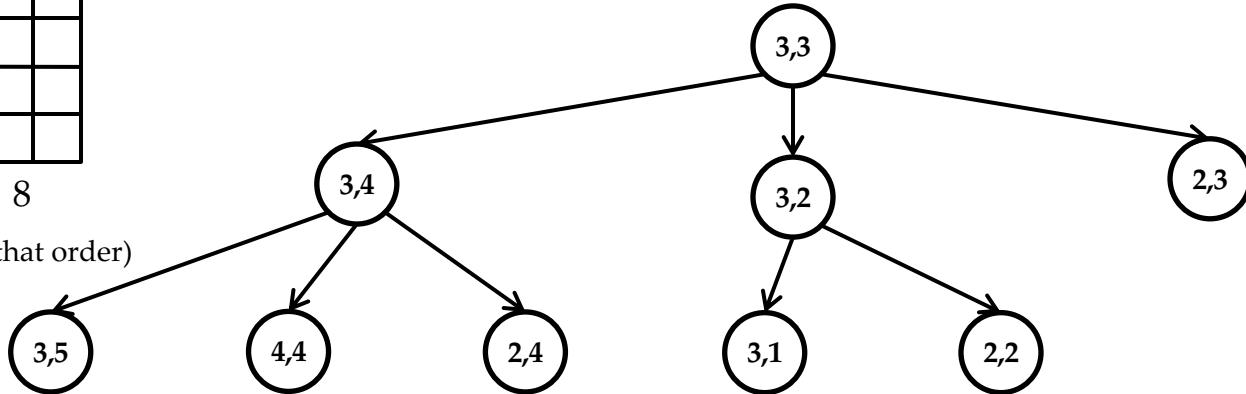
Search Tree: Breadth-first Search (no duplicates in fringe)

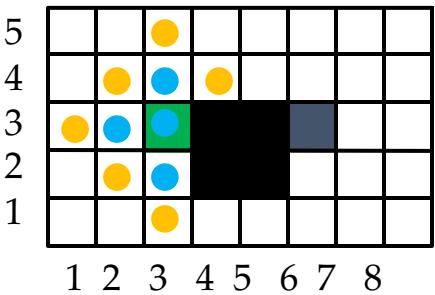




Search Tree: Breadth-first Search (no duplicates in fringe)

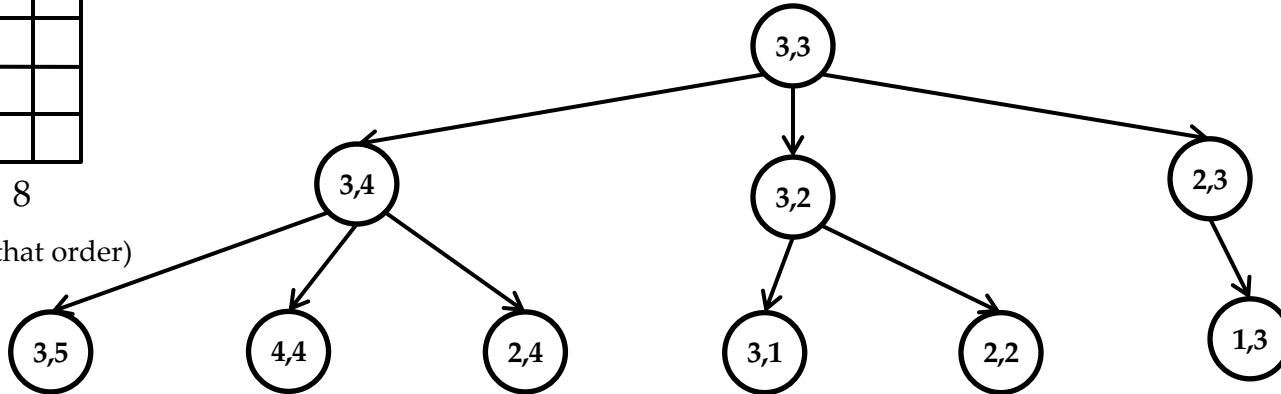
Neighbors = N,E,S,W (in that order)

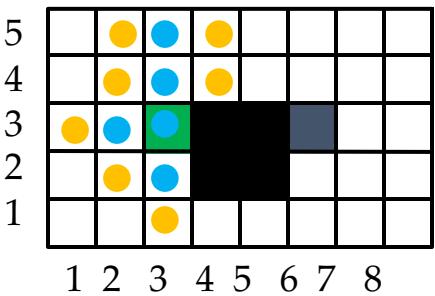




Search Tree: Breadth-first Search (no duplicates in fringe)

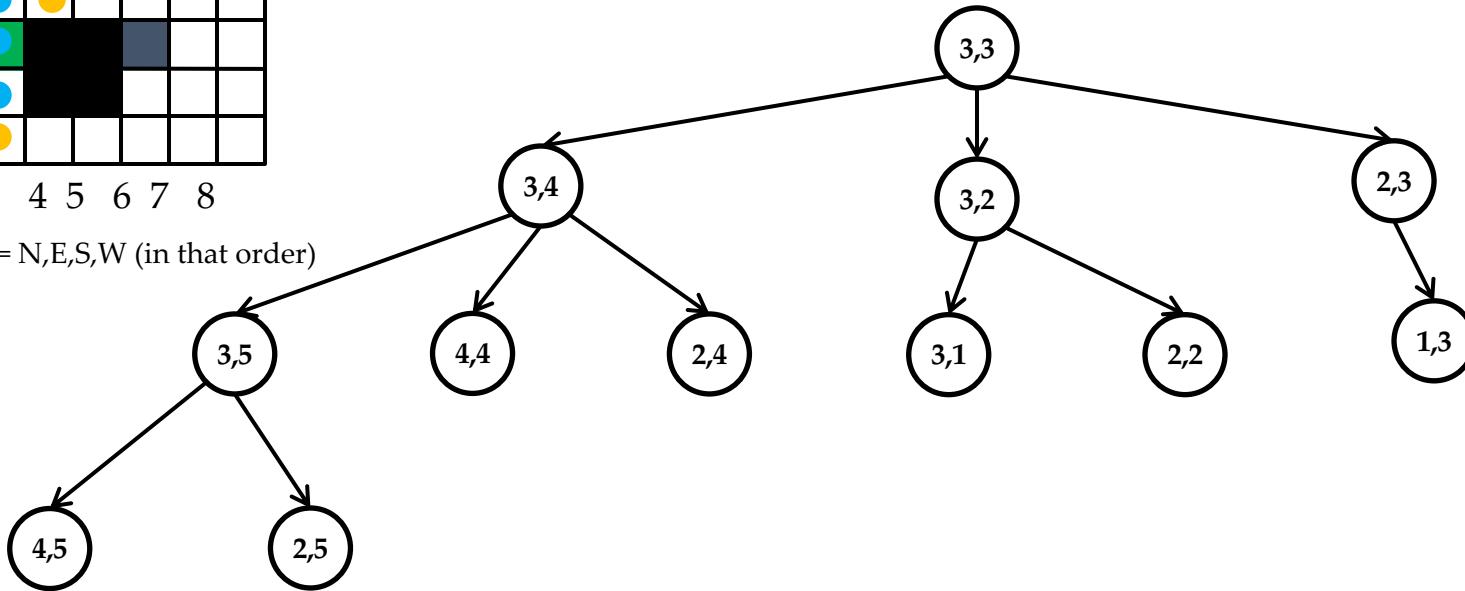
Neighbors = N,E,S,W (in that order)

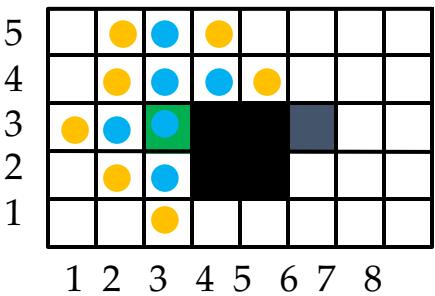




Search Tree: Breadth-first Search (no duplicates in fringe)

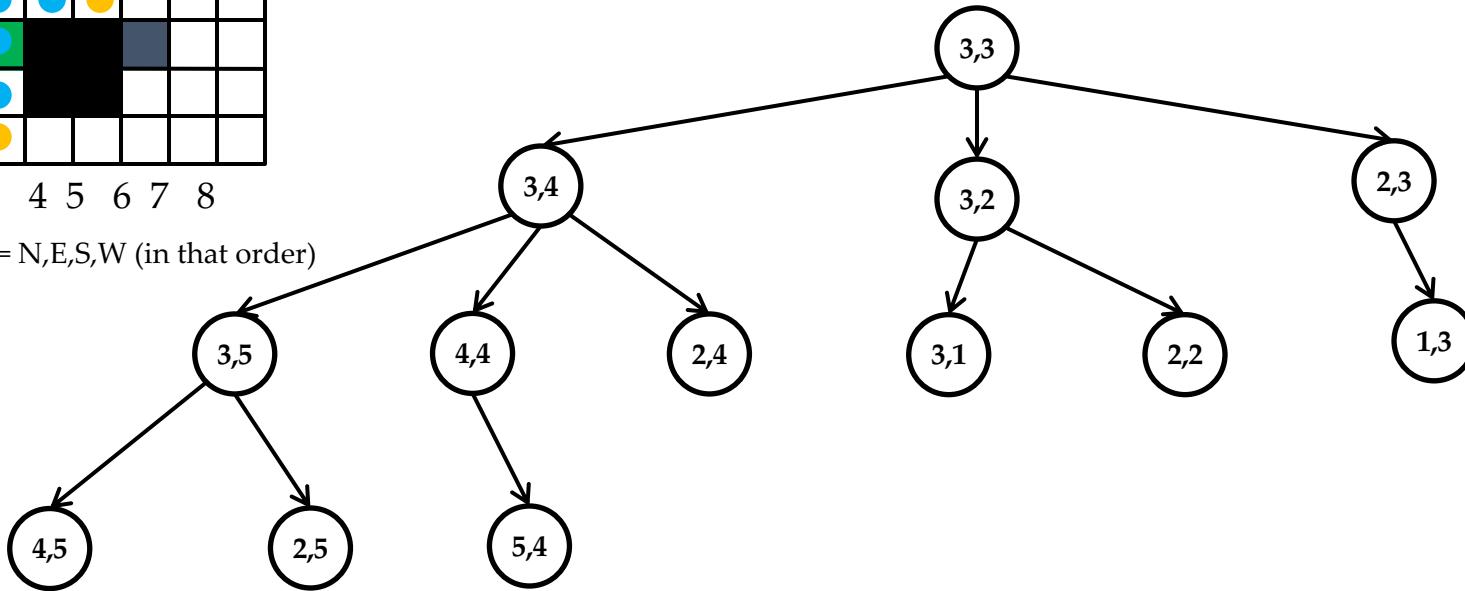
Neighbors = N,E,S,W (in that order)

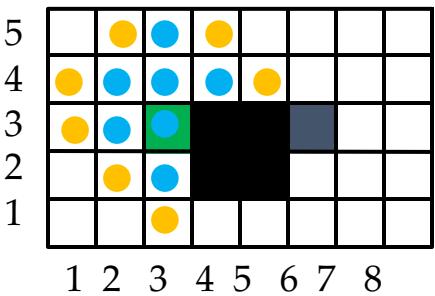




Search Tree: Breadth-first Search (no duplicates in fringe)

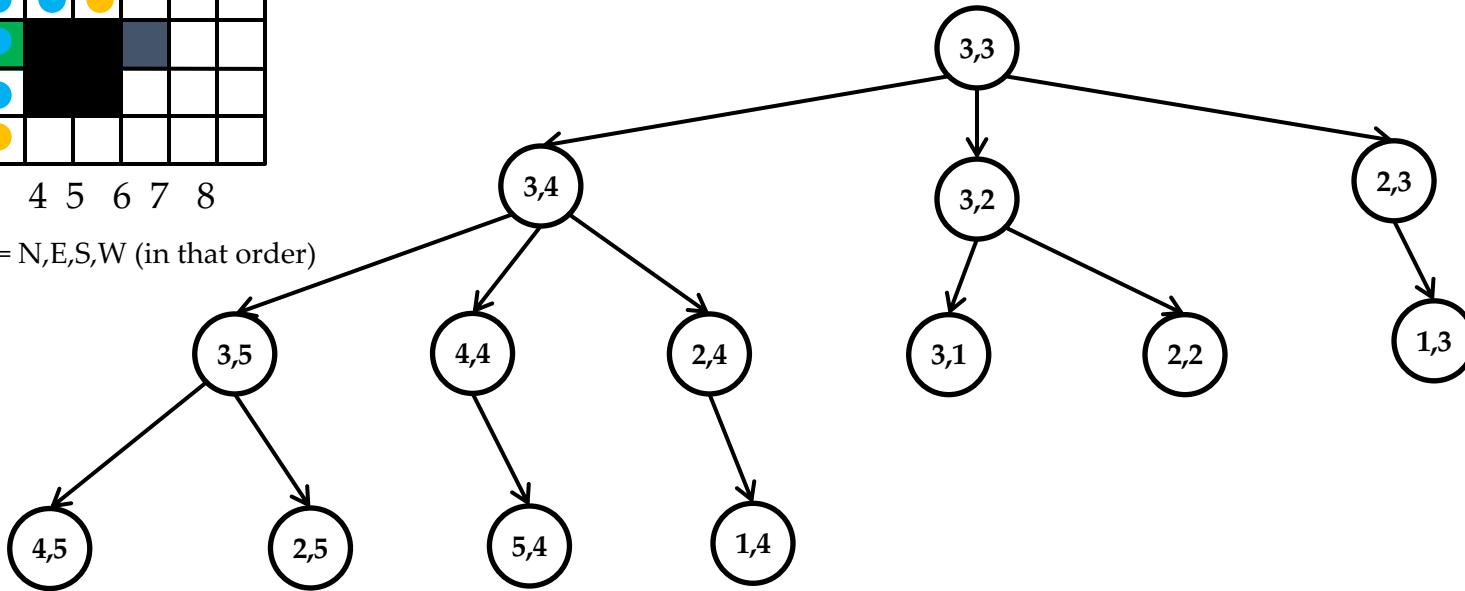
Neighbors = N,E,S,W (in that order)

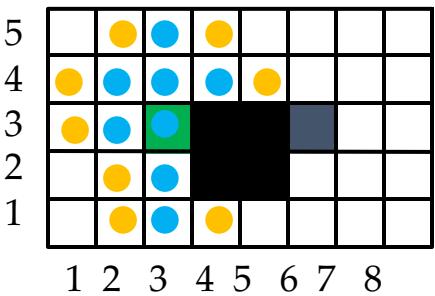




Search Tree: Breadth-first Search (no duplicates in fringe)

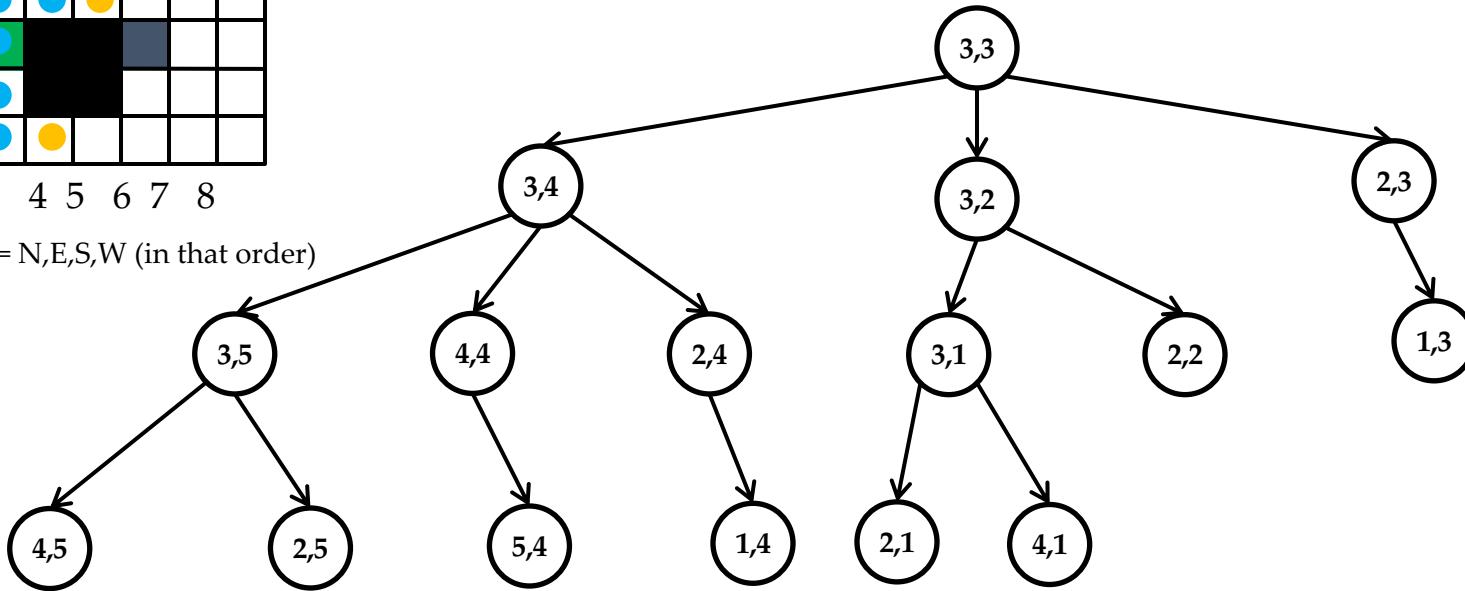
Neighbors = N,E,S,W (in that order)

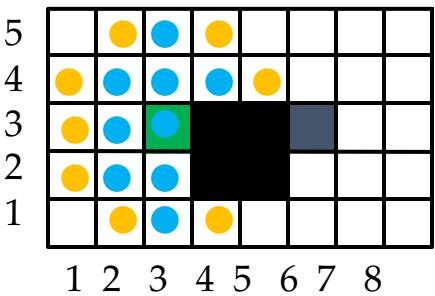




Search Tree: Breadth-first Search (no duplicates in fringe)

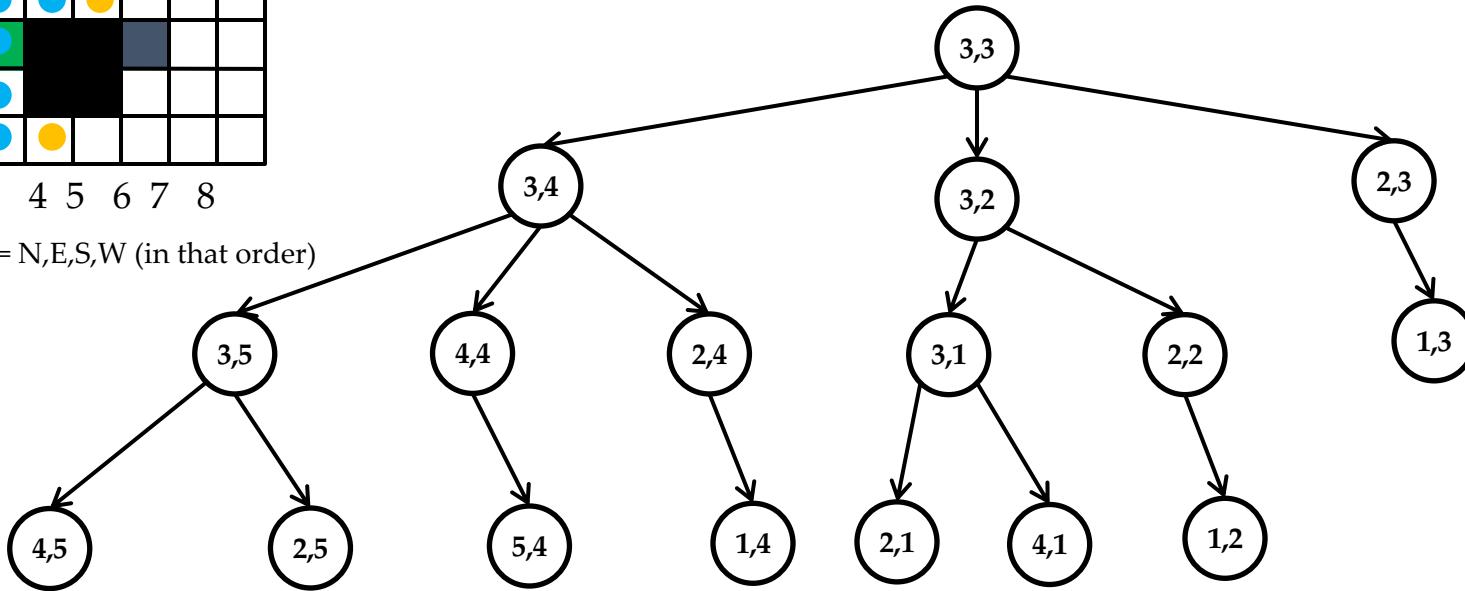
Neighbors = N,E,S,W (in that order)

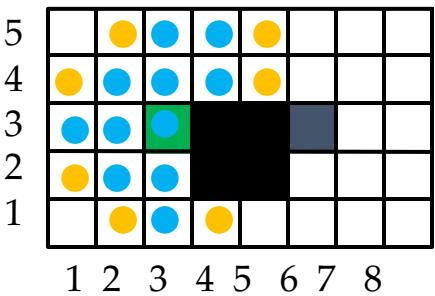




Search Tree: Breadth-first Search (no duplicates in fringe)

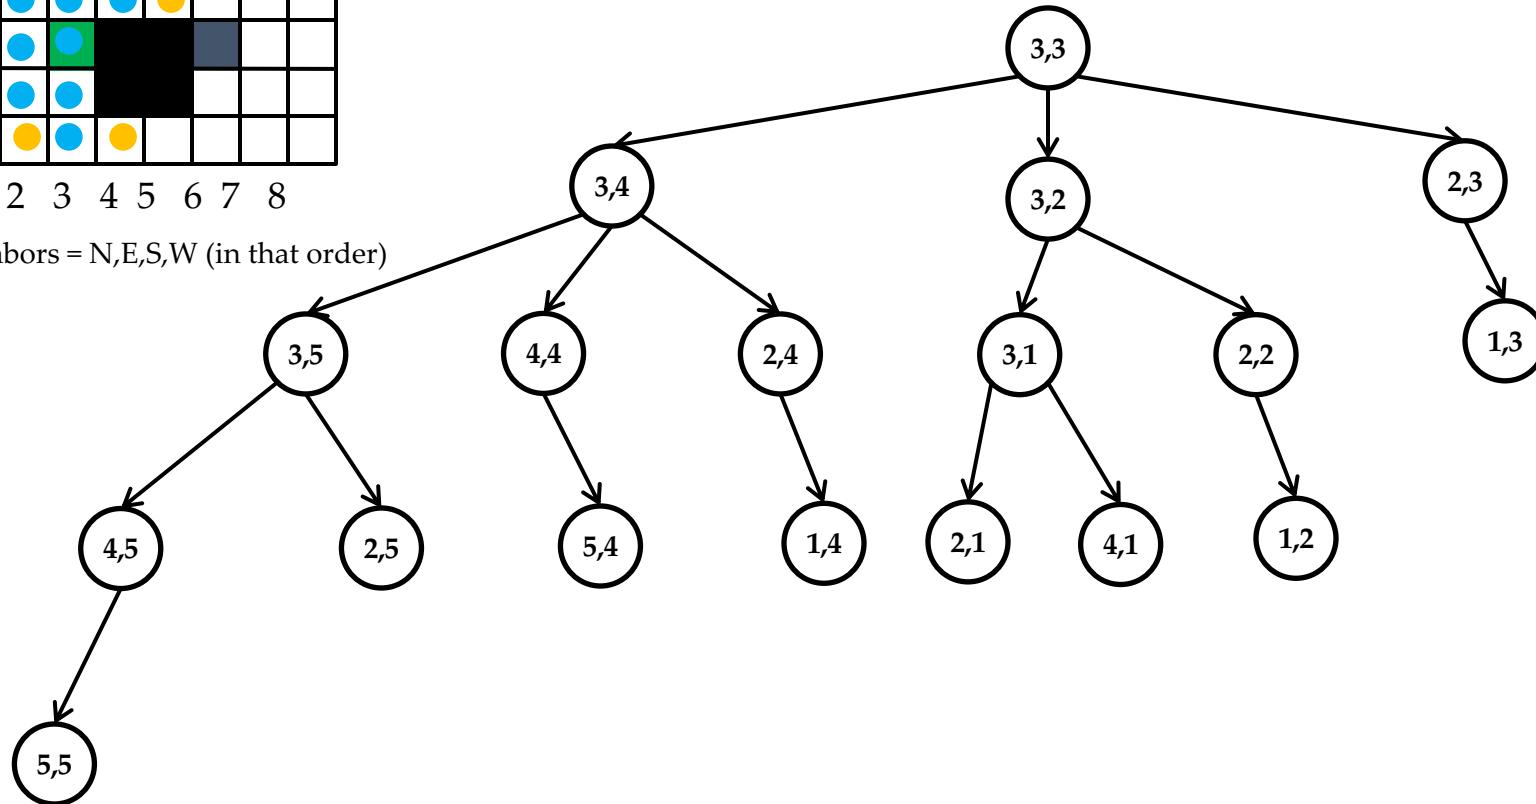
Neighbors = N,E,S,W (in that order)

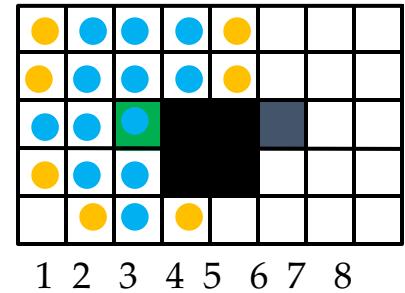




Search Tree: Breadth-first Search (no duplicates in fringe)

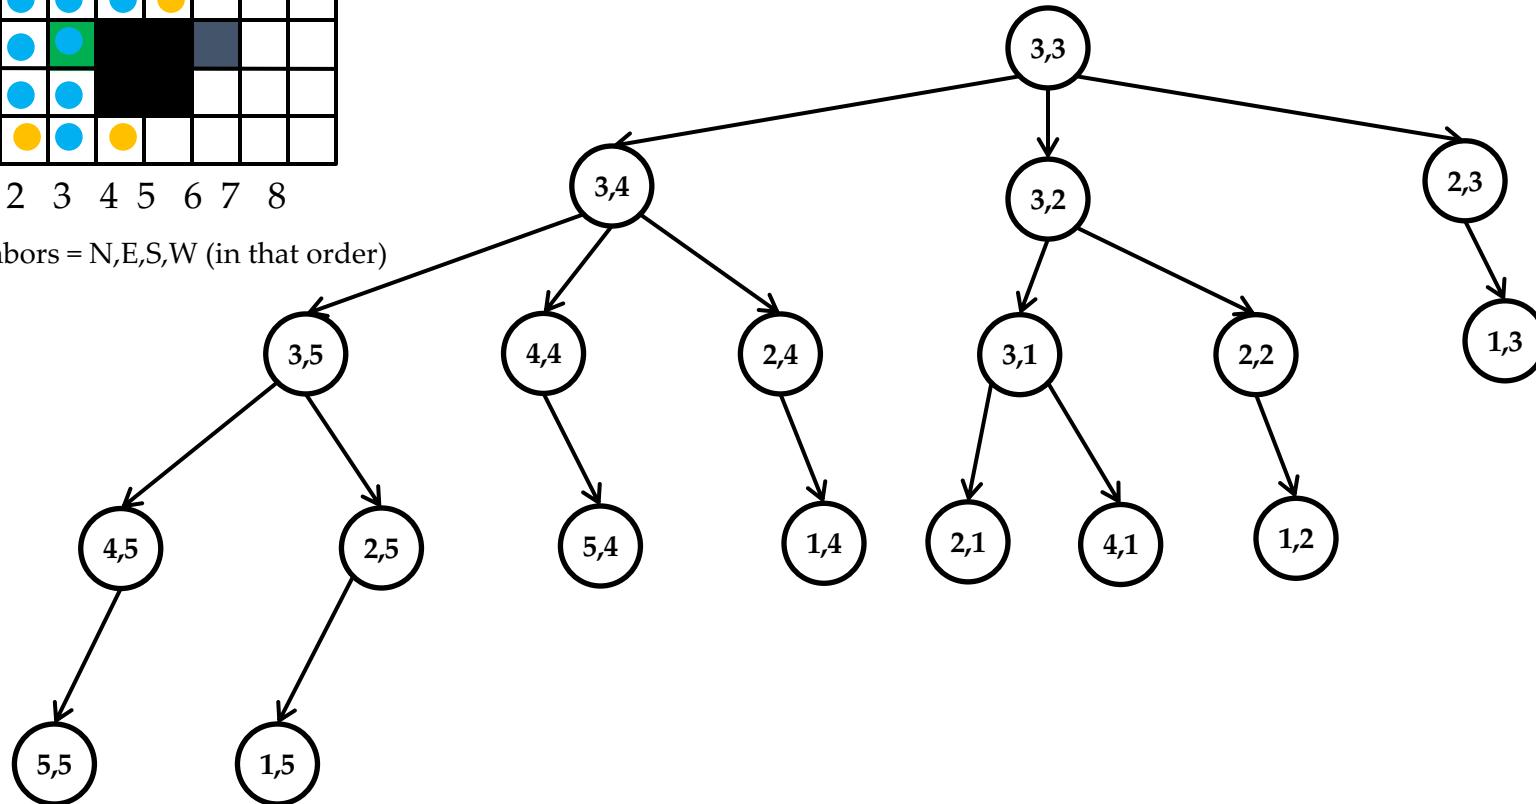
Neighbors = N,E,S,W (in that order)

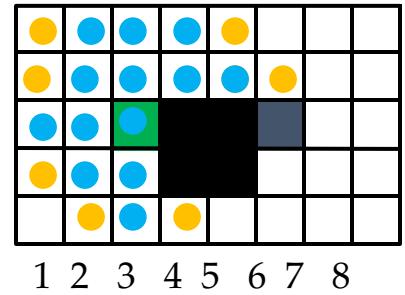




Search Tree: Breadth-first Search (no duplicates in fringe)

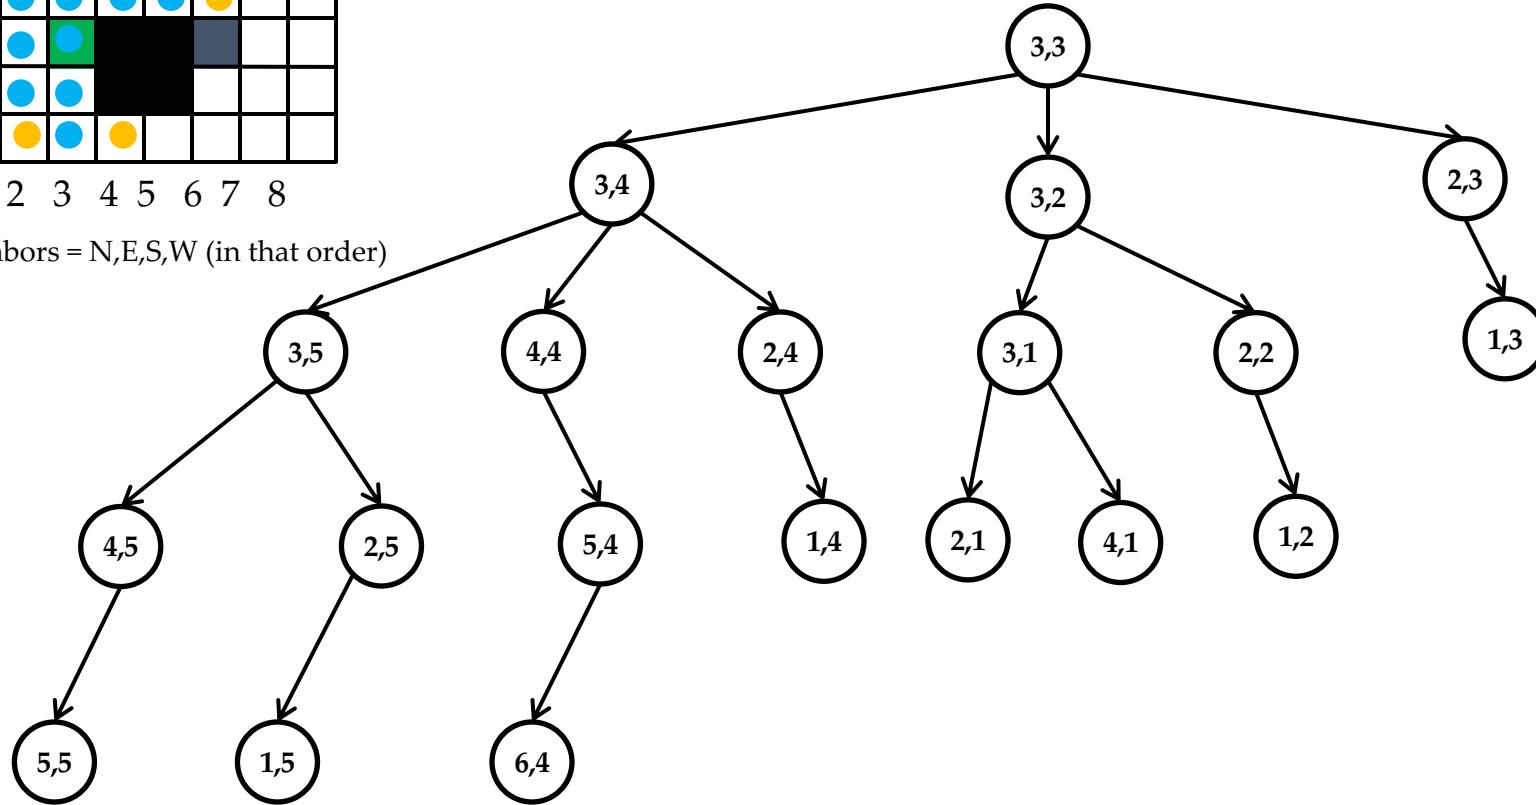
Neighbors = N,E,S,W (in that order)

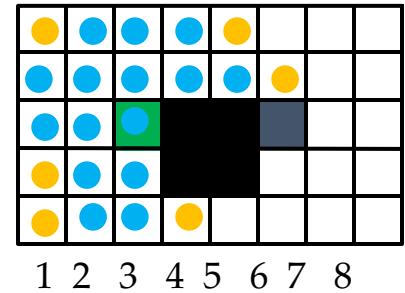




Search Tree: Breadth-first Search (no duplicates in fringe)

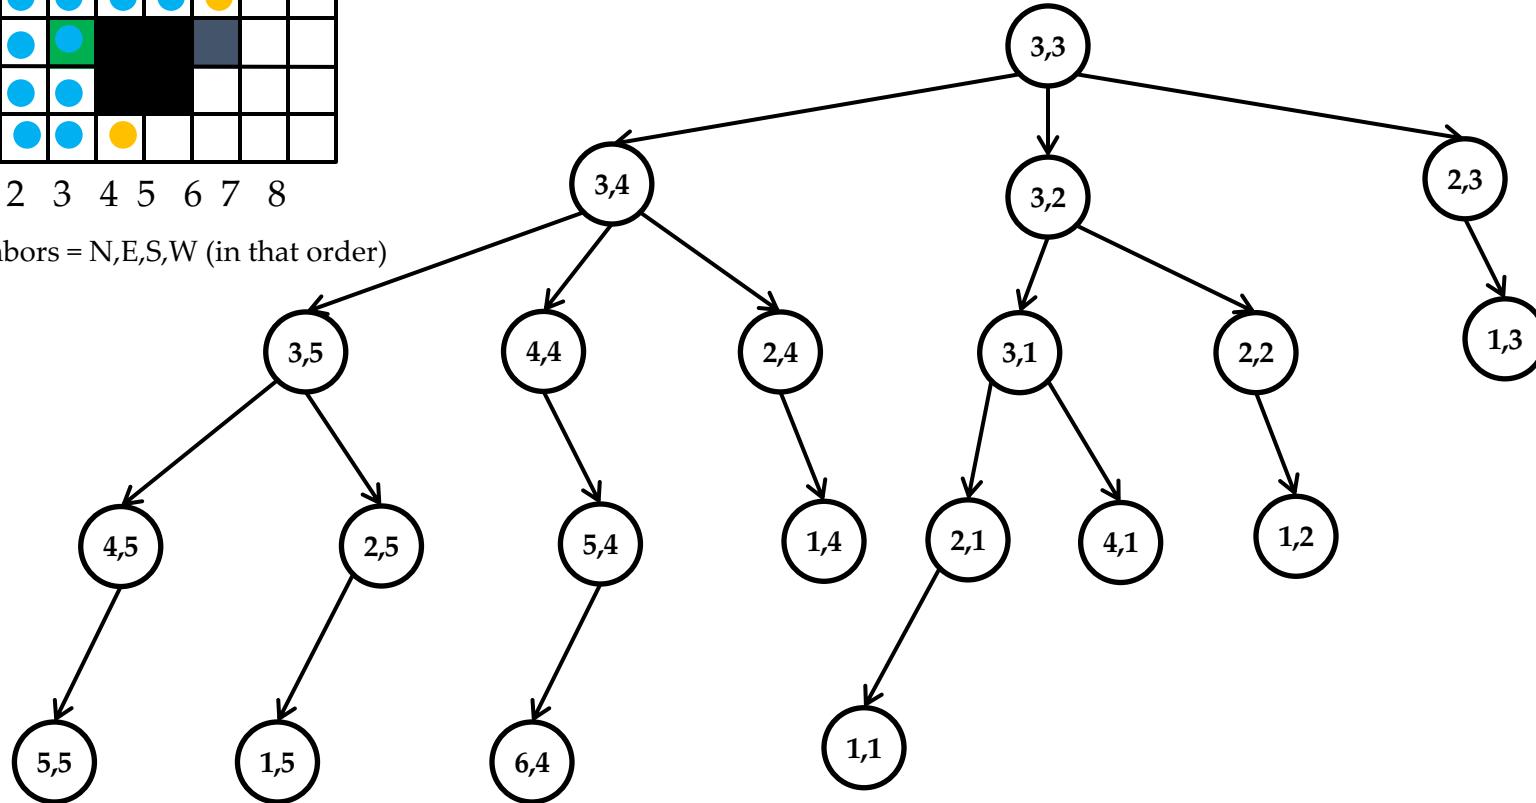
Neighbors = N,E,S,W (in that order)

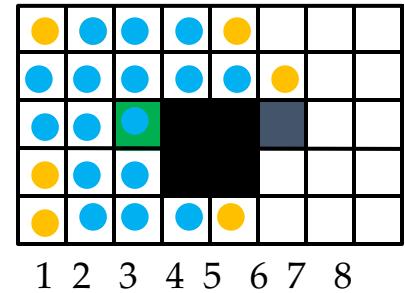




Search Tree: Breadth-first Search (no duplicates in fringe)

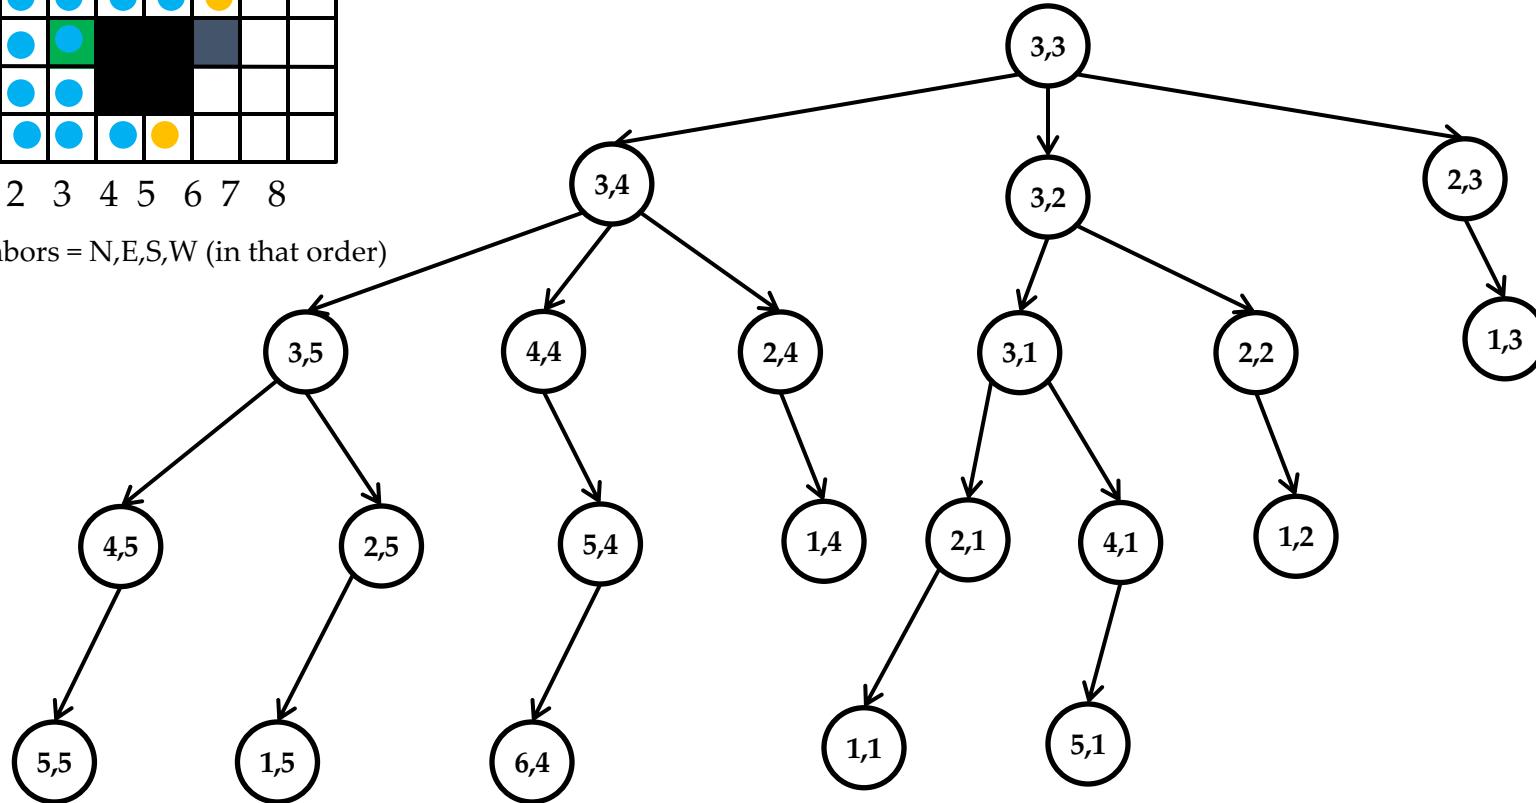
Neighbors = N,E,S,W (in that order)

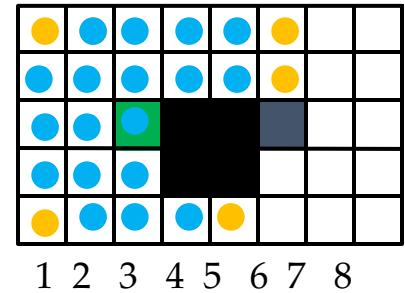




Search Tree: Breadth-first Search (no duplicates in fringe)

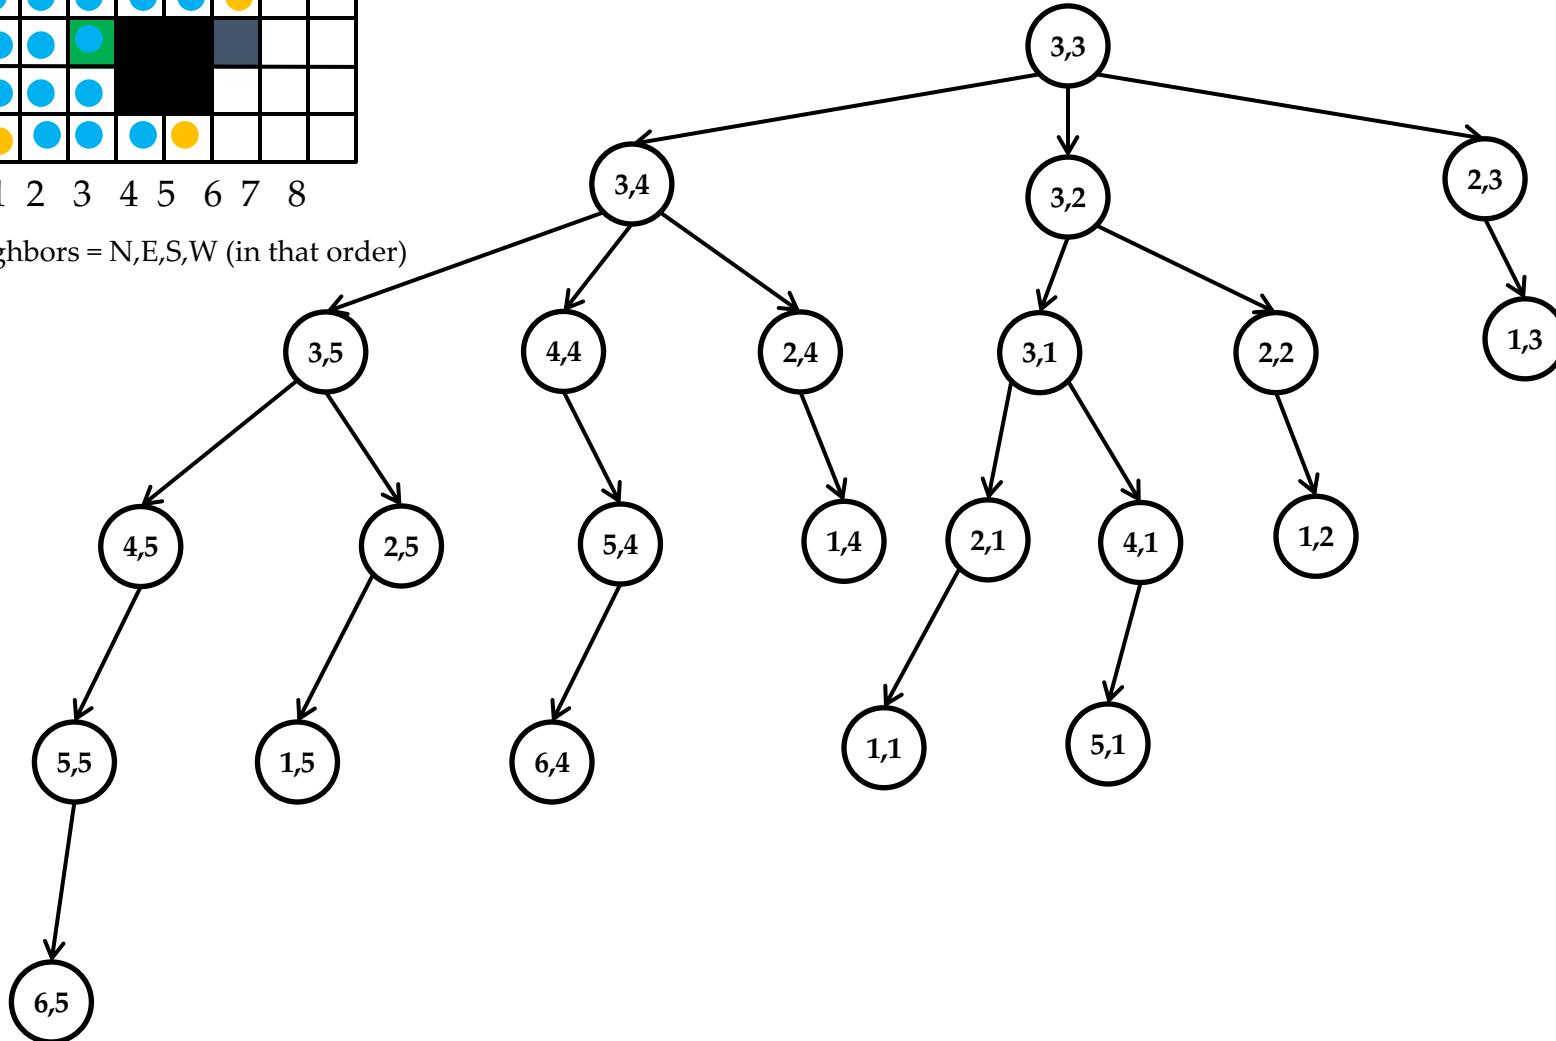
Neighbors = N,E,S,W (in that order)

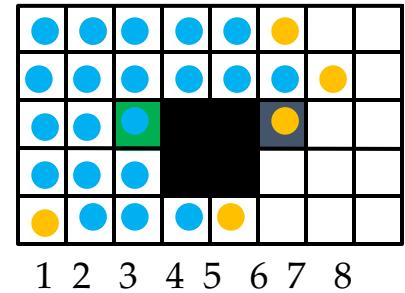




Search Tree: Breadth-first Search (no duplicates in fringe)

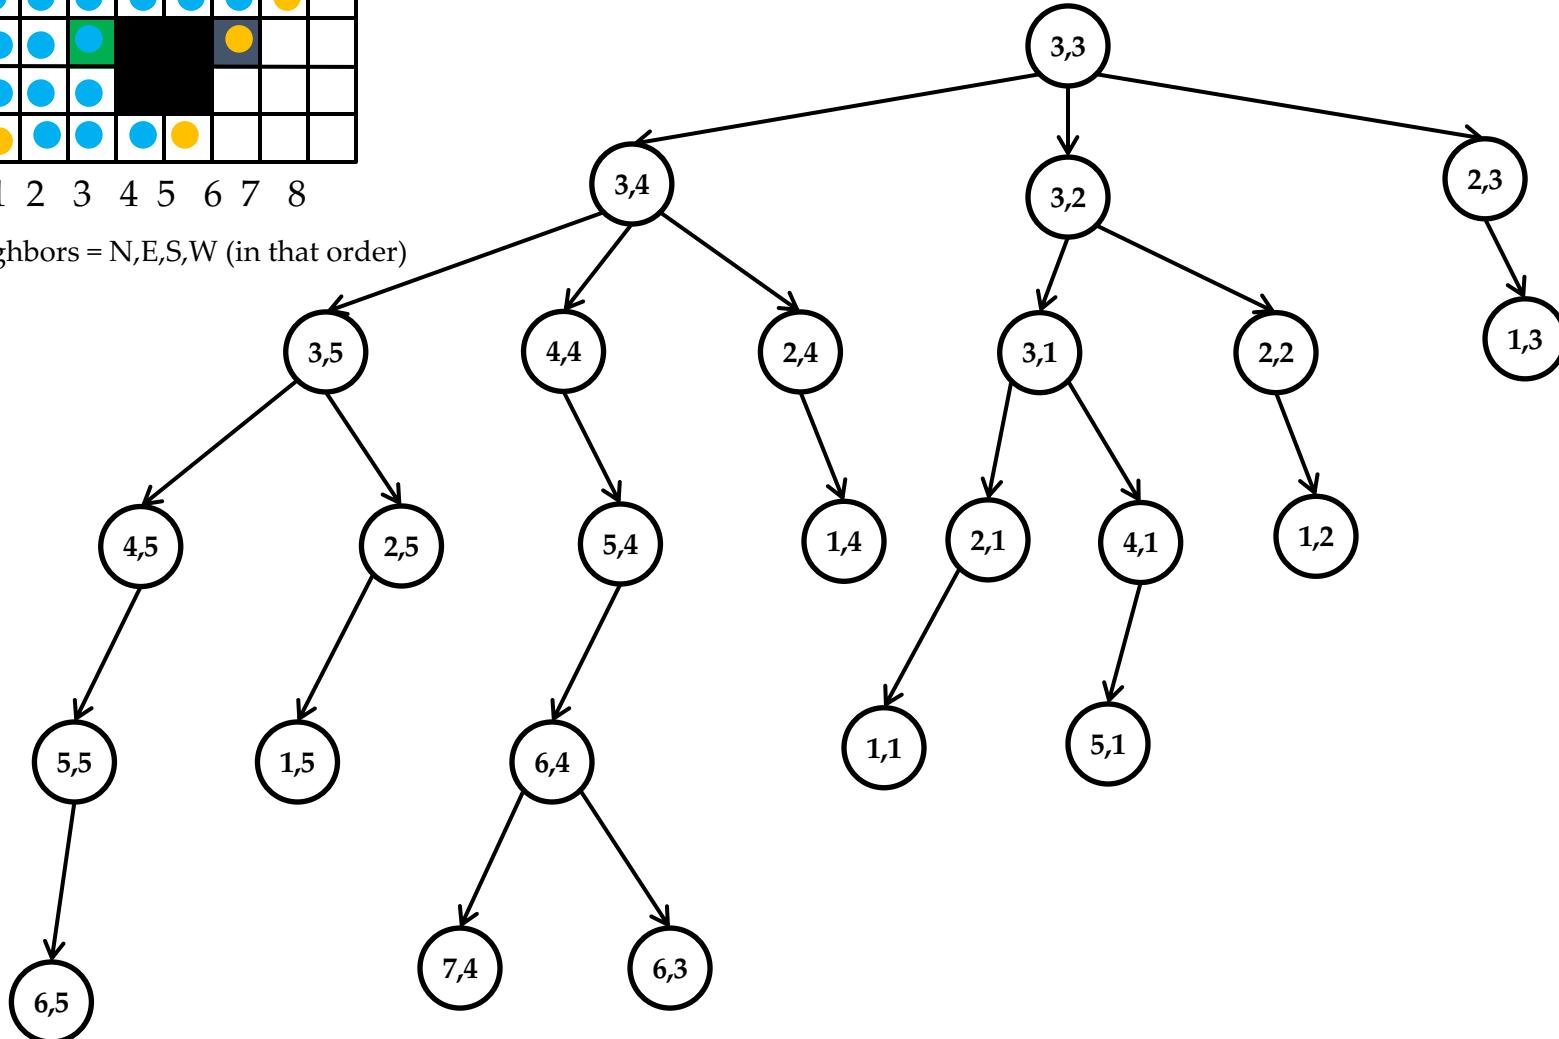
Neighbors = N,E,S,W (in that order)

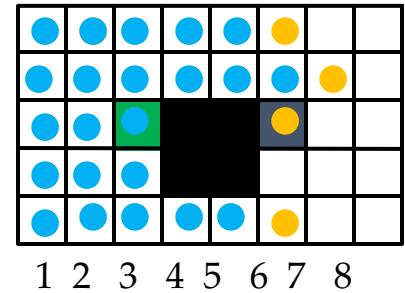




Search Tree: Breadth-first Search (no duplicates in fringe)

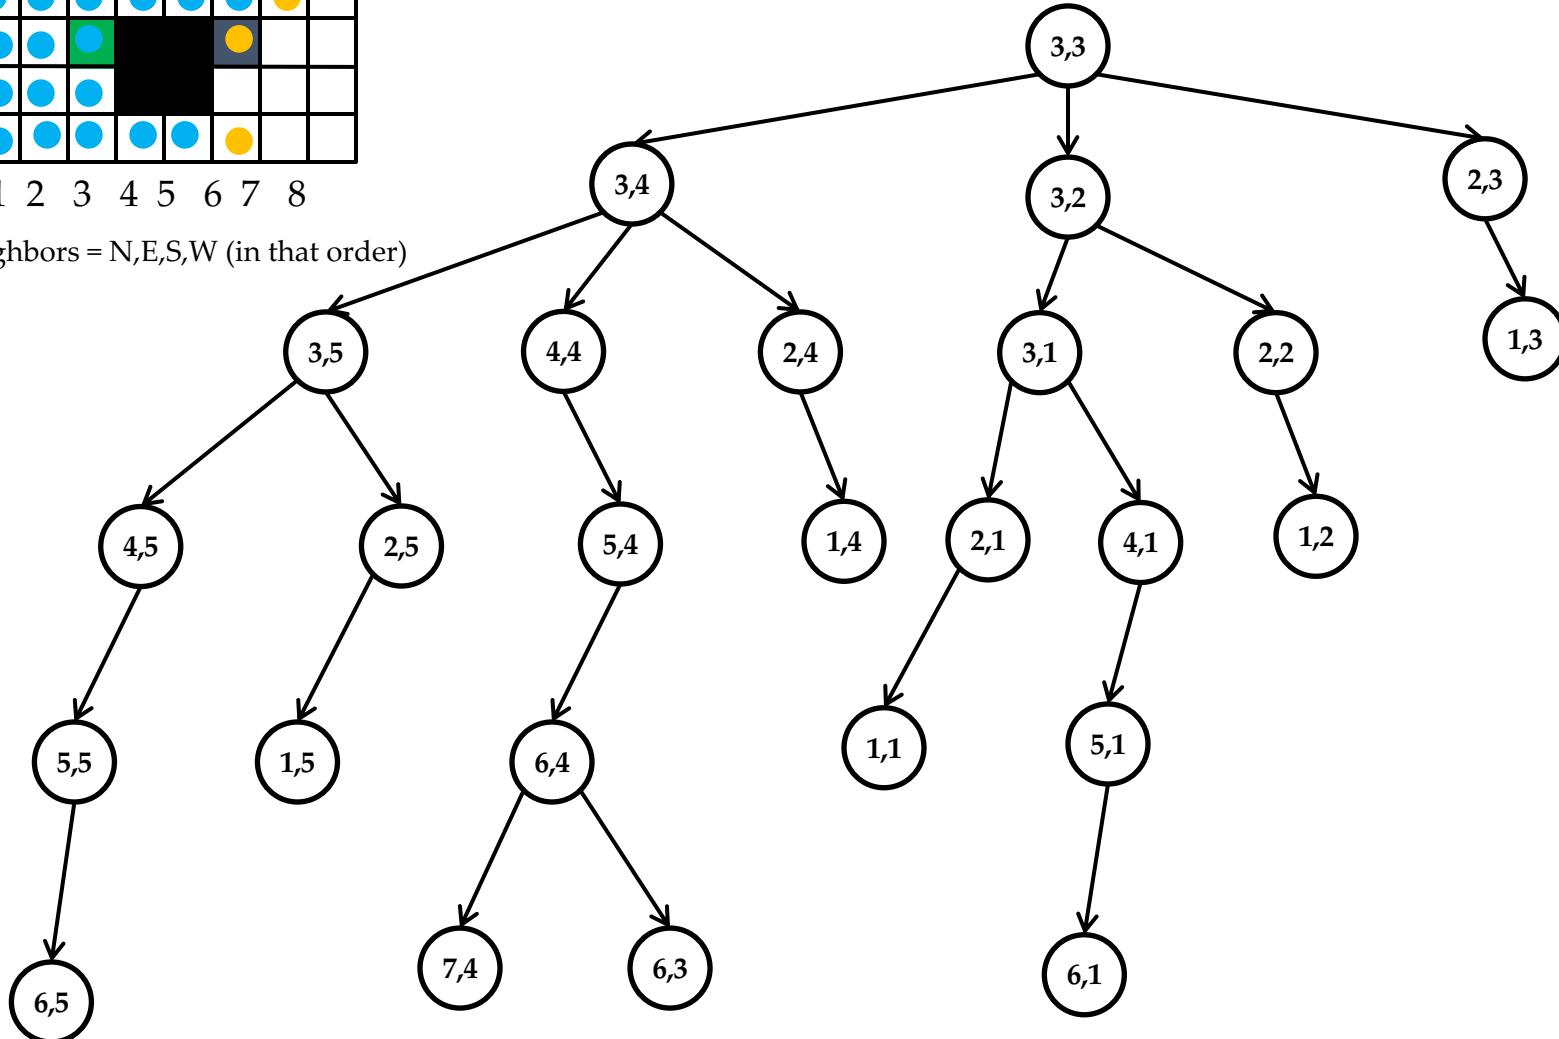
Neighbors = N,E,S,W (in that order)

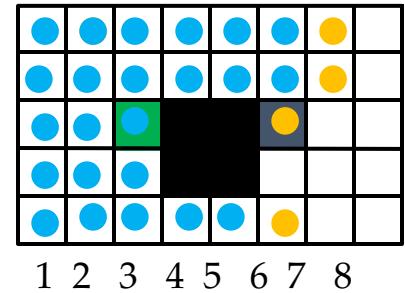




Search Tree: Breadth-first Search (no duplicates in fringe)

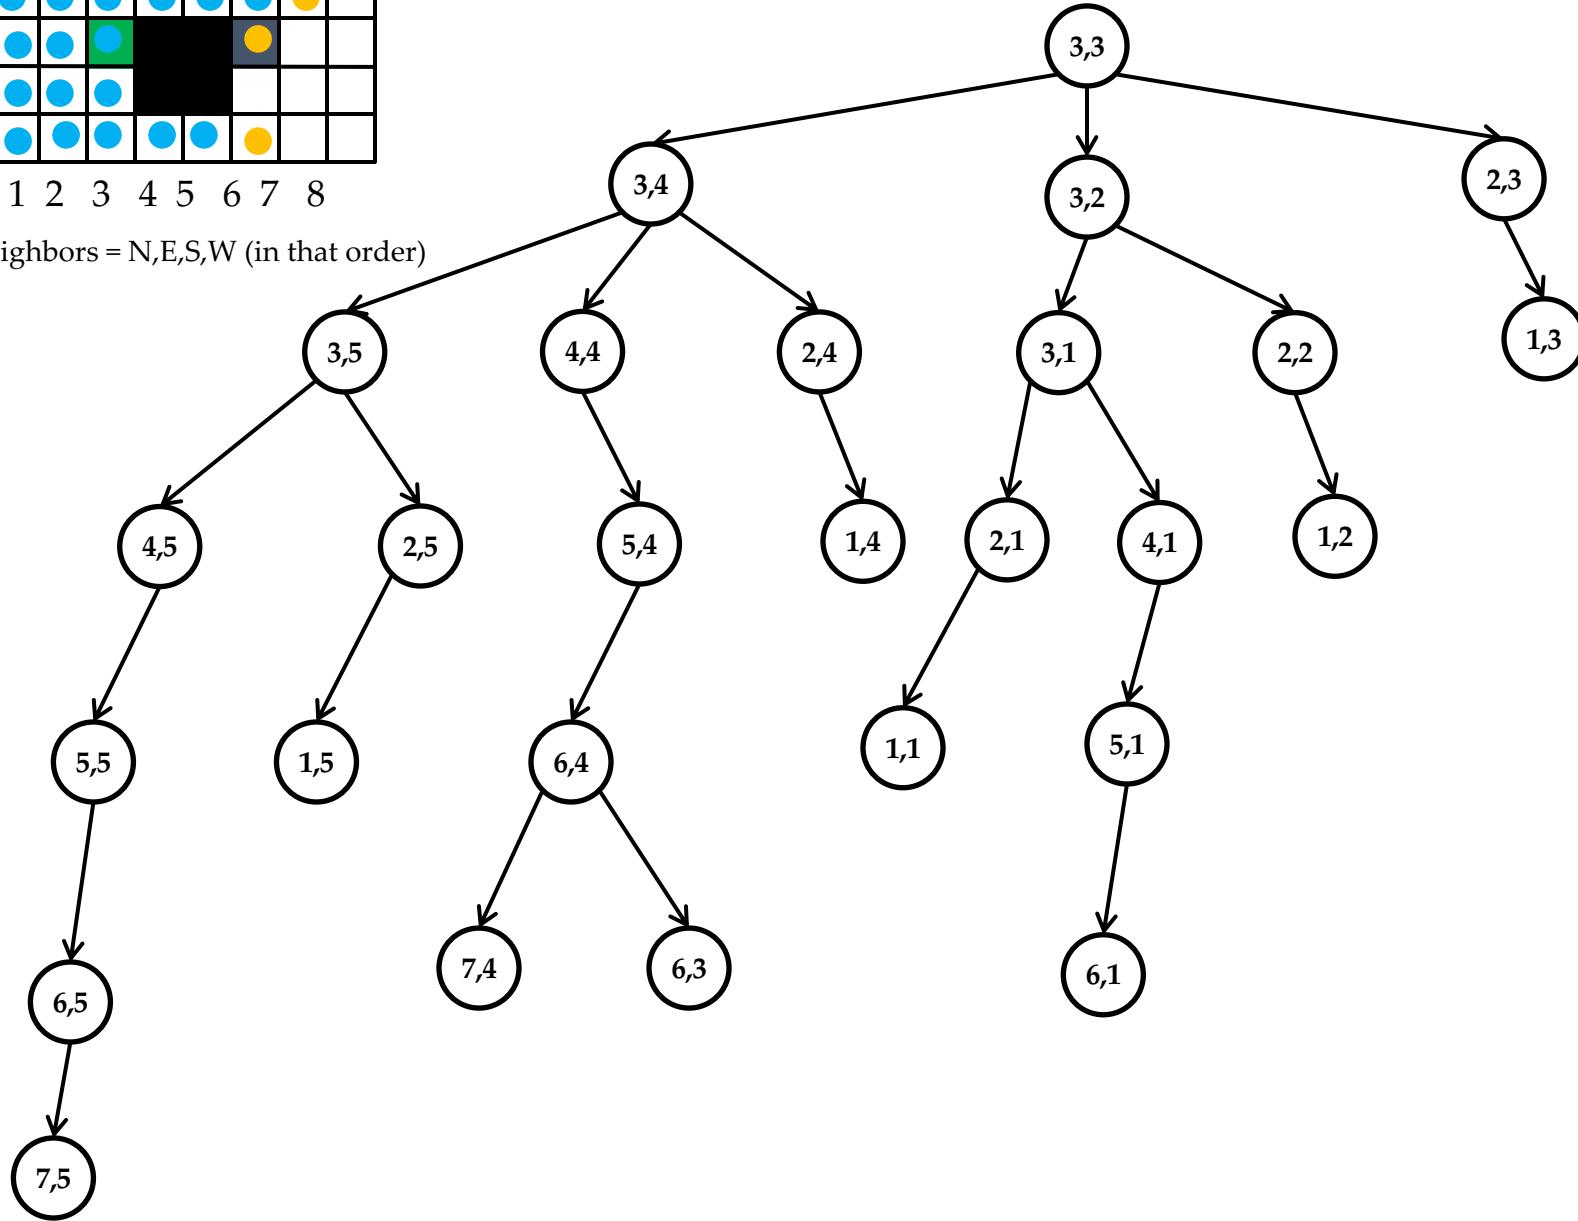
Neighbors = N,E,S,W (in that order)

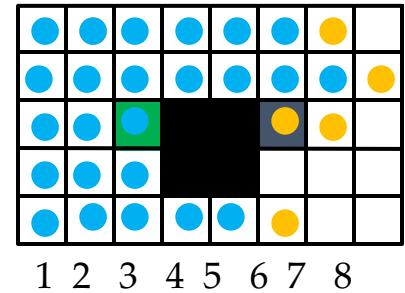




Search Tree: Breadth-first Search (no duplicates in fringe)

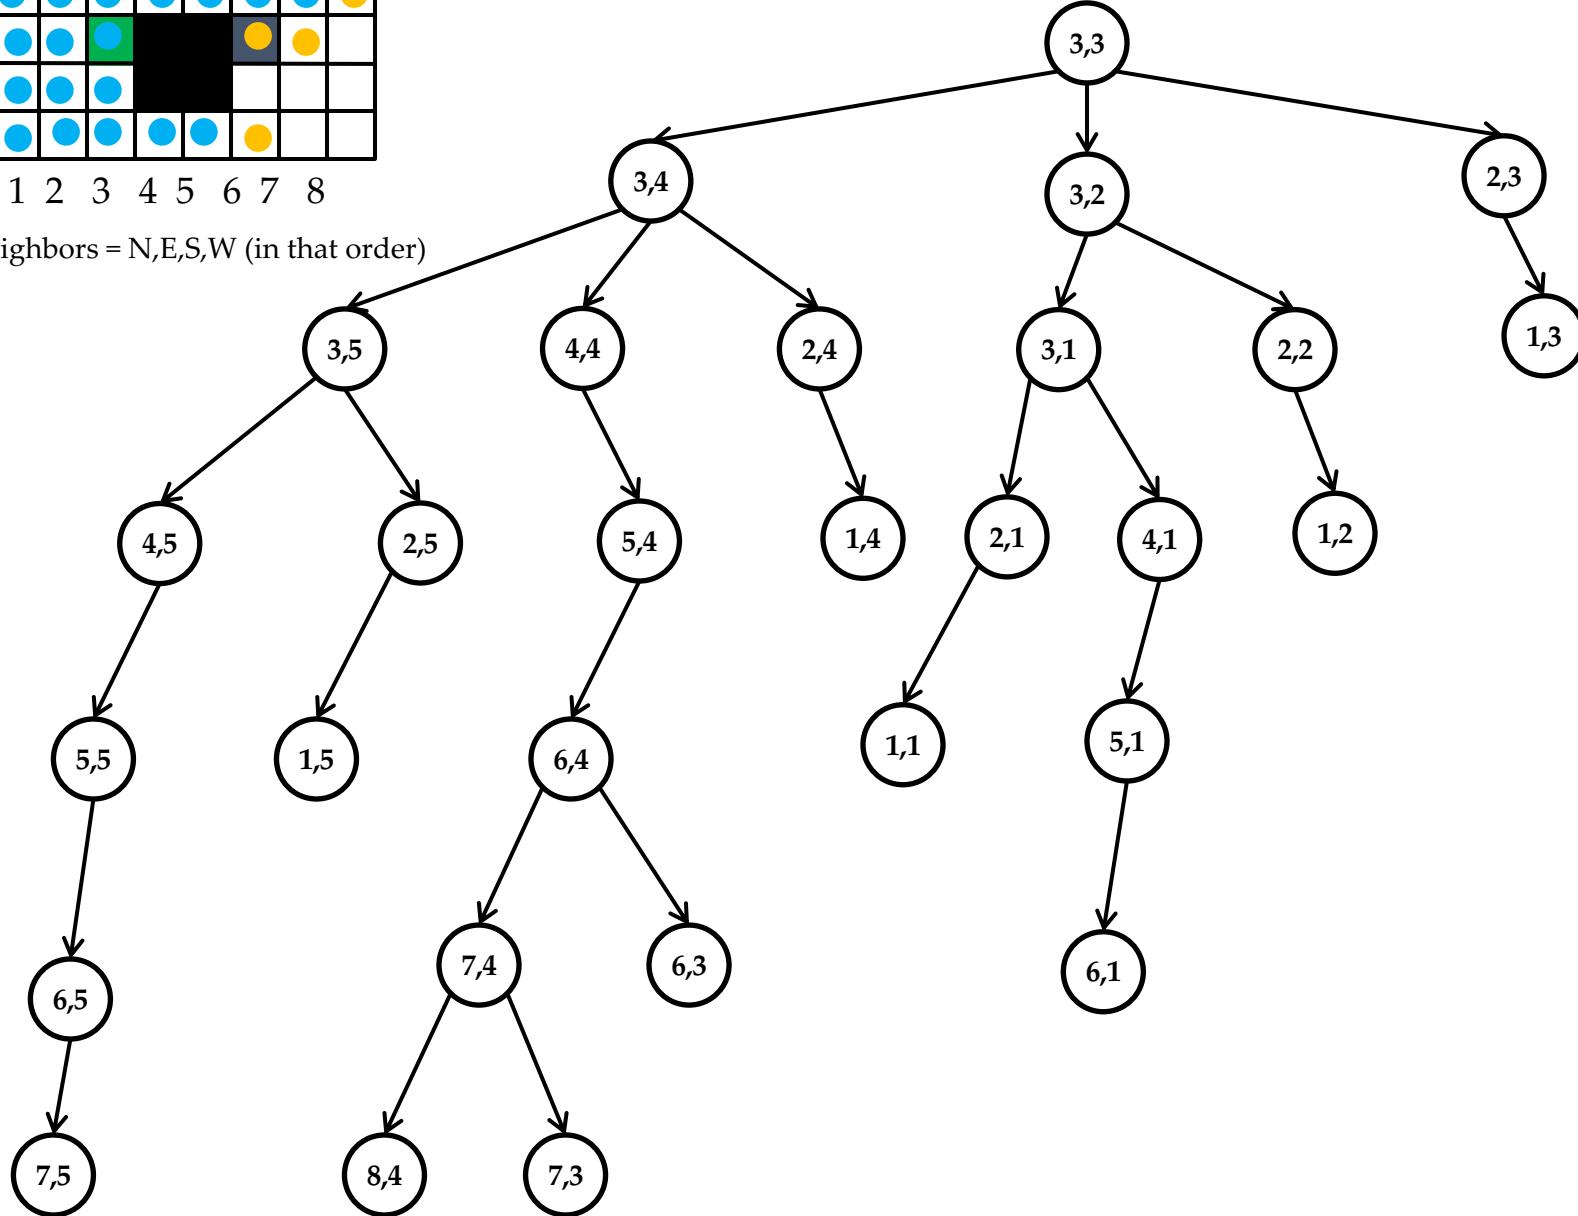
Neighbors = N,E,S,W (in that order)

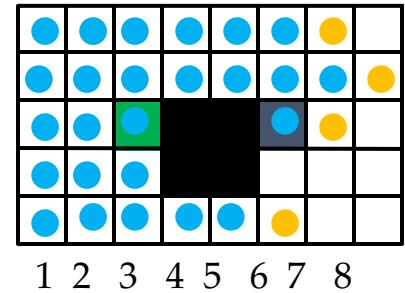




Search Tree: Breadth-first Search (no duplicates in fringe)

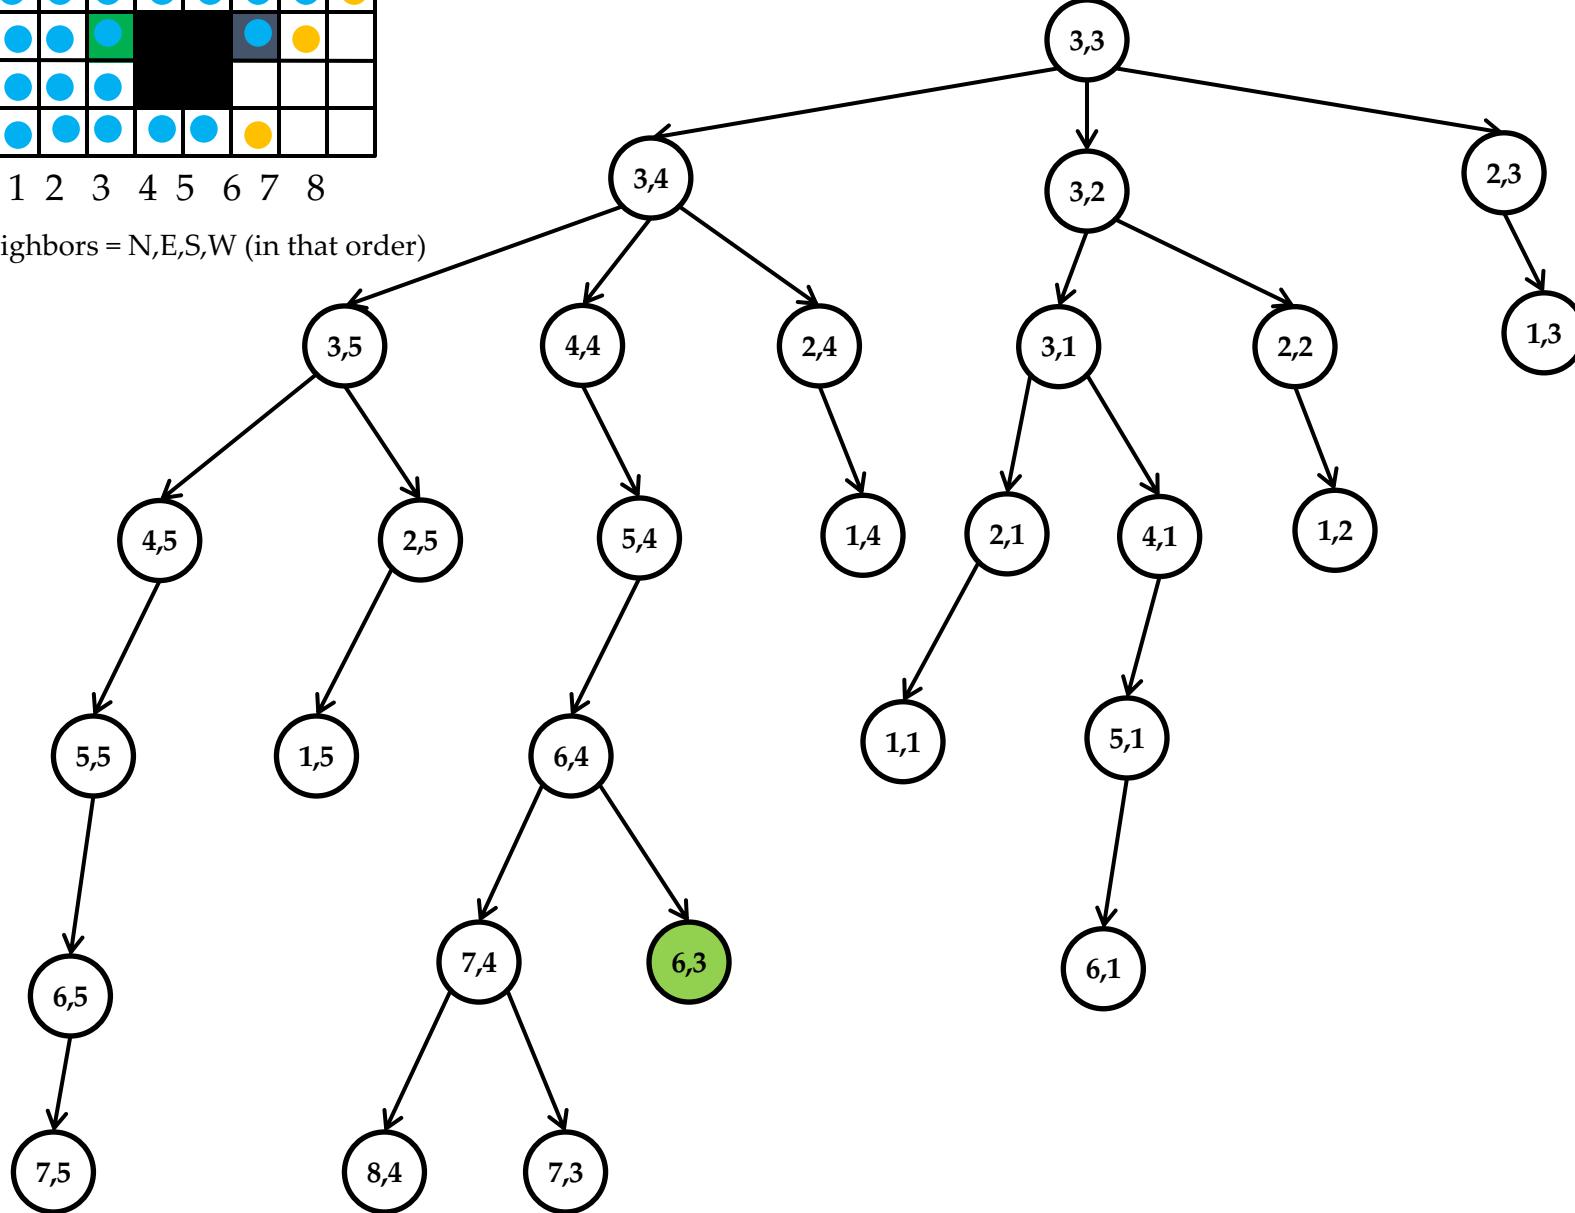
Neighbors = N,E,S,W (in that order)

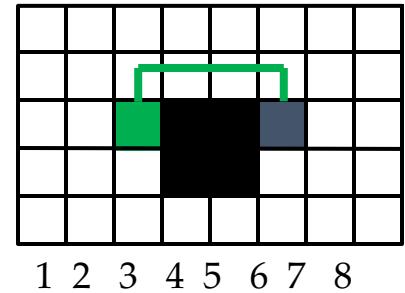




Search Tree: Breadth-first Search (no duplicates in fringe)

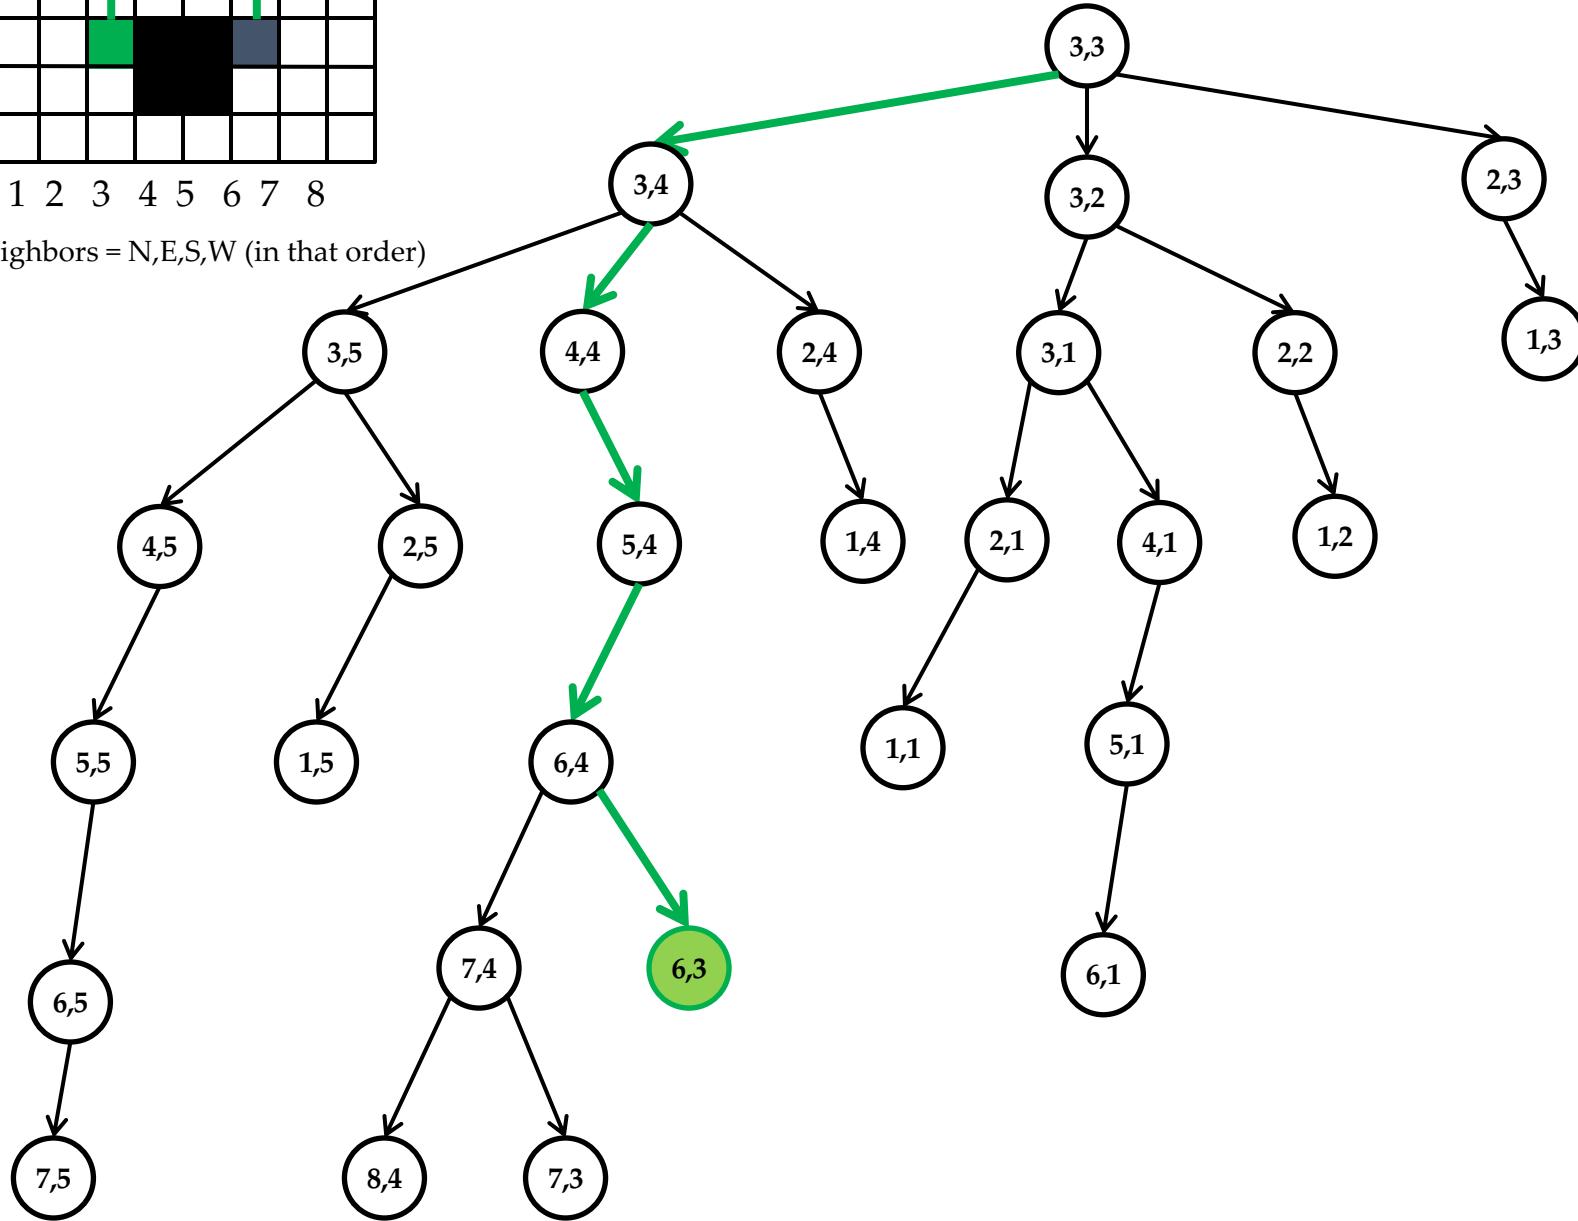
Neighbors = N,E,S,W (in that order)





Search Tree: Breadth-first Search (no duplicates in fringe)

Neighbors = N,E,S,W (in that order)



Uniform Cost Search (a.k.a. Dijkstra's Algorithm)

function UNIFORM-COST-SEARCH(*problem*) **returns** a solution, or failure

node \leftarrow a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0

frontier \leftarrow a priority queue ordered by PATH-COST, with *node* as the only element

explored \leftarrow an empty set

loop do

if EMPTY?(*frontier*) **then return** failure

node \leftarrow POP(*frontier*) /* chooses the lowest-cost node in *frontier* */

if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)

 add *node*.STATE to *explored*

for each *action* **in** *problem*.ACTIONS(*node*.STATE) **do**

child \leftarrow CHILD-NODE(*problem*, *node*, *action*)

if *child*.STATE is not in *explored* or *frontier* **then**

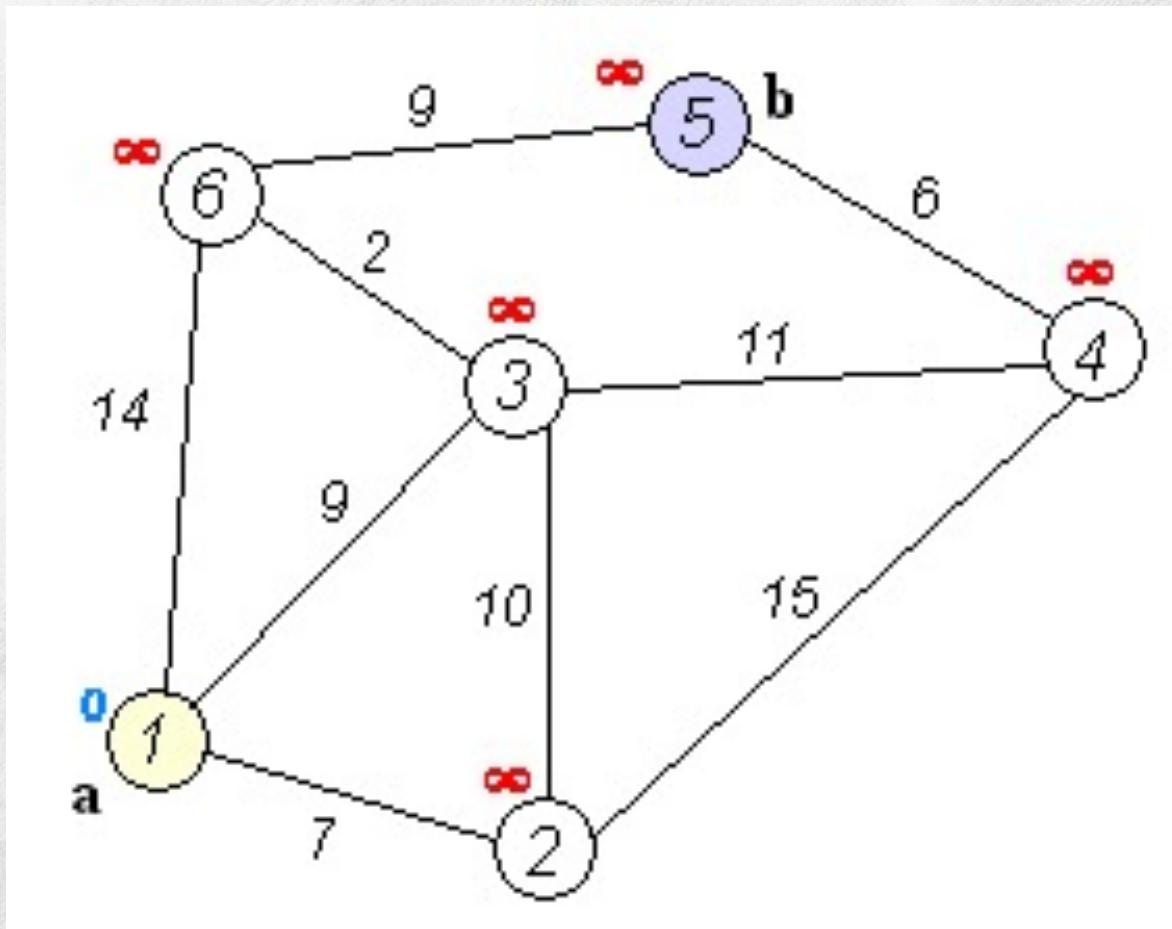
frontier \leftarrow INSERT(*child*, *frontier*)

else if *child*.STATE is in *frontier* with higher PATH-COST **then**

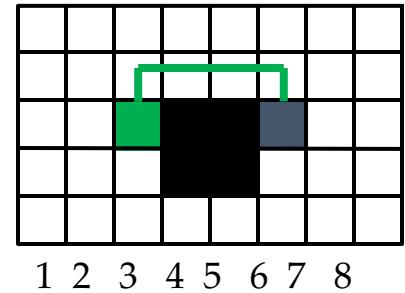
 replace that *frontier* node with *child*



Uniform Cost Search (a.k.a. Dijkstra's Algorithm)

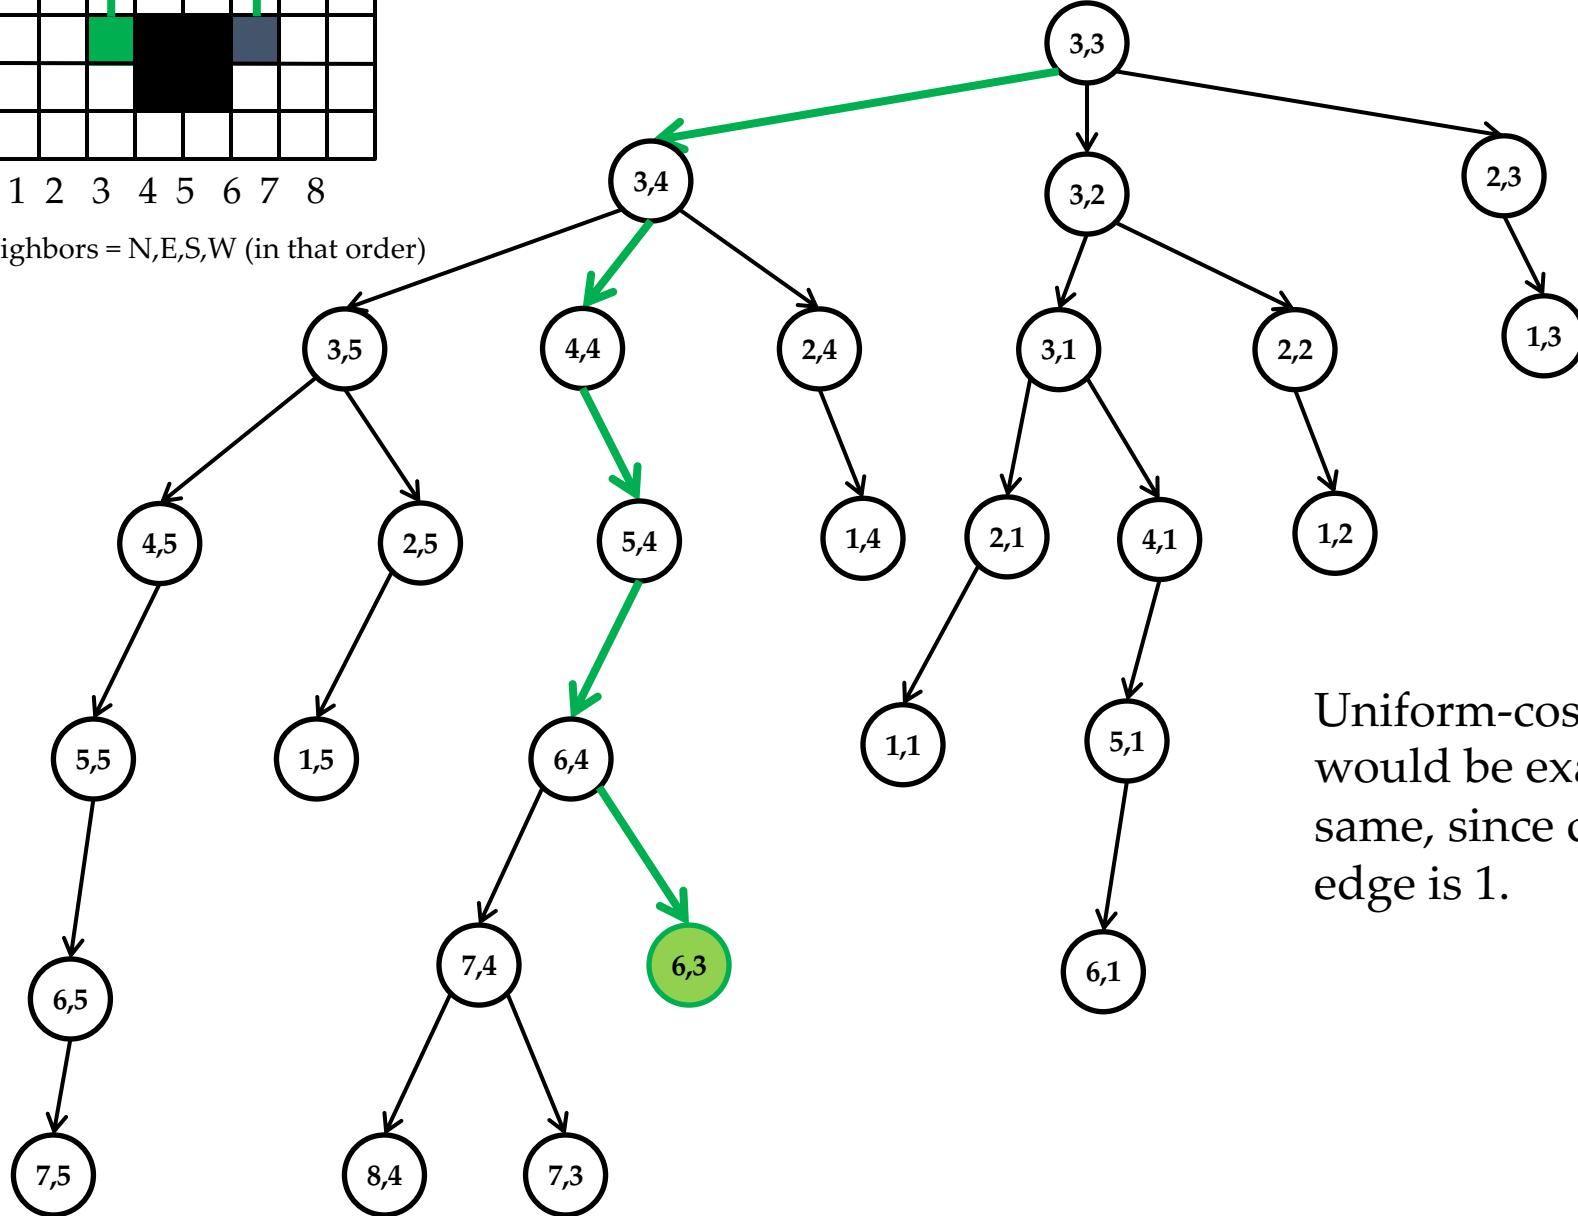


MISSISSIPPI STATE
UNIVERSITY™



Search Tree: Uniform-cost Search (no duplicates in fringe)

Neighbors = N,E,S,W (in that order)



Uniform-cost search would be exactly the same, since cost per edge is 1.

Informed Search



MISSISSIPPI STATE
UNIVERSITY™

Informed Search

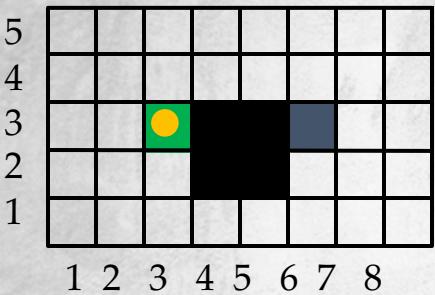
- To bias the search toward the goal we can tell the algorithm how close a state is to the goal, i.e. a *heuristic*
- A *heuristic* returns an estimate of the cost to get to the goal from any state
 - Must return 0 at a goal
- Most common heuristic for robotics – distance to goal (ignoring obstacles)
 - Manhattan distance: distance when only allowed to travel along the x or y axes
 - Euclidean distance: straight line distance



Informed Search

- FRINGE is a priority queue, with priority of each node calculated using the heuristic h
- A first idea: expand the node that is closest to the goal!
- **Greedy Best-First Search**
 - Priority = $h(node.state)$
 - Always expands the closest node to the goal
- Let's see how this works!





Search Tree: Greedy Best-First (no duplicates in fringe)

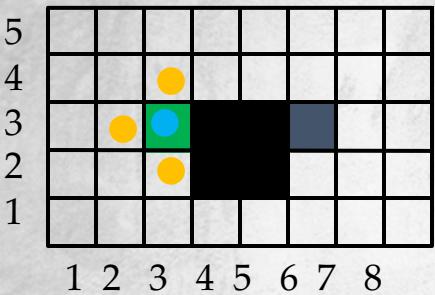
heuristic = manhattan distance



Neighbors = N,E,S,W (in that order)



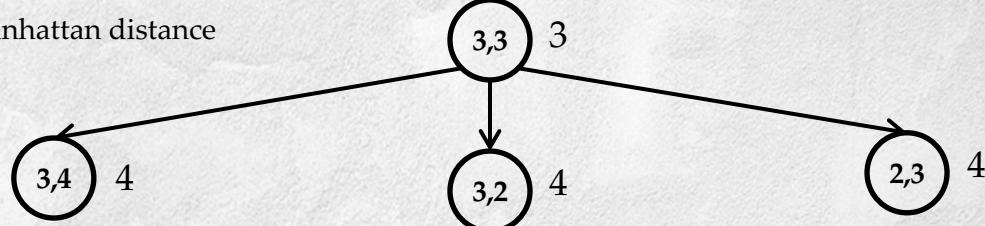
MISSISSIPPI STATE
UNIVERSITY™



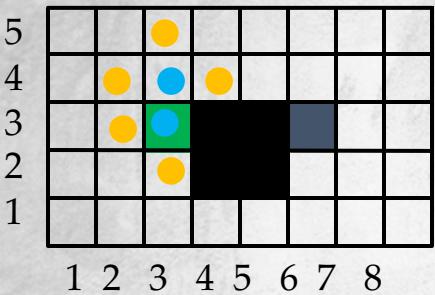
Neighbors = N,E,S,W (in that order)

Search Tree:Greedy Best-First (no duplicates in fringe)

heuristic = manhattan distance



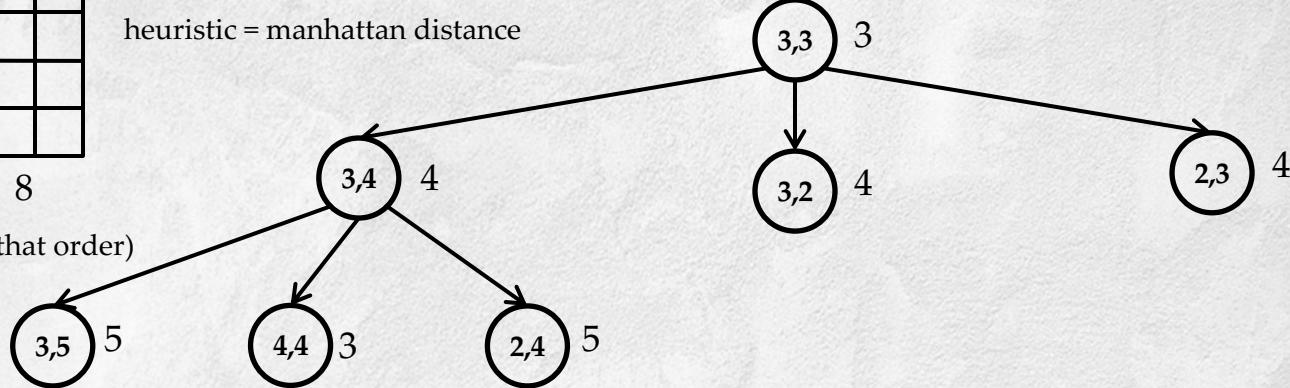
MISSISSIPPI STATE
UNIVERSITY™



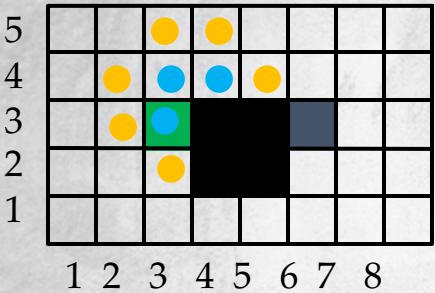
Search Tree: Greedy Best-First (no duplicates in fringe)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



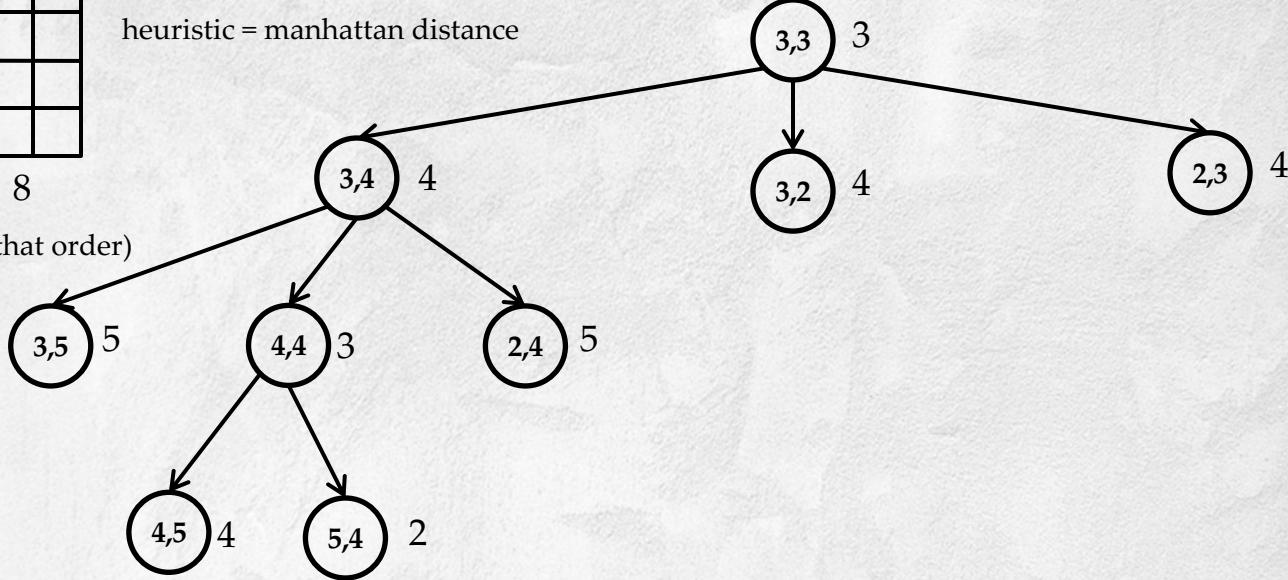
MISSISSIPPI STATE
UNIVERSITY™



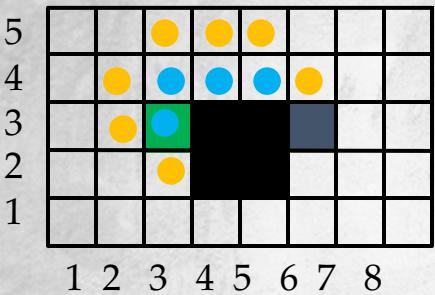
Search Tree: Greedy Best-First (no duplicates in fringe)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



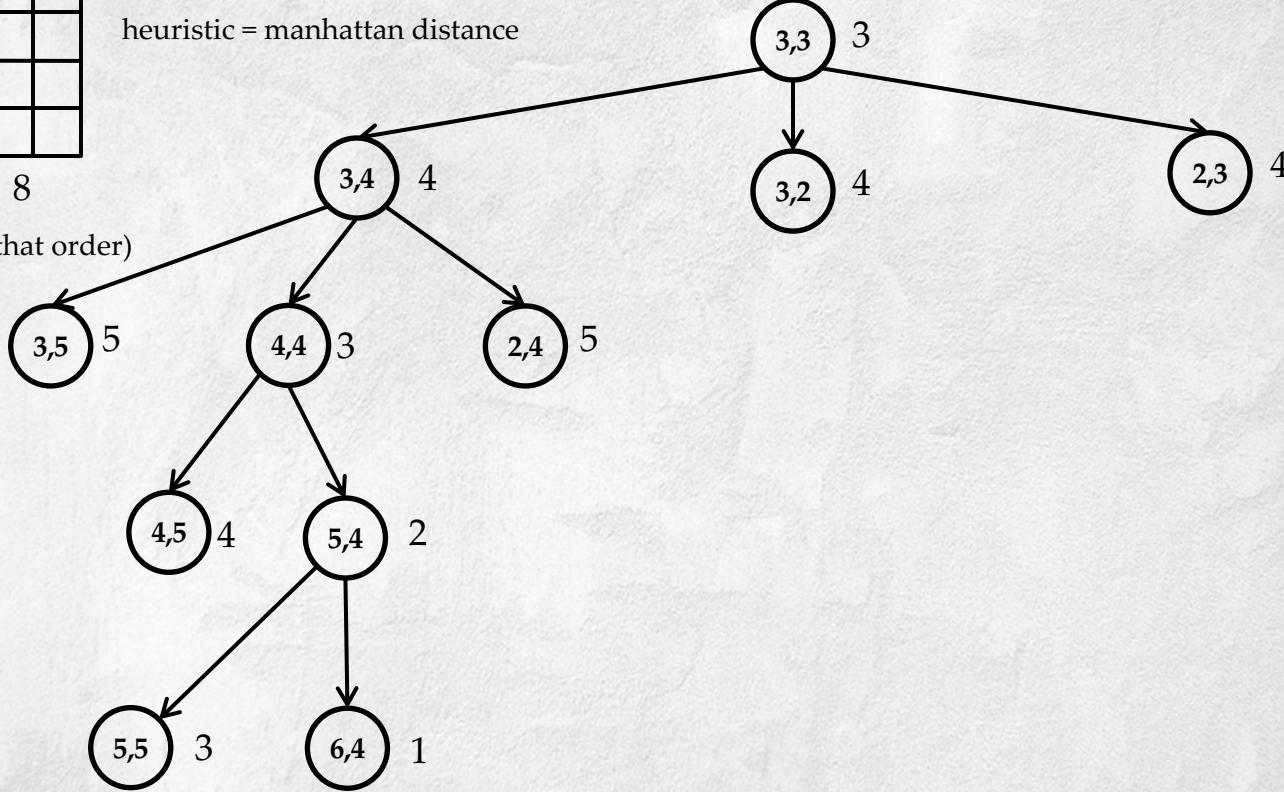
MISSISSIPPI STATE
UNIVERSITY™



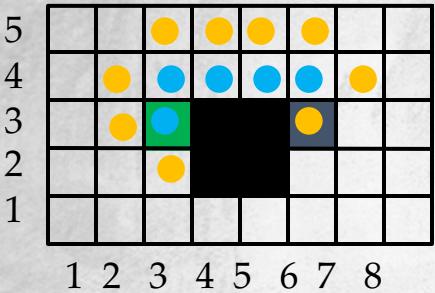
Search Tree: Greedy Best-First (no duplicates in fringe)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



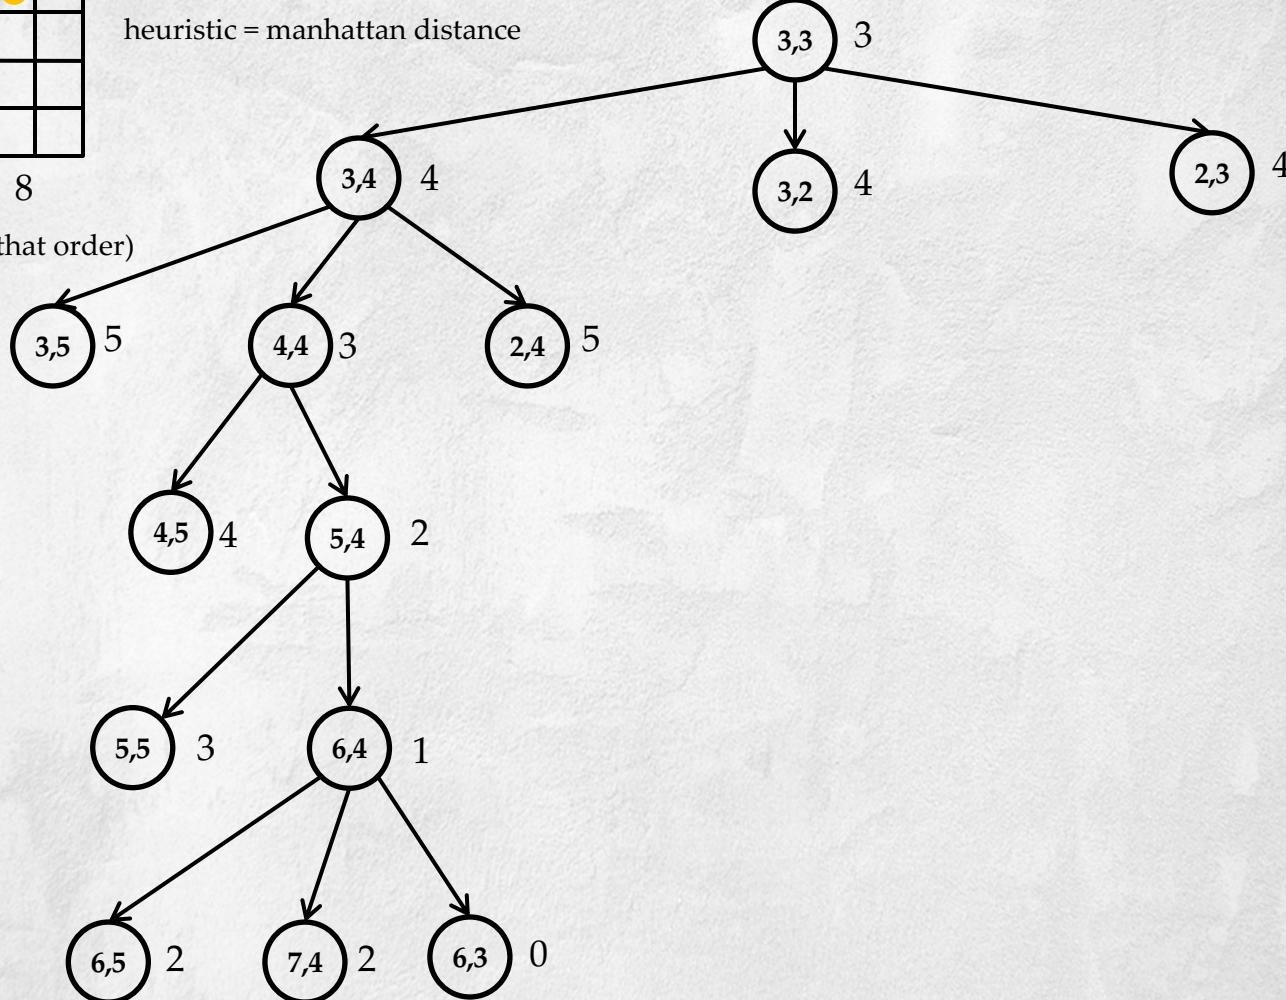
MISSISSIPPI STATE
UNIVERSITY™



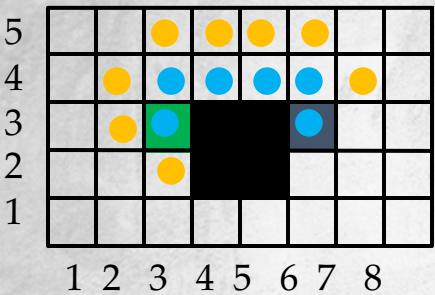
Search Tree: Greedy Best-First (no duplicates in fringe)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



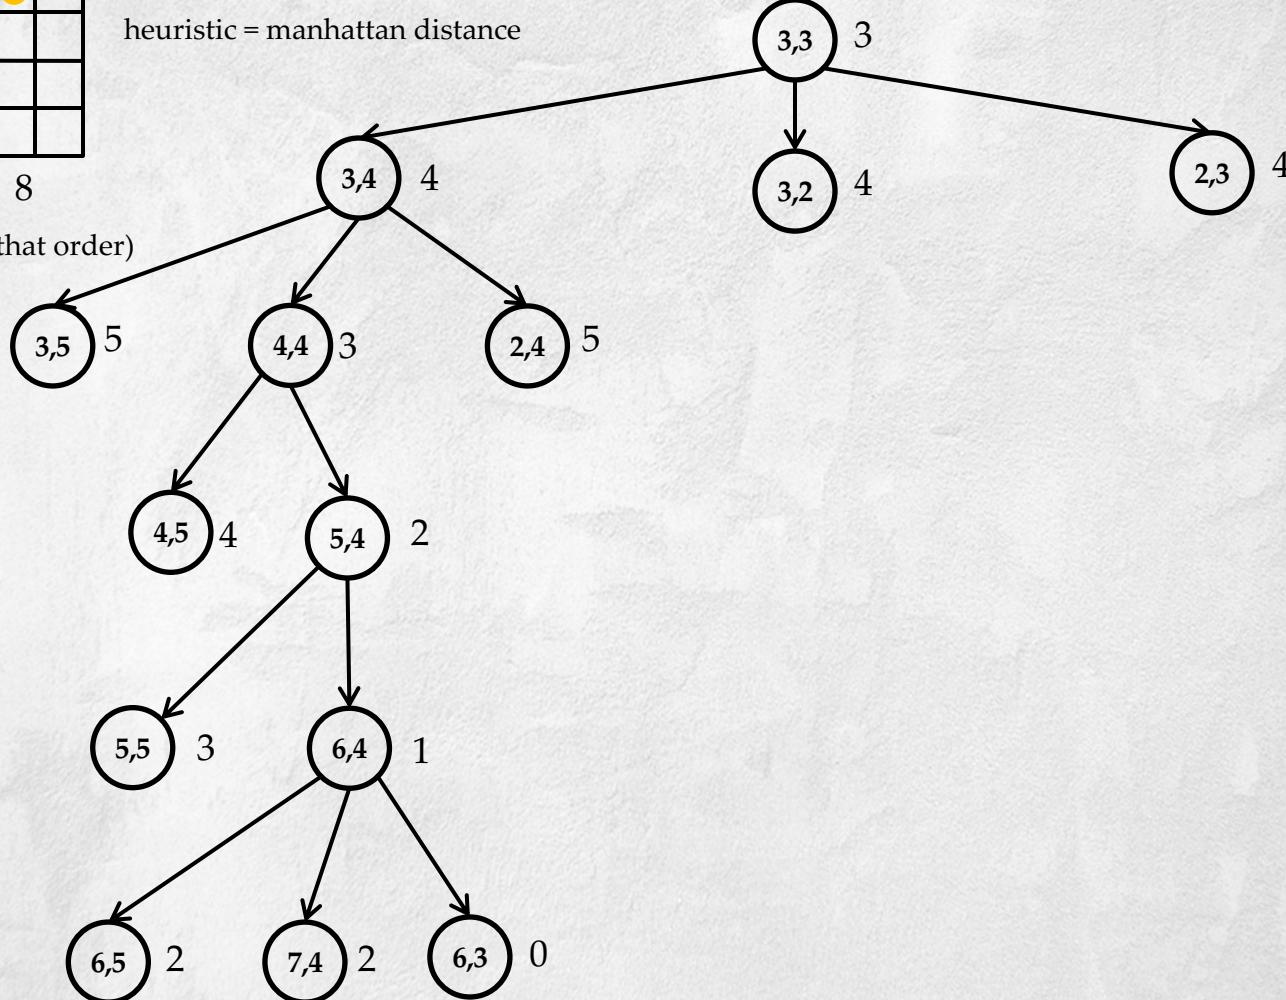
MISSISSIPPI STATE
UNIVERSITY™



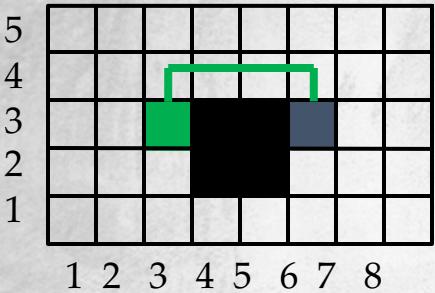
Search Tree:Greedy Best-First (no duplicates in fringe)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



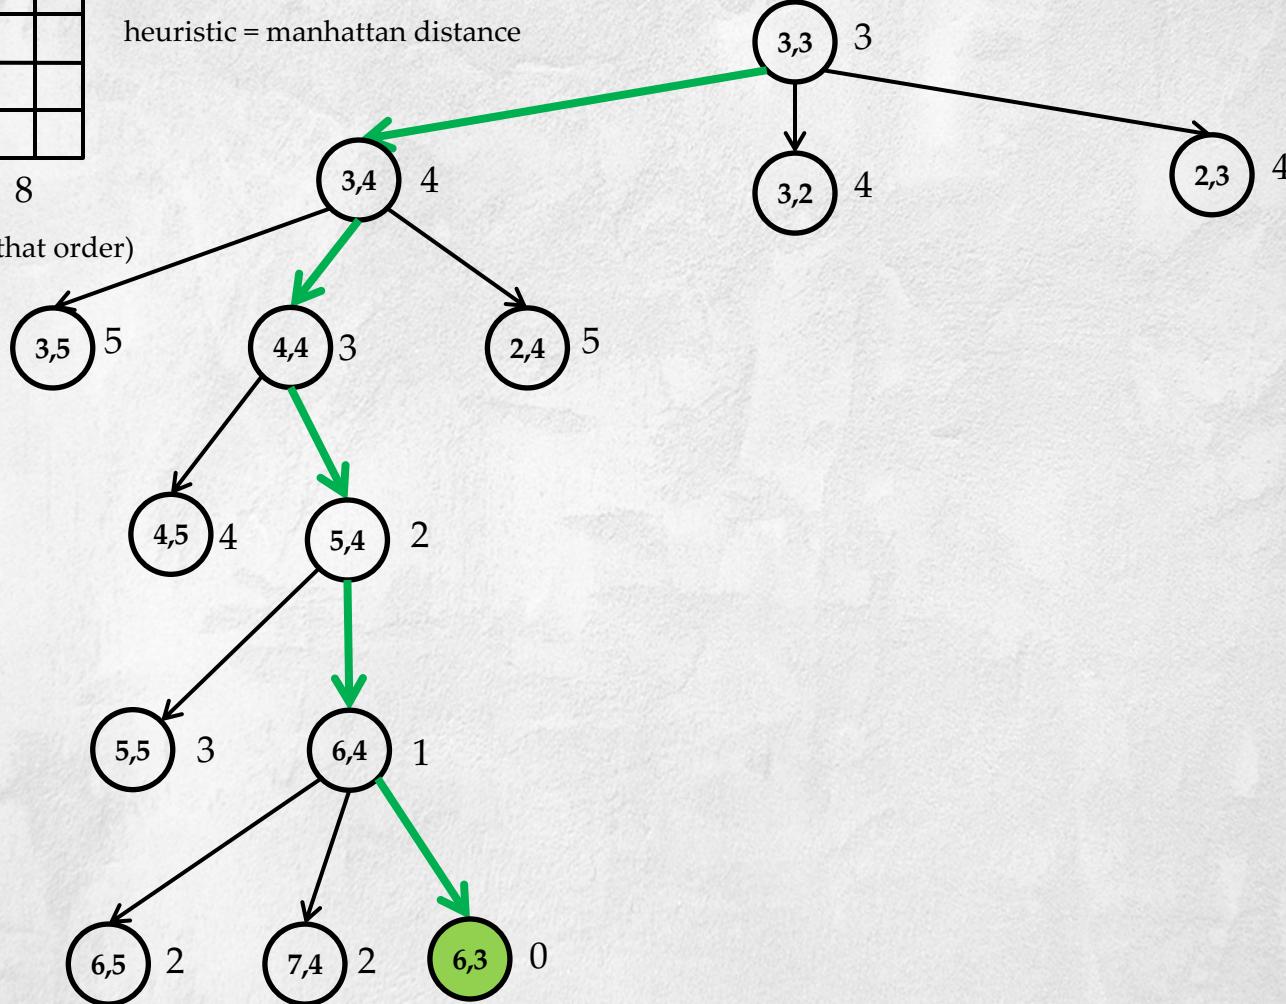
MISSISSIPPI STATE
UNIVERSITY™



Search Tree: Greedy Best-First (no duplicates in fringe)

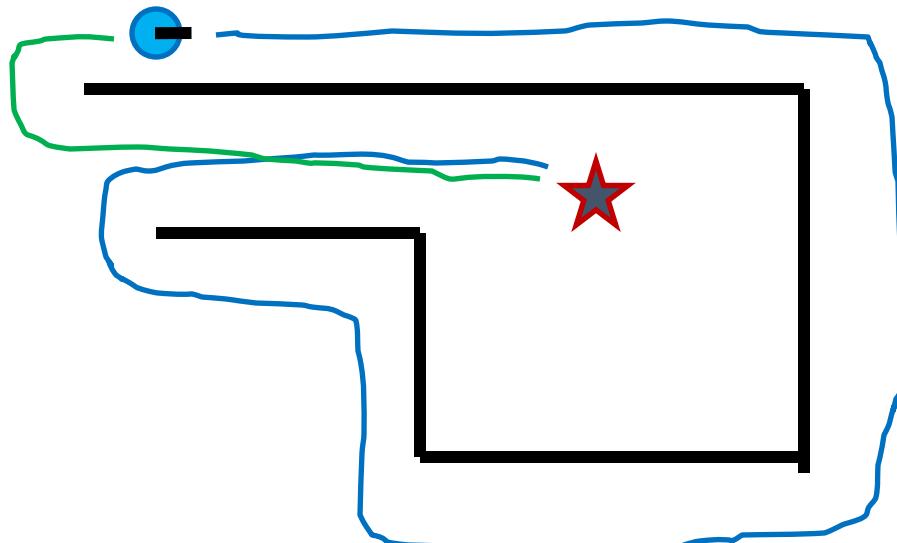
heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



MISSISSIPPI STATE
UNIVERSITY™

Another greedy example



Doesn't find the optimal solution



MISSISSIPPI STATE
UNIVERSITY™

Informed Search

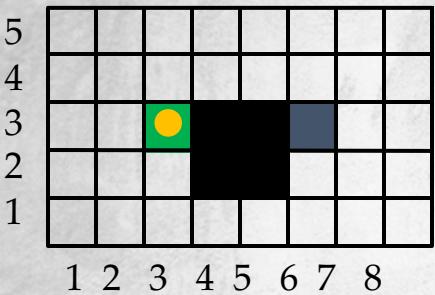
- FRINGE is a priority queue, with priority of each node calculated using the heuristic h
- **Greedy Best-First Search**
 - Priority = $h(node.state)$
 - Always expands the closest node to the goal
 - Optimal?
 - No, ignores cost of path. Could be too much
 - Properties?
 - Usually very fast, small search tree



Informed Search

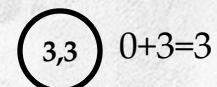
- FRINGE is a priority queue, with priority of each node calculated using the heuristic h
- A second idea: look at cost and heuristic
- **A* Search**
 - Priority = $node.cost + h(node.state)$
 - Takes into account both cost of path so far and cost to get to goal
 - Estimates *total path cost* for that partial path
 - Optimal?
 - Yes, if used with *admissible heuristic* (never overestimates cost to goal)





Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

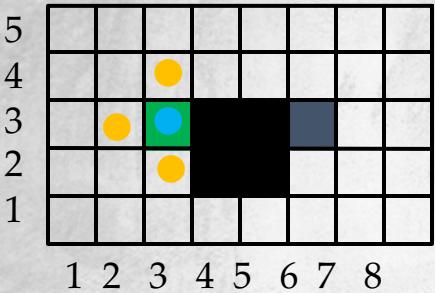
heuristic = manhattan distance



Neighbors = N,E,S,W (in that order)

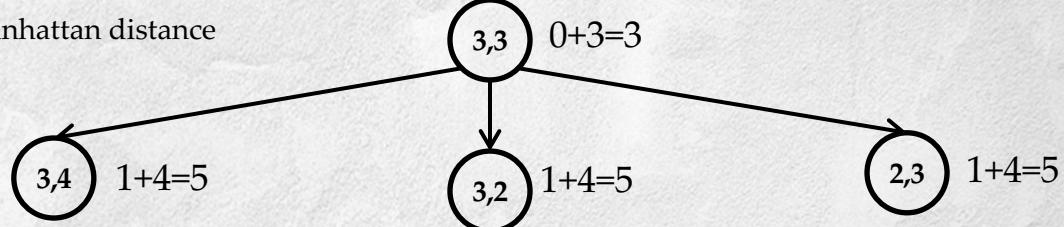


MISSISSIPPI STATE
UNIVERSITY™

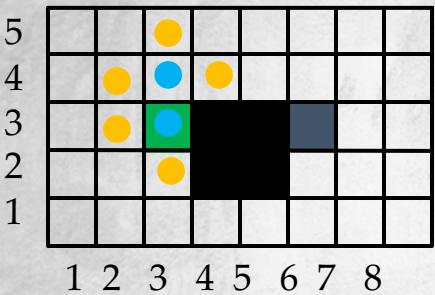


Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

heuristic = manhattan distance



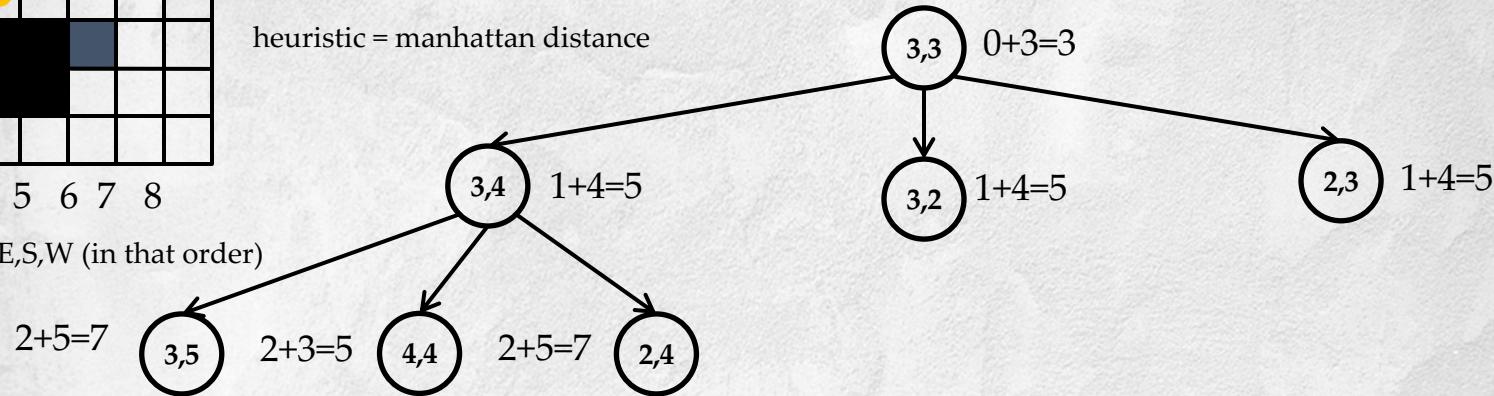
MISSISSIPPI STATE
UNIVERSITY™



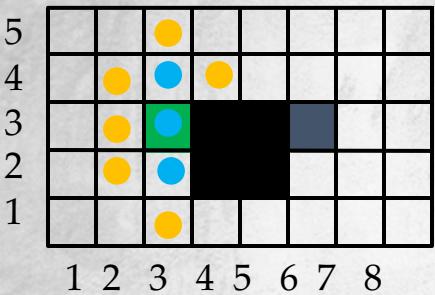
Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



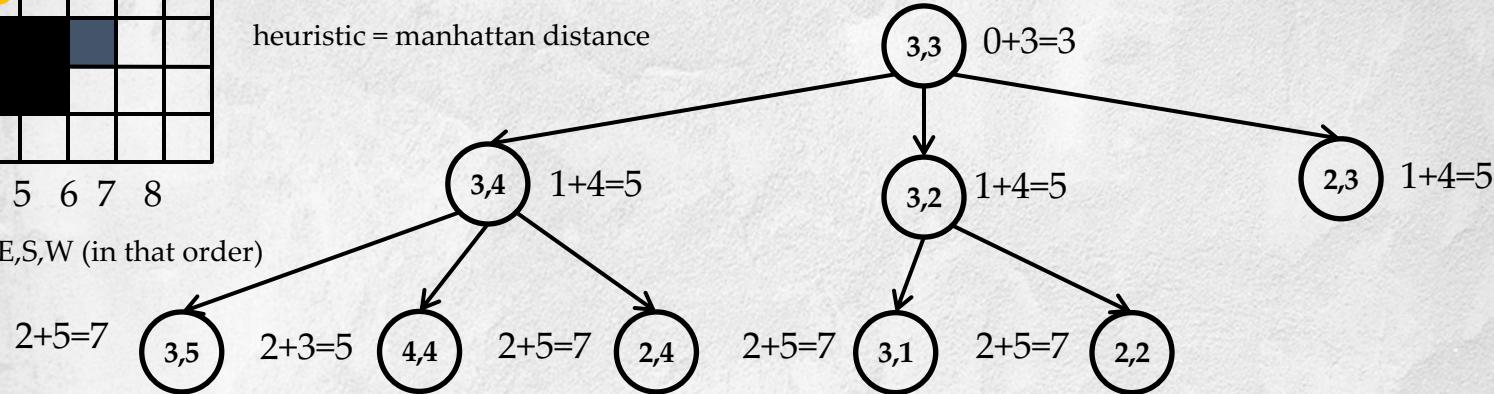
MISSISSIPPI STATE
UNIVERSITY™



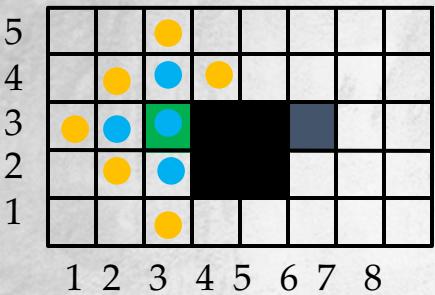
Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



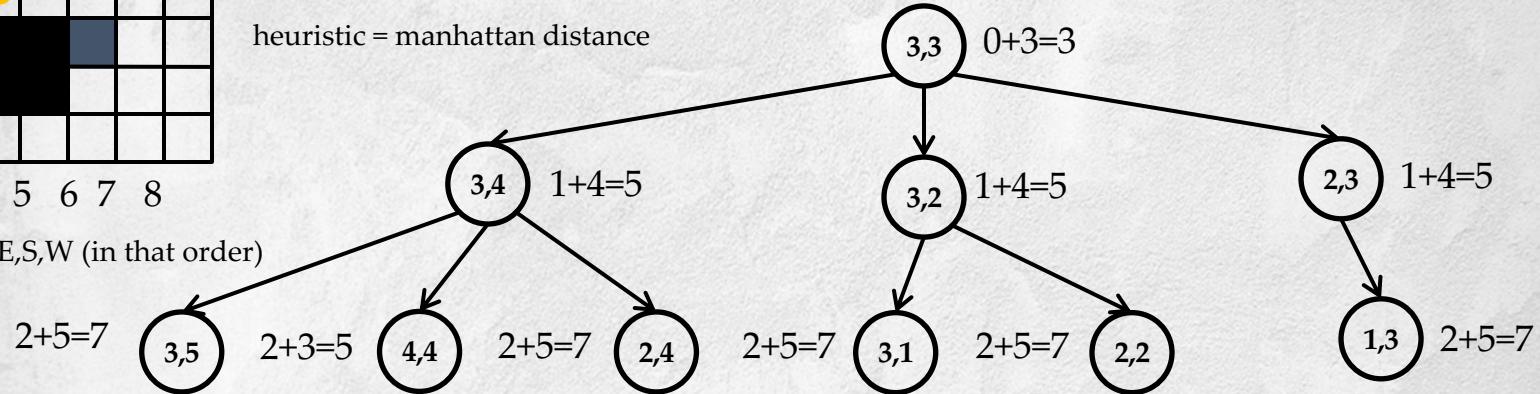
MISSISSIPPI STATE
UNIVERSITY™



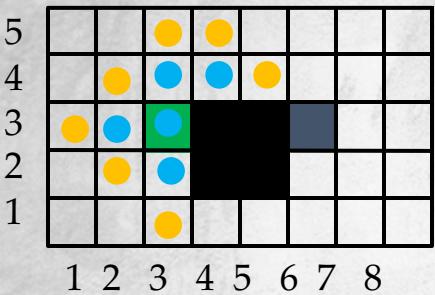
Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



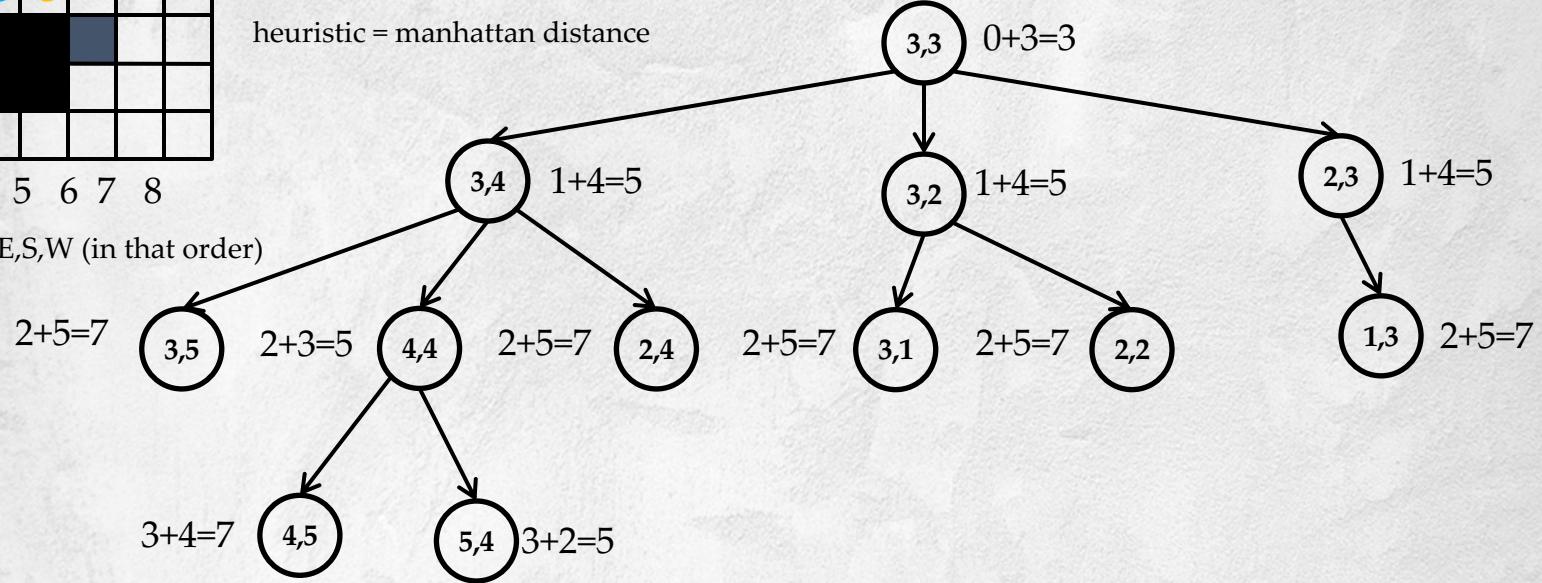
MISSISSIPPI STATE
UNIVERSITY™



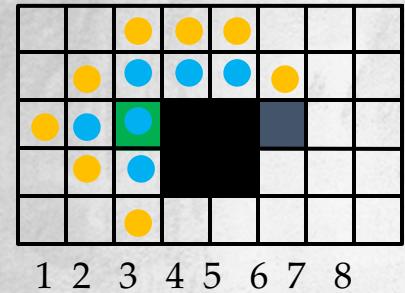
Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)

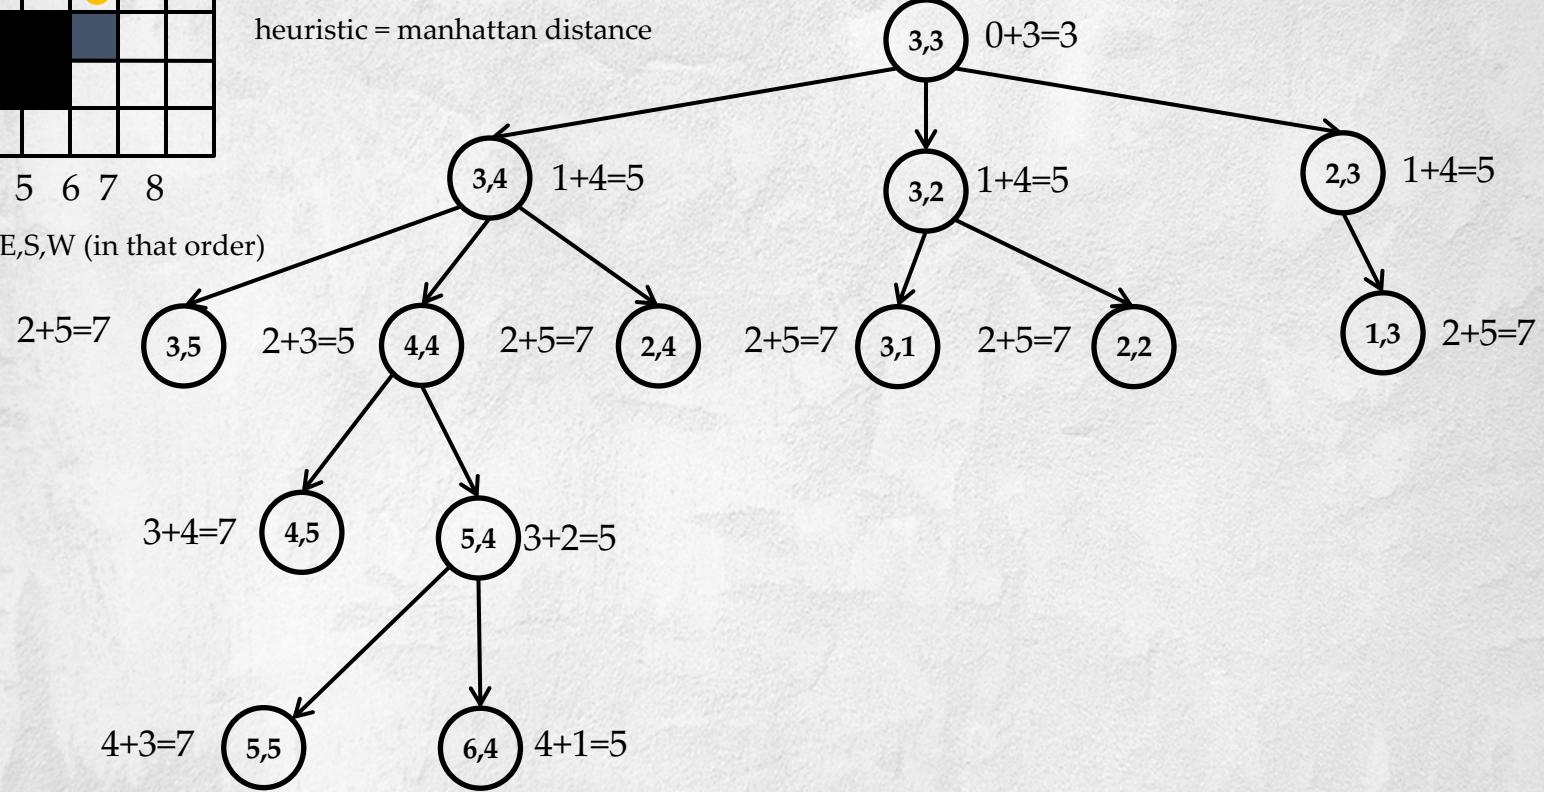


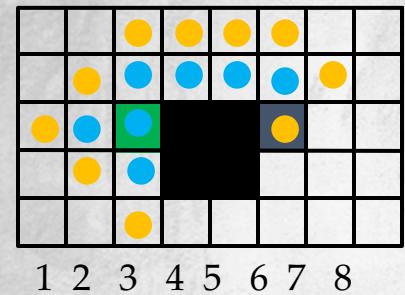
MISSISSIPPI STATE
UNIVERSITY™



Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

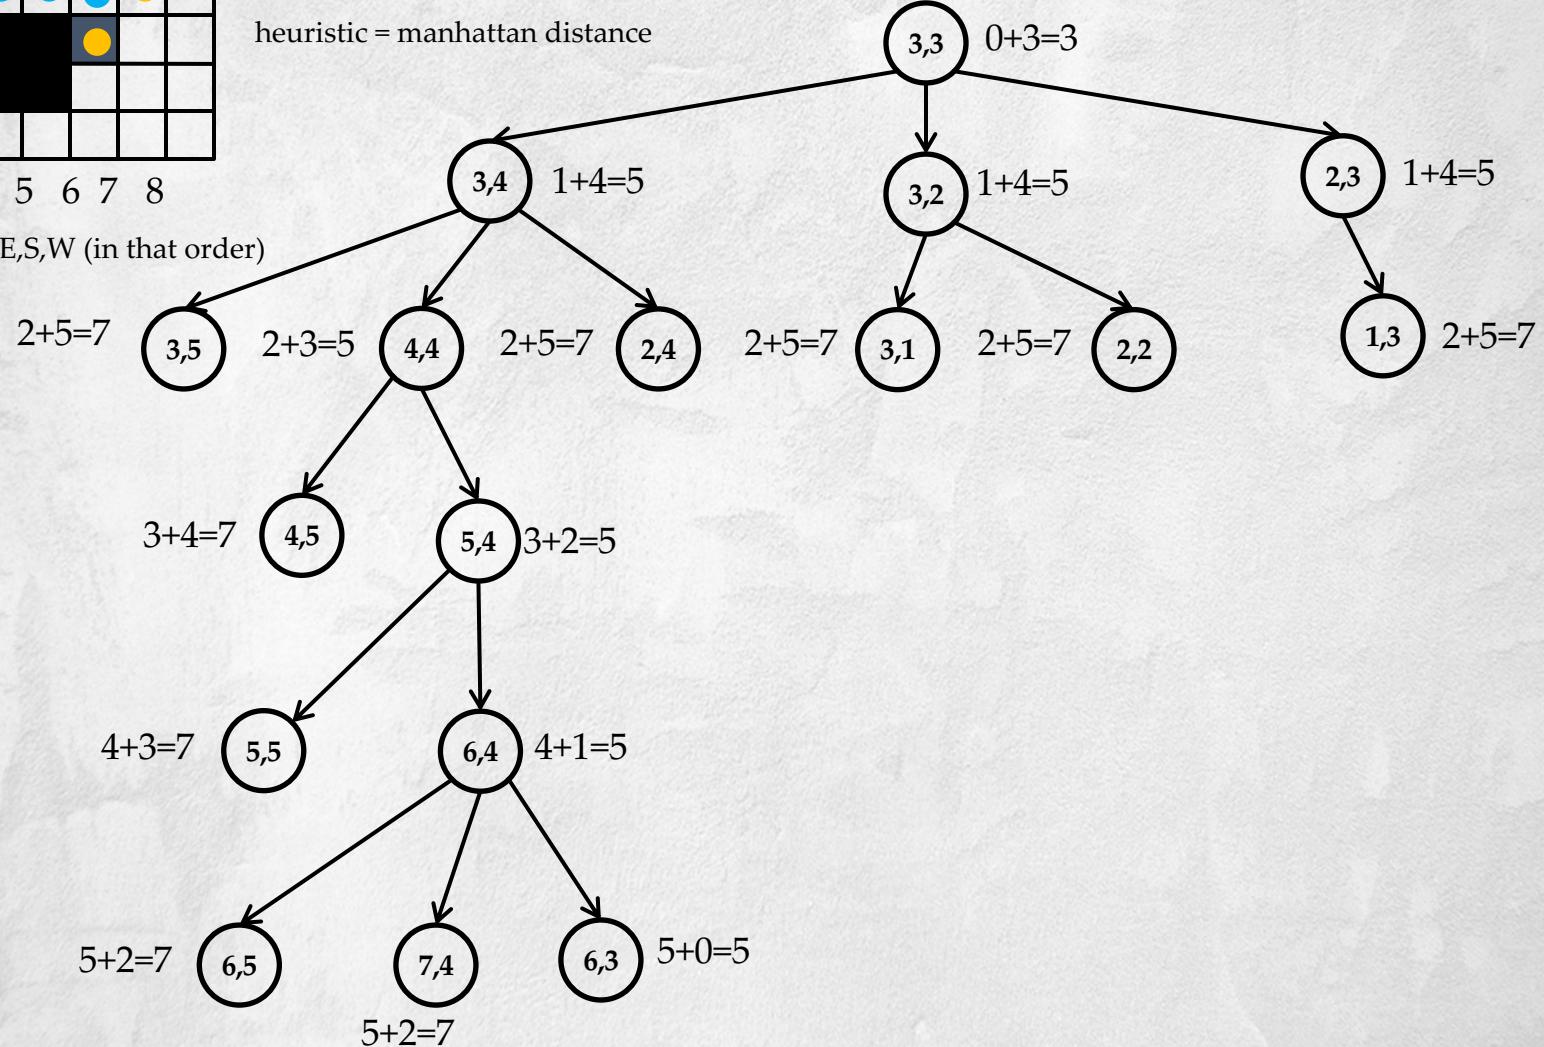
heuristic = manhattan distance

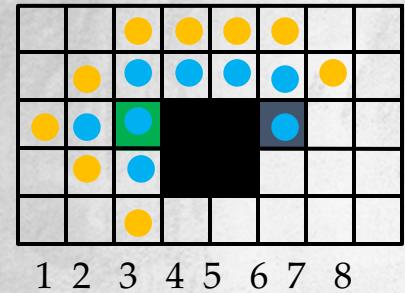




Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

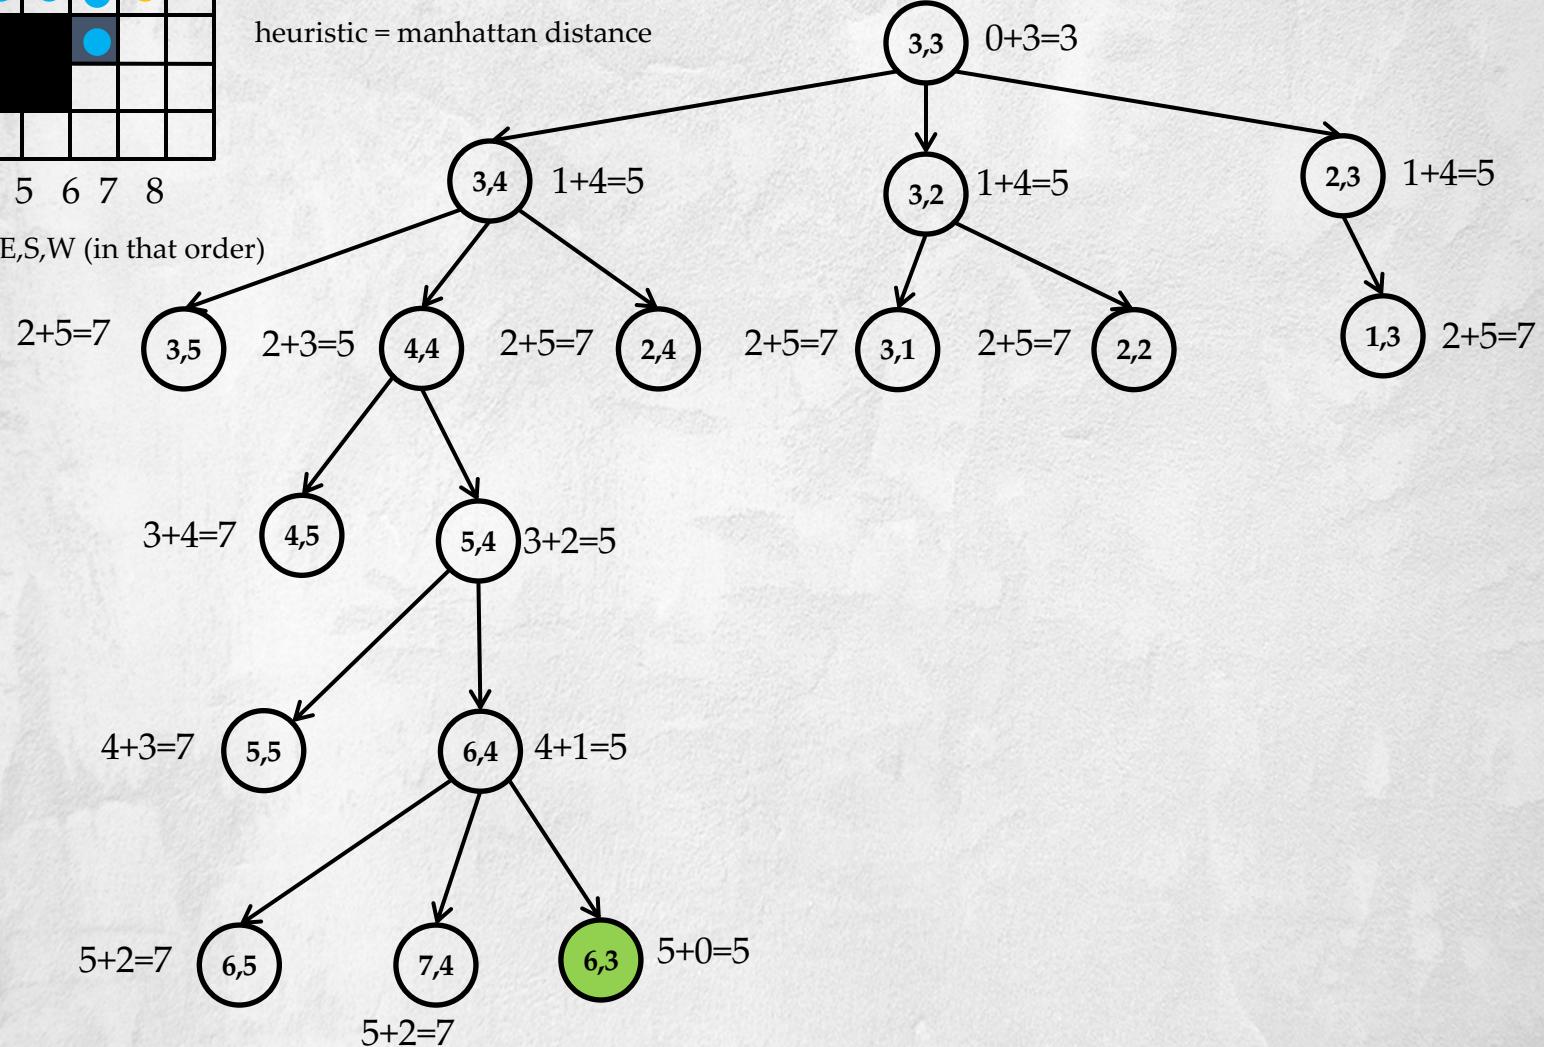
heuristic = manhattan distance

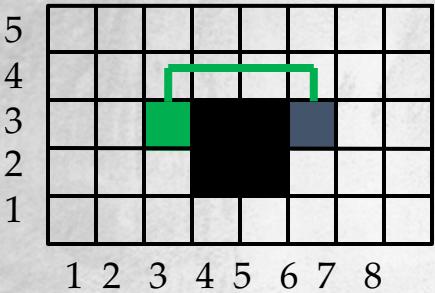




Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

heuristic = manhattan distance

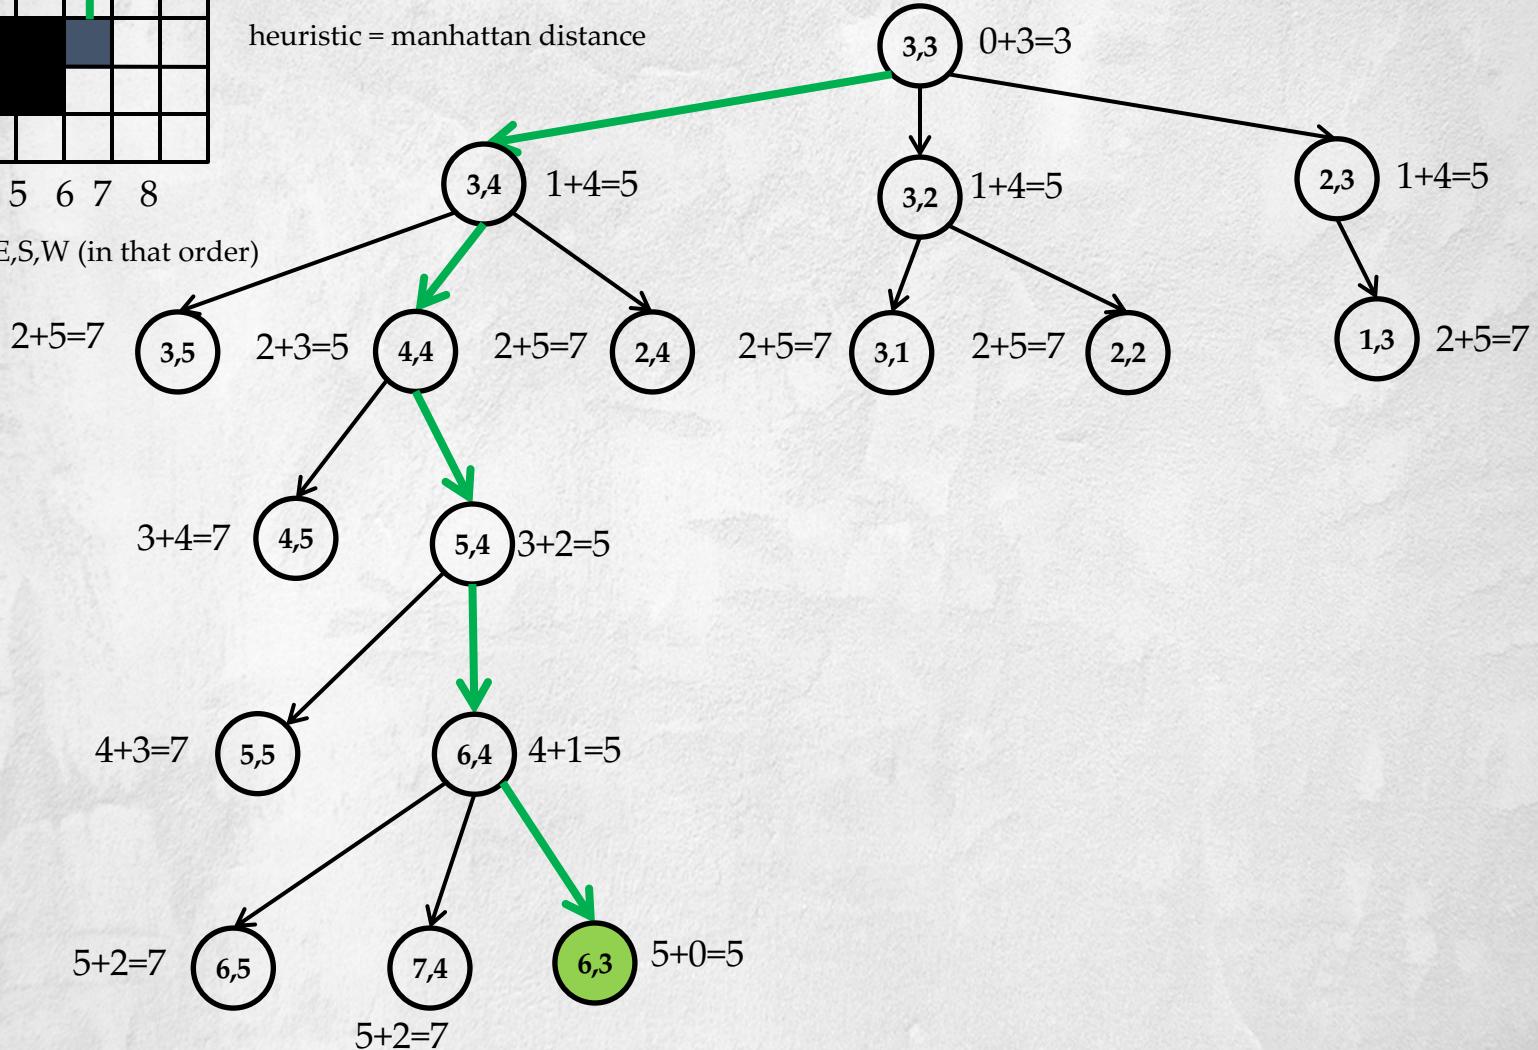




Search Tree: A* search (no duplicates in fringe) (FIFO tiebreaking)

heuristic = manhattan distance

Neighbors = N,E,S,W (in that order)



MISSISSIPPI STATE
UNIVERSITY™

Why is A* optimal?

- Could there be a shorter path in the tree?
 - Not if h doesn't overestimate cost to get to the goal
 - Cost + h is the cost of the best possible path to get to the goal along that partial path
 - If that is higher than the cost to get to the goal, then it can't be better



Analysis



MISSISSIPPI STATE
UNIVERSITY™

Uninformed Search

Criterion	Breadth-First	Uniform-Cost	Depth-First
Complete?	Yes ^a	Yes ^{a,b}	No
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$
Optimal?	Yes ^c	Yes	No

b : branching factor

d : depth of the shallowest solution

m : maximum depth of search tree

C^* : length of shortest path

ϵ : minimum edge length



Informed Search

Criterion	Greedy Best-First Search	A* Search
Complete?	No	Yes
Time	$O(b^m)$	$O(b^d)$
Space	$O(b^m)$	$O(b^d)$
Optimal?	No	Yes

b : branching factor

d : depth of the shallowest solution

m : maximum depth of search tree

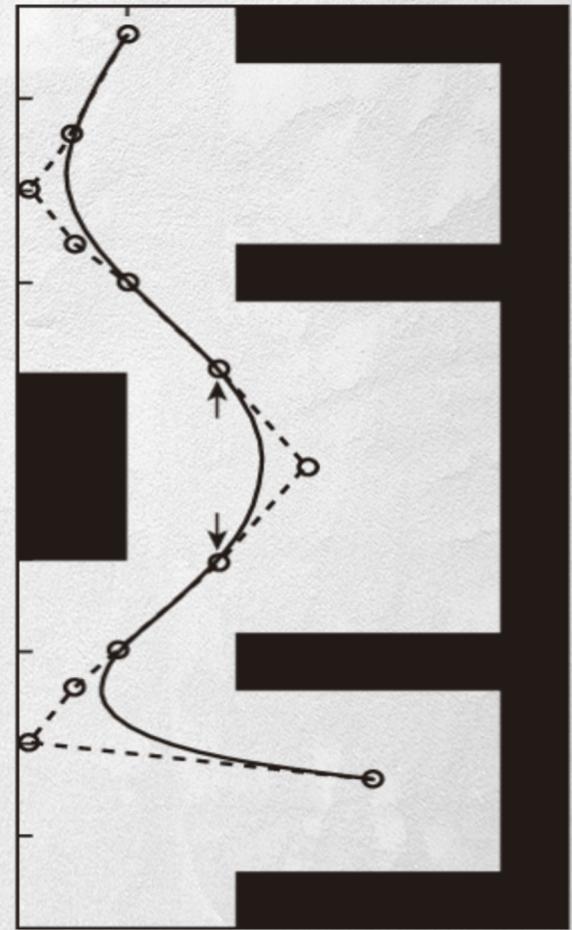
C^* : length of shortest path

ε : minimum edge length



Path Smoothing

- We can try to “fit a curve” to the points along the planned path
- Bezier curves work nicely
- *Curve may hit an obstacle where the straight line did not*



Velocity Profiling

- Accelerate as the robot goes along the path, then decelerate near the end
- Limit velocity when curvature is high
- For each segment of the path, calculate distance and curvature:
 - Iterate forwards and set velocity at each point based on the max velocity given acceleration (from kinematics) and curvature
 - Do the same thing in reverse to handle deceleration

Fig. 13: Velocity Pass

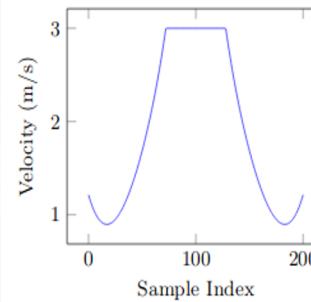


Fig. 14: Forward Pass

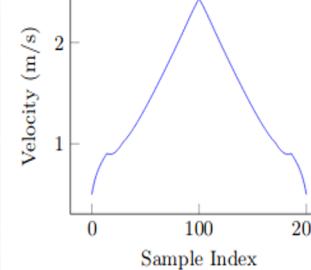
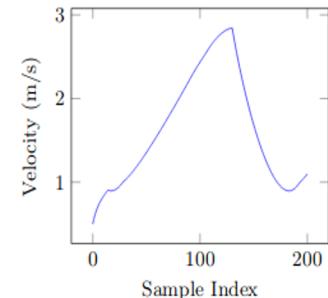


Fig. 15: Backward Pass

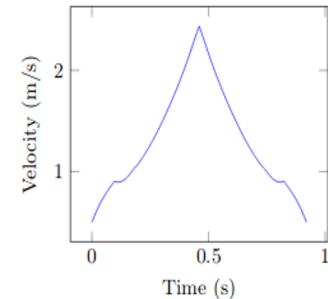


Fig. 16: Velocity vs. Time



References

1. <https://cs.stanford.edu/people/eroberts/courses/soco/projects/1998-99/robotics/basicmotion.html>
2. <https://www.geeksforgeeks.org/search-algorithms-in-ai/>
3. <https://medium.com/kredo-ai-engineering/search-algorithms-part-1-problem-formulation-and-searching-for-solutions-28f722b7a1a6>

